

CCPS 209 Computer Science II

Lecture Notes

Ilkka Kokkarinen

Chang School of Continuing Education
Toronto Metropolitan University

Version of April 14, 2023

<i>Lecture 1: Imperative Java for Pythonistas</i>	7
1.1. Syntax and structure of a Java program	7
1.2. Classes as wrappers for functions	7
1.3. Control structures	8
1.4. Primitive types	9
1.5. Arrays	10
1.6. Iterating through one or more dimensions	11
<i>Lecture 2: Abstraction and Encapsulation</i>	13
2.1. Classes and objects	13
2.2. Instance fields to represent and store properties	14
2.3. More on instance fields, static fields, and local variables	15
2.4. Constructors	16
2.5. Primitive types versus references	17
2.6. References as method arguments	18
2.7. Immutable versus final	20
2.8. The implicit reference this	20
2.9. Method overloading	21
<i>Lecture 3: Inheritance and Polymorphism</i>	23
3.1. Types in programming languages	23
3.2. The need for inheritance to avoid repeated code	24
3.3. Inheritance between classes	25
3.4. Overriding inherited methods	26
3.5. Polymorphism and dynamic binding	27
3.6. Abstract superclasses and methods	28
3.7. Visibility of class members	28
3.8. The universal superclass Object	29
3.9. Overriding the equals method	30

3.10. Dynamic type inference	31
3.11. Decorator pattern	32
3.12. Constructors and inheritance	33
3.13. Overloading vs. overriding: technical interlude	34
Lecture 4: The Java Collection Framework	35
4.1. Interfaces as perfectly abstract superclasses	35
4.2. The Java Collection Framework	36
4.3. Operations on collections	37
4.4. Iterators	37
4.5. Strategy pattern	38
4.6. Java Collection Framework as a teachable moment	39
4.7. Association maps	40
Lecture 5: Swing	42
5.1. AWT and Swing	42
5.2. Graphics rendering	43
5.3. Nested classes	44
5.4. Object factories	45
5.5. Event listeners	46
5.6. Turtle graphics (extra material)	46
5.7. Turtles for the new millennium (extra material)	47
Lecture 6: Exceptions	49
6.1. Pre- and postconditions	49
6.2. Liskov Substitution Principle	50
6.3. The exception machinery	52
6.4. Catch or release	52
6.5. Checked and unchecked exceptions	53
6.6. The Throwable class hierarchy	54

6.7. Assertions	54
6.8. Some exception patterns and anti-patterns	55
Lecture 7: I/O Streams	57
7.1. Streams and their decorators	57
7.2. Readers and Writers	58
7.3. Dependencies between classes	59
7.4. Reflection	60
7.5. Serialization of arbitrary object structures	61
Lecture 8: Generics And Other Nifty Features	63
8.1. Minor language improvements back in Java 5	63
8.2. Autoboxing	63
8.3. Annotations	64
8.4. Basic generics	65
8.5. Erasure and its consequences	66
8.6. More on raw types and reification	67
8.7. Inheritance relationships over generic types	67
8.8. Bounded types	68
8.9. Enumerated types	69
Lecture 9: Concurrency Fundamentals	71
9.1. Concurrent execution of threads in JVM	71
9.2. Java execution model	72
9.3. Threads in Java	73
9.4. Critical sections and race conditions	73
9.5. Mutual exclusion locks	74
Lecture 10: Concurrency Controls	76
10.1. Condition variables	76
10.2. Crowding inside a condition variable	76

<i>10.3. Fairness in concurrency</i>	77
<i>10.4. Asking another thread to terminate</i>	78
<i>10.5. Semaphores</i>	79
<i>10.6. Additional concurrency tidbits</i>	79
<i>10.7. Concurrency and Swing</i>	80
<i>Lecture 11: Computation Streams and Lambdas</i>	82
<i>11.1. The increasingly important laws of Moore and Amdahl</i>	82
<i>11.2. Fork-Join framework</i>	83
<i>11.3. Lambdas as syntactic sugar for anonymous classes</i>	84
<i>11.4. Functional programming</i>	85
<i>11.5. Meeting in the middle of the two paradigms</i>	86
<i>11.6. Computation streams</i>	87
<i>11.7. Lazy evaluation</i>	88
<i>Bonus Lecture 12: Integers and Floats in Java</i>	90
<i>12.1. Positional number systems</i>	90
<i>12.2. Bits and bytes</i>	91
<i>12.3. Signed integers and integer arithmetic</i>	91
<i>12.4. Bitwise arithmetic</i>	92
<i>12.5. Floating point</i>	93
<i>12.6. IEEE 754 floating point numbers</i>	94
<i>12.7. Imprecision of floating point</i>	94
<i>12.8. Floating point arithmetic</i>	95

Lecture 1: Imperative Java for Pythonistas

1.1. Syntax and structure of a Java program

- Unlike in Python whose syntax is much more freeform, every Java program must consist of **classes**, without loss of generality written one class per file. Class names conventionally start with a capital letter, whereas method and variable names start with lowercase characters. When a name consists of several words, **camelcase** is used to improve readability.
- The source code for the class `Foo` must be in the file named exactly `Foo.java`.
- Java source code must be **explicitly compiled** into equivalent Java **bytecode** that the **Java Virtual Machine** can directly execute, unlike in Python where this conversion is done silently when the script is read in before execution. The source code for the class `Foo` is compiled into the machine-readable bytecode file `Foo.class`.
- Java was the first major programming language designed for **portability and security** for the World Wide Web. Java bytecode is executed inside a **virtual sandbox** that prevents malicious programs from accessing the actual computer that simulates the imaginary computer that runs Java applications in the form of the Java Virtual Machine.
- The original grand idea of Java was that any web page can contain **applets**, little executable Java programs that are part of that web page. Due to various screw ups in the nineties and later, these days Java is dead as a doornail in this niche, and has been replaced by **JavaScript** (no relation, despite the name) so that today's web browsers will just flat out refuse to execute and display Java applets at all.
- Java still chugs on in the **back end services**, and also faces the audience from inside **Android**, even though even over there it has since been officially replaced with its backwards-compatible but more Pythonic successor **Kotlin**.

1.2. Classes as wrappers for functions

- Java classes are organized into hierarchical **packages**, some of which comprise the **Java standard library**. An `import` statement is used to make names of the given package visible in the current compilation unit.
- Every class defines a new **data type**. Classes may contain **field declarations**, **methods**, **nested classes** and **initializer blocks** inside them to define how this data type behaves.
- In this first module, we look only at classes that implement data types that are not intended to be used to create new objects, but are syntactic containers for `static` methods that correspond to **functions** inside a Python script.
- Whereas Python relies on the convention of **underscores** to begin some name to document that that name is not supposed to be accessed from the outside, Java enforces proper **encapsulation** at the compile time with keywords `public` and `private`. The `private` methods do not even show up outside of the class, and can be accessed only from the methods in the very same class.

- The Java syntax is case sensitive but **insensitive to whitespace**, even though the proper white-space hygiene should be followed to clarify the program structure to the human reader. Structural nesting of code is indicated by **curly braces** that delimit the **body** of that structure.
- Both Java and Python are **strongly typed** languages. At runtime, every object has a **type** that determines its essential nature and available capabilities. However, unlike Python, Java language is **explicitly typed** in that every name used to refer to data in the program must be **declared** along with this type.
- The compiler enforces that any assignment to a variable has a **right hand side expression** whose type is compatible with the type of the **left hand side variable**. Once a Java program has passed this hurdle of compile time type checking, there cannot exist a **first time** (and therefore **no time**) that a variable is assigned some value whose type does not match the type of that variable, regardless of which actual route the execution takes through the program.
- This requirement of explicit typing is in full effect throughout the source code, even if the intended type would be obvious from the surrounding context. Compared to Python, the syntax of Java seems more bureaucratic when programmers are constantly forced to say things out loud, just to prove to the compiler that they are also aware of those things, yes sir, I can boogie.
- Unlike a Python program that can basically do anything it wants at any time and whose future behaviour therefore allows few guarantees to be done at compilation, **a Java program that compiles cleanly cannot crash at runtime because of a type error**.
- These programs can still crash at runtime due to a host of other reasons, though. (Besides, every self-respecting Pythonista will surely know to write enough **unit tests** to guarantee that their Python program will not be crashing in front of customers because of some silly untested call.)
- Java syntax is more rigid on use of **parentheses** even in situations with no ambiguity.
- Furthermore, every statement must end with a **semicolon**.

1.3. Control structures

- **If-else decisions, while-loops, for-loops**; all these familiar control structures of imperative programming can also be found in Java. However, their details are slightly different both in syntax and use so that, for example, Java uses `else if` instead of `elif` in Python.
- In Java conditions, only proper **boolean** truth values are acceptable as conditions. Other types cannot be used in a “truthy” and “falsy” manner for convenience.
- Instead of clear and intuitive keywords `and`, `or` and `not`, Java uses operators `&&`, `||` and `!` coming all the way from the seventies **C programming language**. Bah.
- The work of the one-liner decision `e1 if cond else e2` of Python is done by Java **ternary selection** with the terse syntax `cond ? e1 : e2`.
- Unlike if-else that is a **statement** but not an **expression**, ternary selection is a proper expression that can be **evaluated** to produce a result, and can therefore be embedded inside some larger expression to make a small local decision that we don’t want to waste a couple of lines of separate if-statement and its associated rigmarole of curly braces.
- It is never wrong to use parentheses to clarify your intention, even if these parentheses are technically redundant. (Modern IDE’s will easily clean up any such redundancies anyway.)
- Java is not Python, so feel free to use parentheses liberally whenever you have a ternary selection inside some larger expression. The potential of a silent mishap is just too great.
- Unlike in Python, the **scope** of local variables in Java spans precisely **the innermost block of curly braces** that surround the declaration of that local variable.

- Once the execution exits the code block that comprises the scope of a local variable, that variable is automatically removed from top of the **stack**, with no effort required from the programmer. In Python, each name comes to existence at the first assignment and continues to exist inside the namespace until explicitly deleted regardless of the execution path taken afterwards.
- Original Java uses a **C-style for-loop** that iterates through an integer sequence, usually a simple **arithmetic progression**. Its head consists of three parts that correspond to the three things that define an integer sequence.
- **Initialization of the loop variable** is executed once when the execution reaches the beginning of the loop. The syntax allows a new loop variable to be declared in that part so that its scope spans the body of that loop.
- Same as in the while-loop, the **condition** determines whether the loop should go on for at least one more round, and is therefore the negation of the goal that the loop is supposed to achieve. For example, if some for-loop is supposed to count up to 100, inclusive, with the loop counter variable `int c`, the condition to keep going would be `c < 101`.
- The **update** is executed each time after executing the body to update the value of the loop counter. In vast majority of for-loops, the update happens with one of the operators `+=`, `-=`, `++` or `--`, but this update could be an arbitrarily complex expression or a heavy method call.
- The **increment** and **decrement** operators `++` and `--` can be applied either from left or right. These correspond to **pre- and post increment and decrement** of that variable.
- The syntax of the for-loop separates the logic of iterating through the values in the sequence from the logic of processing each such value in the body of the loop.
- The slightly impure **break** and **continue** statements can be used to exit a loop altogether, or jump directly to the next round of the loop. Statement **labels** can be used to make these operations work through more layers of nesting, not just the innermost one.
- Same as in Python, both for- and while-loops will correctly do nothing if the condition is `false` to begin with. **When there is nothing to do, the correct thing is to do nothing.**
- Java offers a third type of loop that Python does not have, the **do-while**. Such loops behave otherwise like while-loops, except that they start looking at the condition only after first executing its body once. The do-while loop does not have the option of doing nothing.
- In practice, do-while is by far the least common of the three loop structures of Java, accounting for less than one percent of all loops in actual programs. The do-while loop is the most natural choice only in situations where both (a) it follows from the rules of the current situation that the loop must execute its body at least once to reach the goal, so doing nothing is not an option, and (b) evaluating the condition of the loop does not make sense until the body has been executed once, such as in the classic **input validation** loop.

1.4. Primitive types

- The Java language makes a fundamental distinction between its eight **primitive types** (`byte`, `short`, `int`, `long`, `float`, `double`, `char`, `boolean`) that are defined and set in stone in the language itself, versus all **classes**, any number of which can be freely written in the future.
- These eight are the only primitive types that exist in Java. No new primitive types can be defined by the user. They are, in a sense, the “elementary particles” that all higher level data consists of, analogous to the way that molecules are built up from atoms from a small fixed set of elements. (Or how arbitrary complex and long proteins and other biological structures and superstructures

are built from a small number of different amino acids. Or how words, sentences, paragraphs and arbitrary complex novels up to *War and Peace* can be built up from a fixed alphabet.)

- In fact, analogous to how the physical phenomena that were once erroneously thought as indivisible “atoms” were later discovered to be further divisible to even smaller parts, our primitive types also are not the actual elementary particles of computing. They are all composed of **bytes**, which are in turn composed of **bits** in memory, each one being either 0 or 1. One can't really get any simpler than that.
- These bits and their operations are what all computing is made of, since in principle all computation is merely **syntactic sugar** for **propositional logic**.
- Even when the element type is some primitive type, the corresponding **array** types of any dimensionality are **classes**, not primitives.
- The primitive types correspond to the data types that the processor machine code can directly and natively operate on, so that the operations on these types translate directly to the machine code instructions for efficiency.
- To guarantee portability between different architectures so that every Java program produces the exact same answers for any sequence of integer arithmetic operations, Java **guarantees the sizes, bitwise encoding and potential overflow behaviour for all four primitive integer types**.
- The same guarantee is **not** in effect for **floating point arithmetic**. For the extremely rare cases that require portability of results, the use of **IEEE 754 floating point standard** can be enforced with the keyword `strictfp` applied to one variable, one method or the entire class.
- Java uses double quotes for string literals. Single quotes stand for **character literals**.
- Like every serious system created over the past couple of decades that has to deal with text, Java is fully standardized to use **Unicode character encoding** in the `char` and `String` types to guarantee that all **textual data** is fully portable between systems and never unintentionally corrupted in these computations.
- The Java standard library offers a **wrapper class** for each primitive type, for example `Integer` for `int`, to allow a primitive value to be wrapped inside an object. **Autoboxing** can convert automatically between the primitive type and the corresponding wrapper as needed, with a few caveats explained later, mostly related to the notion of object equality versus value equality.
- These wrapper classes also offer a bunch of useful `static` methods for common operations for that particular type, especially the methods to determine the Unicode properties of arbitrary characters implemented as `static` methods in `Character`.

1.5. Arrays

- Java does not offer **lists, tuples, sets and dictionaries** in the core language. However, these essential aggregate types are available as classes in the **Java Collection Framework** in the package `java.util`. Unfortunately, the sore lack of **operator overloading** in Java makes the everyday use of these aggregate types more painful than the clean and uniform Python syntax.
- The fundamental aggregate type in Java is an **array**, a **fixed-length sequence of homogeneous elements**. The closest Python equivalent to Java arrays would be the `ndarray` type in the practically standard **numpy** numerical computation extension.
- For (almost) any legal type `Foo` in Java, the corresponding array type exists and is syntactically denoted by `Foo[]`, which we should read out loud as “array of `Foo`”.

- Even if the element type in an array is some primitive type such as `int` or `boolean`, the array type is still a class. New array objects must therefore be created with the operator `new`, same as when creating new objects from other classes.
- Same as in Python, square brackets are used to extract an individual element from the array based on its position number. However, Java does not allow negative indices that would be counting from the end, nor **slicing** of an entire subarray into a separate array object. (Slicing can be done with the utility methods `Arrays.copyOfRange()`.)
- The length of the array object is not part of the type (thank the Lord for small favours), so a variable whose type is `int[]` can refer to any such array object regardless of whether it contains ten or ten million integers.
- The length of the array object must be decided and set in stone at the object creation, and cannot change later for that array object. If array resizing is needed, you need to create a new larger array, copy your existing elements there, and continue to use the new array as if it had been the original array all along.
- The length of the array object `arr` is available at runtime as `arr.length`. Note that `length` is a `public final` data field, whereas in the `String` type, `length()` is a method.
- If both the length and the elements of the array object are known, an array object can be **initialized** by listing these desired elements between curly braces. This handy syntax can only be used in initialization but not in a later **assignment**, despite the fact that both these operations are syntactically denoted by the same equal sign character.

1.6. Iterating through one or more dimensions

- Since Java 5, the handy **syntactic sugar** of **for-each loops** can iterate through the elements of the array in order. When applied to an empty array, the loop correctly does nothing.
- For technical and historical reasons, Java array objects don't have any of the useful basic methods that people coming to Java from more reasonably designed languages would expect them to have. Even though the blatantly obvious operations such as **equality comparison** and **conversion to string** are included, their implementations are useless in practice.
- To rectify this flat out embarrassment, the utility class [java.util.Arrays](#) contains these operations, along with many other useful operations such as **sorting and searching**, as `static` methods that receive the array they operate on as the first parameter. Also, the methods for array **copying, slicing and resizing** are also implemented in this utility class.
- For example, to acquire the `String` representation of the array object `arr`, use the expression `Arrays.toString(arr)` instead of the plain old `arr.toString()`. (This tip should make it easier to debug your methods that operate on arrays.)
- The Java type system **only really recognizes one-dimensional arrays**. However, an array whose elements are themselves arrays will in practice behave as if it were a genuine **two-dimensional array**. (If Python has taught us anything, “*If it walks like a duck and quacks like a duck...*”)
- Inside an `int[][]` instance used to represent a two-dimensional grid of integers, each element is an `int[]` that contains the elements of a single one-dimensional row. No law or nature or man dictates that these rows must have exactly the same `length`. Unlike in **numpy arrays** of Python, two-dimensional Java arrays can be **ragged** with different row lengths, even zero.
- A **three-dimensional array** can similarly be represented as an array whose elements are two-dimensional arrays, and so on up the line towards infinity. However, as turns out to be the case with

many other such infinite staircases of computability, only the first couple of steps seem to be ever used in actual programs. (The author's first visceral reaction to seeing some Java program define a ten-dimensional array would surely be to ask which sillyhead has this time misunderstood his need to have an array with ten *elements*...)

Lecture 2: Abstraction and Encapsulation

2.1. Classes and objects

- Back in the old times before the ascent of object-oriented programming, **imperative programming languages** such as **Fortran, C and Pascal** forced the programmers to write **functions** that operate on passive data laid out in the memory. For example, `slice(bread)` ;
- In a more modern, and in many ways far superior **object oriented programming** approach, an **object** is an entity that combines together both the data and functionality that operates on that data. The object internally contains its own functionality, so now it is `bread.slice()` ;
- Or, depending on the way concepts are divided into classes, this might alternatively be `slicer.slice(bread)` ;
- (In fact, to better correspond to the linguistic concepts that these terms refer to when talking about natural languages, is “object-oriented programming” perhaps better called “subject-oriented programming”?)
- When writing Java code, you always write **classes**, which are **blueprints** for objects. From a class that has been successfully **compiled** into Java **bytecode**, any number of structurally identical **objects** can then be **constructed** during the program execution.
- Objects are said to be **instances** of the class that they were constructed from.
- Each object contains some **capabilities (methods)** that the outside world can **invoke (call)** to ask the object to perform that particular action. The functionality has been baked into the “live” data instead of being a separate, qualitatively different thing that exists outside it.
- Objects are intended to represent and model the entities of the **problem domain** that the program was written to perform useful computations on. For example, a bank account object that represents an actual bank account (that we understand to be another “object” inside a different computer) might recognize the methods “deposit” and “withdraw”.
- However, it would not recognize the methods “fly”, “get length” or “honk”, since in the problem domain of banking, those operations would make no sense.
- An object also contains internal **data fields** to remember things that it needs during its lifetime. For example, a toy bank account object would have to remember its **balance** and **owner**. These data fields would then greatly affect the behaviour of the **withdraw** and other methods.
- Local variables describe the reality at the present moment, whereas data fields describe how the reality was in the past. Values of data fields inside the object are called the **state** of that object.
- We will not try to store and remember all of the past (if you think about it a bit, you will realize that the entirety of the past would not all fit inside our present moment anyway), but encode into data fields only those aspects that can affect our future decisions and actions. (This tradeoff determines which character in the movie *Memento* your program is closer to; the one played by Guy Pearce, or the one played by Joe Pantoliani.)
- **Each object should contain exactly as much state as is necessary and sufficient for its future decisions: no more and no less.**
- A new object is created from the class using the Java operator `new`, passing the constructor arguments to the operator `new`. The object continues to exist in memory as long as it remains reachable from the live variables of your program, after which it gets automatically **garbage collected** after some unspecified time by the JVM background garbage collector.

- An object, **once created, cannot change its type or its location in the heap memory**. As long as that object exists, it will remain an instance of the very class that it was originally constructed from. The values of data fields can change during the object lifetime, but the type and the corresponding structure never will.
- Each class should model one **concept** of the **problem domain**. As a rough software engineering rule of thumb, **nouns of the problem specification become classes, and verbs become methods**. This still leaves you a design choice of whether **transitive** verbs belong to their linguistic subjects or objects. (When a boy throws a ball, is “throw” written as a method in the class `Boy` or in class `Ball`?)
- Some objects constructed from the classes of the **Swing framework** can have a visual presence on the user screen independent of the **IDE** that is being used, and can listen to and react to the user actions done with a keyboard or mouse.

2.2. *Instance fields to represent and store properties*

- Variables declared inside the class are called **data fields** (or **instance variables**) so that each object constructed from the class contains a separate copy of that field.
- The positioning of that variable declaration inside the class but outside any methods by itself makes that variable to be a data field. No additional keywords whatsoever are needed to tell the compiler that some variable is a data field.
- All objects reside in a memory area called **heap**, whereas **local variables** are stored in **stack frames** in the **local variable stack**. Values of data fields of an object persist in the heap as long as that object itself exists in the heap, regardless of the method calls and returns performed by these methods.
- Unlike local variables, data fields can have an **access modifier** `public` or `private` to control their visibility to the outside world. As a general rule, all data fields should always be `private`.
- (Leaving the access modifier out entirely denotes **package access** so that the member is `public` inside the same package, but `private` from everywhere else. This allows classes in the same package and under control of the same entity to have closer access to each other's internals than what is given to the outside world.)
- If some outside access to a private member is desired, you should instead provide an **accessor method** (“getter”) and/or a **mutator method** (“setter”). This allows you to enforce logical constraints on the intended legal values of these methods, or make some data attributes virtual in that they are computed from other data that explicitly exists.
- Good **object oriented design** exposes only the **semantic functionality** that objects implement to satisfy the needs of the outside world, and hides the implementation details of how the objects internally implement their behaviour. Data is **always** an internal implementation detail, and should therefore always be `private`.
- Objects can have arbitrary **properties**, some of which are implemented as fields, while some other properties are **virtual**, meaning that they are **computed on the fly** from other properties at the time that they are requested. In a properly designed class, the outside world should in principle not be able to tell which properties are virtual inside the black box object.
- The Java naming convention for the mutator method to set the value of some property `Foo` is `setFoo`. The accessor method for that property is then conventionally named `getFoo` as a reasonable person might expect.

- If that property is a `boolean` truth value, the synonymous variant `isFoo` is also allowed.
- Sometimes something in the problem domain must always be the same for every object of that class. It would be redundant and inefficient (and also error-prone, should the value of that common data change during execution) to store the same value separately inside every object constructed from that class.
- Making a field `static` causes **one and the same copy to exist in memory** and be simultaneously shared by all objects and methods of that class.
- In retrospect, `shared` would have been a far better keyword for this purpose. The keyword `static` is used for technical reasons (in programming lingo, “static” means compile time and “dynamic” means runtime). Too late to change the Java keywords and syntax now, though.

2.3. *More on instance fields, static fields, and local variables*

- Variables defined inside a method are **local variables** whose lifetime spans the execution of the **scope** in which these variables are defined. Once the execution leaves the scope, that local variable automatically ceases to exist.
- In imperative programming languages, local variables reside inside a **stack frame**. A new stack frame is created on top of **stack** at every method call. This stack frame is removed by the JVM when the execution returns from the method.
- If the same method is called again, its local variables are created from scratch inside a brand new stack frame, so they can’t and won’t remember their values from their “past lives” in the previous calls to that same method.
- **Method parameters** are also local variables stored inside the stack frame along with the proper local variables, except that their initial values are provided by whoever calls the method. The method writer should never try to dictate them; ignoring the argument value given by the caller means that the parameter should have been a local variable in the first place.
- **Local variables do not have any access specifier**, since having one would not make any sense to begin with. Local variables don’t even exist unless the method is being executed, so the whole issue of accessing a local variable when the control is outside the method is moot.
- Variables declared `final` can never again be reassigned after their initialization. This keyword prevents certain types of bugs where you or somebody else modifies something that should not be modified. Knowing that some piece of data will never change also allows for compiler optimizations behind the scenes.
- Java 7+ can recognize variables to be **effectively final** and treat these variables as if the `final` declaration had been in the source code all along.
- Note that the concepts `final` and `static` are orthogonal, so that any field could be both, either one, or neither. All four combinations of taking and leaving these two are legal and meaningful.
- A field that is simultaneously `static` and `final` is called a **named constant**. There are many advantages to using named constants instead of writing **magic numbers** in your code. In addition to improved readability, changing the value of a named constant requires the change in one place before recompilation, instead of having to hunt all the occurrences of that constant in the code and edit them all, probably missing one or two and thus causing weird bugs in your code.
- This serves as a nice illustration of the design principle of **Single Point of Truth** that expects any contingent truth about your program and its behaviour to be defined in exactly one place inside that program.

- Parameters and local variables can also be declared `final`, although this is merely a stylistic guideline that cannot affect the behaviour of other methods. The extra verbosity does prevent some types of silly bugs in that method, though.
- Declaring some item of data to be `static` also constitutes an irrevocable commitment that there can be exactly one piece of that data during the program execution. This might be the easy way to do something quick and dirty, but also closes many doors from future improvements.
- For the same reason, programmers ought to eschew **global variables**. Leave them to the seventies where they belong along with everything else that is beige and hairy.
- More generally, an important purpose of proper software engineering is to prevent the accumulation of **technical debt** that accumulates with **quick and dirty solutions** that may save time at the present moment but will impose exponential costs for the inevitable future redesign to prevent the system made up of “duct tape and bubblegum” from collapsing.
- Modern IDE’s offer **refactoring** tools to automatically and reliably convert programming constructs to other constructs that allow further development possibilities.

2.4. Constructors

- For reasons of both **safety** and **security**, the Java Virtual Machine (JVM) makes the following solid guarantee for every time any new object is allocated somewhere in the object heap memory: Before anything else can happen to that object, **those memory bytes are first filled with zeros** to erase every possible remnant of the objects that previously resided in those bytes.
- This policy also conveniently guarantees that every primitive numerical field starts out its life initialized to zero, and that all `boolean` fields start out as being `false`.
- This helps make Java **portable** between different computer architectures by guaranteeing that every Java program runs the same way every time in every computer running a standard-compliant JVM. This would not be the case at all if the object fields were not initialized but were kept with whatever those particular memory bytes contained at the time of object allocation.
- **Local variables** inside a method are not similarly initialized to zero values. However, the Java language requires that every local variable must be explicitly initialized to some value before it is used, even if that variable were honestly intended to start with a zero value.
- Occasionally you might want some fields to be initialized to some non-zero values. If you know that value at compile time, you can **initialize** that field at its declaration. Otherwise, usually because the value of the field will be provided by the entity who needs the new object at the moment of object creation, it is necessary to write a **constructor** in the class to allow initialization of the fields to values that cannot be determined until the actual runtime.
- A constructor is a special method that you never invoke explicitly. (“Hardly ever.”) Instead, the JVM invokes the constructor automatically for a new object. The constructor contains the statements needed to initialize that object and make it ready, willing and able to handle all the action that follows.
- Java recognizes your method to be a constructor from the rather silly syntactic rule that the **constructor has no return type**, not even `void`, and its **name must be exactly the same** as the class itself, including capitalization.
- (Some more modern languages have a separate keyword `constructor` or similar to make it clear what is what and who is who, instead of pursuing the false economy of trying to minimize the number of keywords in the language.)

- **Default constructor** refers to a constructor that takes no parameters. If you don't write any constructors into your class, **the compiler automatically synthesizes a do-nothing default constructor**, to make it possible to create objects from that class.
- However, as soon as you explicitly define even one constructor, synthesis of the default constructor no longer takes place, and new objects can be created using only the constructors that you have explicitly provided. This is useful in situations where the existence of an object simply does not make sense without some particular initial information that the caller must provide as that constructor argument.
- You should prefer using constructors rather than **initializers** for the fields, unless the field is initialized with a constant literal. Since the field initializers are executed in the order that these fields are defined inside the class source code, initializing a field with a method call can cause this method to be executed before the rest of the fields have been initialized (and are therefore still zeroes due to the earlier guarantee of object memory allocation).
- This would impose a silent **order dependency** on the fields defined inside the class that may cause the class to silently break when somebody rearranges its members in the future.

2.5. *Primitive types versus references*

- A **variable** is defined in computer science as “a named memory location that contains a value”. Those bytes will continue to exist in memory separately of how our code assigns names to them according to the **scope** and **lifetime** of those variables.
- Every variable whose type is one of the eight primitive types **contains the stored value inside the variable itself**, in the very bytes that comprise the variable itself.
- However, every variable whose type is any class (including `String`, which could otherwise *almost* be treated as if it were the “ninth primitive type”) is a **reference** that contains the **memory address** of the object that is stored in the **heap**.
- Java does not have **pointer arithmetic** that would allow the programmer to treat a memory address stored inside a reference as an integer that it really is at the level of bytes, and sneakily store it in some integer variable or perform integer arithmetic on it. Java therefore does not have any of the myriad safety or security problems normally associated with pointers in languages such as C.
- It's not the pointers themselves that cause the famous harms, but pointer arithmetic.
- A reference can also be assigned the special value **null** to indicate that the reference is not currently pointing at any object. Any attempt to access the contents of an object through a `null` reference crashes the program with a `NullPointerException`.
- References serve as names for things so that we can talk about those things. This is no different from how names are used in other walks of life. For example, if you mentally assign the name “Bob” to some unknown guy who walks past you in the street, your forgetting that name has no effect on the existence or identity of the person who you referred to by that name.
- This distinction of objects and the variables that refer to them has many important consequences that would baffle anyone who incorrectly assumes that each variable contains the object itself the same way as the primitive types do, instead of actually just containing the directions for JVM to find that object using its memory address.
- Each primitive type takes a different number of bytes to store and represent its value. However, since every memory location is uniformly numbered regardless of what sort of value the bytes in that location happen to currently semantically represent, every reference in Java takes the same

number of bytes of memory to store and represent. In modern 64-bit architectures, a memory address is stored in eight bytes.

- For primitive types, the **equality comparison** operator `==` compares the actual values stored in them. For references, the operator `==` still compares these values, but those values now happen to be memory addresses, not the object contents.
- **The result of comparison of two references using `==` will never depend on the object contents**, because this operator checks only whether both objects are the exact same object in the same memory address. For this reason, either one of the two references compared in this manner can even be `null` and nothing bad could still happen.
- If you want to compare the contents of the actual objects that `a` and `b` refer to, use the expression `a.equals(b)`. Since equality is **transitive**, `b.equals(a)` should work equally well.
- The memory address gives every object a **unique identity** that separates it from all other objects inside that are currently live the heap. Even if two objects were 100% identical clones of each other in every possible way with respect to both their types and values of their data field, the fact that they reside in two separate memory addresses allows us to tell these two objects apart.
- After some object has become unreachable from your code after you have lost all references to it, and the **garbage collection** has made that part of the heap memory available for future allocations, another object might well be created in the same memory address. However, your program won't be able to tell even in principle that this has happened, since your program no longer has any live references to that memory location that it could compare its new reference to!
- Garbage collection allows programs to run for an arbitrary long time without worrying about running out of memory. In more primitive languages where the memory must be released explicitly back to the system, the Scylla and Charybdis of **memory leaks** and **dangling references** cause all kinds of difficult and unpredictable bugs in programs.
- A Java program can still have an **indirect memory leak** in that some reference that you will never again use keeps the object alive. To prevent this, you should set a reference to equal `null` once you know that you will never again use the object that it points to.
- Since all process memory is released when the process terminates, programs whose runs are short could theoretically run without any garbage collection. **Embedded systems** are usually meant to run uninterrupted for a long time, even for years or decades, and therefore should not leak any memory over time, not even an almost unnoticeably slow drip.
- (On the other hand, there is the tale of a guidance system inside a missile, where the designers simply removed the garbage collection altogether and put in twice as much physical RAM as was needed to cover the theoretical maximum of memory allocations during the maximum flight time of that missile. After the execution of the guidance system program had completed, there was no longer any need for garbage collection, at least not within the program that no longer seeks to conquer the human, nor to be conquered by the human. [This is all totally “warrior Zen”, or Tao, or something else vaguely exotic of that nature as long as it makes me look cultured and smart.](#))

2.6. References as method arguments

- When executing some method call of the form `obj.method(args)`, the JVM follows the reference `obj` to the actual object in the heap. The object type, stored as an invisible field alongside the object, determines which version of the `method` is called with the given `args`.
- Two references are said to **alias** if they contain the memory address to the very same object.

- No-one can be faulted for asking why anyone would create two variables `clarkKent` and `superman` and assign them both to refer to the same object. Alas, such aliasing must inevitably happen every time you call a method whose parameter type is a class, since that method expects some reference given to it as an argument at the time of the call.
- [All parameter passing in Java is always done by value](#). However, passing a reference to a method by value causes the caller and the method to share the same object. If the method goes on to modify the internal state of that object, that modification of the object state will persist even after the method has terminated and the control has returned to the caller.
- Passing a reference to an object by value can be used to simulate passing its referred object to a method by reference. The method can modify the internal state of the object, but cannot modify the original reference to the object.
- This phenomenon is most commonly seen with **methods that operate on arrays** that can contain (even hundreds of) millions of elements. It would be highly inefficient to create a separate copy of that array object for every method call.
- Instead, the method is given only the memory address where the entire array object resides. It therefore takes the exact same amount of execution time to pass a million element array to a method as it would take to pass a tiny ten element array.
- This is most commonly seen in **recursions** that operate on arrays, since it would be inefficient to extract the particular subarray under interest into a separate array object to be passed down the line in recursion. Rather, the entire array is passed down by reference at every level of recursion, and additional **index parameters** tell the method which particular subarray it receives as its turf to restrict its operations in.
- (These parameters are typically called `start` and `end`, or sometimes just `n` to denote using the first `n` elements. The `end` index is conventionally **exclusive**, the same way as in Python.)
- As seen earlier, references can also be given the special value `null`. This value does not exist for primitive types, because all their possible bit value combinations are already in use to mean actual possible values of that type.
- Inside the JVM, `null` is guaranteed to be encoded as bytes that are all zeros **to indicate that the reference is not really pointing to any object at the time**. All reference fields are therefore guaranteed to become `null` at the object creation.
- The existence of `null` references is the infamous “[billion dollar mistake](#)” of software engineering. Languages such as Swift created in the 2010’s fight this scourge with source-level guarantees that some reference cannot be `null`, so that the method receiving such a reference does not have to go through that *kabuki* of checking for `null` at runtime.
- Programming in Java is guaranteed to respect **memory safety**. Try as you might, the lack of unsafe **pointer arithmetic** in Java makes it impossible to produce a reference that points to some memory location that does not currently contain a live object in the object heap! Throughout its entire lifetime, every reference is either `null`, or is assigned a value from some other expression that points to a live object in the heap.
- This maintains the global **invariant** of every reference always having a legal value in this sense, since there cannot exist the first time that an illegal reference gets initialized!

2.7. Immutable versus final

- Same as with variables whose type is one of the eight primitives, a reference can be declared to be `final` so that it cannot be later reassigned to point to some other object.
- Because a reference is a separate entity from the object that it points to, the keyword `final` affects only that reference, and does not reach into the actual object to magically encase it in carbonite. That same object can be simultaneously pointed to from any number of separate references anyway, some of which are `final` while some others are not.
- Declaring a reference to be `final` does not prevent any method calls through that reference from modifying the state of the object. For this reason, `final` should not be confused with **immutable**, a related concept that means that the class has been **intentionally designed** so that **it has no public mutator methods**, but accessors only.
- (In practice, a couple of other restrictions are needed to guarantee that no outsider gets to put their fingers into this pie to mess it up. Interested students can see “Effective Java” for details.)
- An object constructed from an immutable type cannot change its internal state later during execution in a way that could later be detected from outside using the `public` methods of that object.
- Objects that are immutable from the point of view of the outside world can still optimize their execution time by internally **caching** the results of some computations that they might again need later, that way modifying their internal state. However, this does not affect the externally observable results that the methods of that **black box** object return to their gentleman callers, except that those same results will be provided to them faster.
- Immutability provides a whole bunch of surprising advantages in programming, which is why we **favour immutability** and consciously aim for it whenever possible and convenient. However, immutability [has also one giant downside](#) in situations where the code needs to iterate through a large set of different values and would therefore require the creation of a separate object for each such value, which is why we don't automatically make all data immutable.
- For example, `String` is intentionally defined to be an immutable type in Java, for the same reasons as in Python. To allow us to sidestep the previous downsides, `StringBuilder` represents a **mutable** sequence of Unicode characters.
- `StringBuilder` is optimized for the common operation of **appending** more characters into the end of the string by keeping extra **slack space** in the underlying character array.
- In computing, we can often **trade memory for time and vice versa**. Since these days we usually have more memory available than we know what to do with it, taking this deal for that direction usually makes your program run faster.
- (For algorithm enthusiasts: Building a long answer string piecemeal using immutable strings would require an $O(n)$ copy of the entire string for every single individual character **append**, making the result building to be at least $O(n^2)$. Since the `StringBuilder` class offers an $O(1)$ **amortized time** operation **append**, the result will be built up in $O(n)$ total time.)

2.8. The implicit reference `this`

- Sometimes it might be greatly beneficial for some instance method to be able to find out which particular object it is currently being executed on. To allow this, every instance method automatically has a special local variable `this` that points to that very object.

- Most of the time using `this` is redundant, because when some method call `foo()` is written without the explicit object reference prefix, the compiler silently rewrites this call as if you had written it as `this.foo()`.
- Using `this` is unavoidable in methods that call other methods that need to know the identity of the original object, passed as a parameter to the said method. (This situation occurs in object-oriented design most often with the so-called **callback** methods.)
- If you give the constructor parameter the exact same name as the field that that parameter will be assigned to, the parameter **shadows** the field with the same name in the statements of the method body. The otherwise redundant reference `this` can then be used to force the compiler to treat that name as a field inside an expression.
- Back in the day, the reference `this` was syntactic sugar. Any **method** call `obj.foo(args)` could be converted by the compiler into an ordinary **function** call `foo(obj, args)`, with the implicit parameter `this` receiving its value from the first parameter that refers to the implicit object for which the method call is being made.
- (Python still uses `self` for this **implicit parameter** of each instance method.)
- Methods declared `static` are called without an object, and therefore do not have the implicit parameter `this`. Instead of prefixing such method call with an object reference, these calls are prefixed with the name of the class that contains the `static` method, such as `Math.sin(2)` or `Character.isWhitespace(ch)`, to resolve the potential ambiguity if more than one class defines a `static` method with an identical signature.
- Same as how the reference `this` can be left out when calling an instance method in the same class, the class name prefix can be left out when calling a `static` method in the same class.
- A method ought to be written to be `static` if it does not need anything from the underlying object, but has everything that it ever needs to compute its result in its parameters and the `static` fields of the class. A non-static method can always be later turned `static` if need be, with the refactoring tools of modern IDE's again eliminating the possibility of human error.
- A class that contains nothing but `static` methods for the common operations for some particular data type and is not intended or designed for instance creation is called a **utility class** for that data type. The Java naming convention is to name such utility classes after that data type but **pluralized**, such as `Arrays` or `Collections`.

2.9. Method overloading

- Unlike in our natural language, the term **overloading** has no negative connotations in object oriented programming. This technical term simply means that the class defines several methods with the exact same name but different parameter lists.
- Using the exact same name for all these methods immediately conveys to the reader the message that these methods are **semantically identical**. That is, they achieve the same purpose even though they happen to operate on different types of arguments.
- The overloaded methods could just as well have been given different names, but in general it is bad style (and even a **code smell**) to have your method name refer to the data types that it receives as parameters.
- **The parameter lists of overloaded methods must be different** ignoring the parameter names so that the compiler can always determine without ambiguity which method is being called. The

overloaded versions may also use different access modifiers and return types, but these have no effect on method overloading resolution done solely on parameter types.

- Having two overloaded methods taking parameters of different primitive types is perfectly legal. Note that adding a second method `foo(int x)` to the class that already has the method `foo(double x)` causes the call `foo(42)` to no longer be bound to the second method.
- It is legal for a class to have two overloaded methods that differ only by the order of their parameters. However, this possibility still does not mean that you *should* do this, as everyone will then have to always remember which one of your two such methods was which, so this mental effort will no longer be available for more important things. You are encumbering the proverbial camel with an extra straw for zero benefit. Eventually such little straws, each one by itself as light as to be almost meaningless, will add up to break that camel's back.
- Constructors can be overloaded the exact same way as the ordinary methods. The choice of which constructor gets executed at object creation is based on the types of the arguments given to `new`.
- Sometimes overloaded constructors contain duplicated code, another classic code smell. Java does allow the odd-looking call to a constructor from another constructor to eliminate such duplication, done by using `this` as a method name. (As one of the [Epigrams](#) of Alan Perlis points out, in English every word can be verbed.) However, such calls are further restricted so that such a call must always be the first statement in the constructor.
- Overloading should not be confused with **method overriding**, a very different but an extremely important concept for the rest of this course for inheritance and polymorphism. We could easily program all of our complex systems without method overloading and really not even miss anything essential. However, without method overriding everything in Java would instantly collapse, so that we might as well put on our brown seventies pants and go back to coding in either Basic or C.

Lecture 3: Inheritance and Polymorphism

3.1. Types in programming languages

- During the execution of a program, some **data objects** exist in the memory for the Java code to operate on. Even though all these objects are nothing but raw bytes, we prefer to impose a convenient fiction of **higher level semantics** upon them to make them easier to reason about.
- We give an intended meaning for those bytes from the outside by operating on those bytes only as if they really had that meaning, which causes them to effectively have that meaning. As every sociologist knows well, [things that are treated as real have real consequences](#).
- Any higher level entities defined in the language are fictional in the exact same sense that, say, hobbits and vampires are fictional. Despite this, these fictional entities such as “functions” or “loops” can still be used as convenient shorthands in expressing ideas whose consequences might as well be real as far as our future actions and decisions are concerned.
- In the **type system**, a **type** defines all those properties that will steadfastly remain fixed during the entire lifetime of each object that was constructed from that type.
- The purpose of types is to prevent **type errors** in programming by setting up voluntary but binding constraints of what can be done with given data, akin to an extremely well-behaved good dog taking itself for a walk while holding the other end of its own leash in its mouth. In **explicitly typed** programming languages such as Java, such errors are revealed at compilation.
- **It is always better to find errors as soon as possible** so that you know to turn back whenever you are hurtling towards inevitable failure, instead of being like that guy in that joke who fell off the top of a skyscraper and was heard saying at every floor “So far, so good...”
- Java is a rather heavy and verbose language for programmers, all this legalese designed to reveal errors at compile time. Contrast this to Python, Ruby and similar languages that instead optimize for ease and brevity, leaving most of the error detection above basic syntax rules for **unit testing**.
- One comment to the article “If programming languages were weapons...” that made rounds on the Internet a few years ago gave Java a much more accurate description than the original article, calling it “a belt fed 240G automatic weapon that tries hard to avoid gun accidents: each round has to be individually authorized by ticking off a detailed questionnaire. By the time you get to actually shoot, someone may have already beaten you to death with a stick.”
- In Java and all similar object-oriented programming languages, new types are defined as **classes**. Any number of separate **instances** of this type can then be created. These different instances of the same type are **structurally identical**, but store different values in their data fields and reside in different memory addresses, which allows us to tell them apart even if the contents of their data field happen to be indistinguishable.
- In **object oriented design**, each class would ideally correspond to one **concept** in the **problem domain**. If the problem domain consists of n different concepts, the program would consist of n classes, with the **verbs** of the concepts becoming the methods of these classes. Difficulties in thinking up short descriptive names for classes and methods tend to strongly suggest bad design choices in choosing these classes and methods.
- The problem level concepts tend to form hierarchies so that both checking accounts and savings accounts are subtypes of bank accounts, same way that both hawks and sparrows are subtypes of birds. Birds in turn are subtypes of animals, and animals are subtypes of objects.

- Types should be used to model static variation in behaviour, never the dynamic variation of the current internal state. **Properties that can change over the lifetime of the object should not be distinguished using the type system.**
- For example, the types `CheckingAccount` and `SavingsAccount` can reasonably co-exist, but having `AccountWith1000Dollars` and `AccountWith500Dollars` existing as separate types would be Bizarro World coding.
- Since the balance of the particular bank account can change over time, the balance should instead be represented as a data field inside the class. Note that this is a completely separate question from asking whether that bank account type should be designed to be mutable or immutable!
- If classes can express types but no inheritance mechanism exists to organize these types (as it were in languages that predate object oriented programming), it becomes impossible to express the notion that the problem domain concepts of `Hawk` and `Sparrow` are somehow “closer” to each other than, say, the concepts of `Hawk` and `Planet` are to each other. Every type would be an island, as far as the type system was concerned.

3.2. *The need for inheritance to avoid repeated code*

- Unlike in the physical world, within the pure Platonic perfection of mathematics there should **never be any need to repeat anything**.
- For example, billions of people on Earth can simultaneously think of and use the same number 7, instead of every person needing to possess his or her own personal copy of that same mathematical object. Since computer programs are mathematical objects, this principle should apply to them equally naturally.
- The **DRY Principle (Don't Repeat Yourself)** instructs us to **never say anything twice** in the same program. This is related to the earlier **Single Point Of Truth Principle** that requires every **contingent truth** about the behaviour of the program to be stated **exactly in one place**.
- An important practical application of the Single Point of Truth Principle is to require that you are not allowed to write any **magic numbers** other than 0 and 1 in your code, and all other constant literals must be given descriptive names as `static final` **named constants**. This makes it easier both to understand what is going on, and to modify these named values later.
- Inheritance allows subclasses to inherit a method implementation defined once in the superclass, instead of the programmer having to copy-paste the same method definition to all subclasses. However, contrary to the common misconception about object oriented programming, even though this is certainly a nice thing to have, this is not even close to being the main benefit bestowed on us by the inheritance mechanism.
- Without inheritance and static type checking, it would not be possible to write a type safe **polymorphic** method `void pluck(Bird b)` that accepts as its argument any past, present or future subtype of bird, but sternly rejects any planet, bank account or car with a compile time error. In absence of inheritance, you would end up writing two essentially identical methods `void pluck(Hawk h)` and `void pluck(Sparrow s)` only to keep the language's type checking formalism happy, eventually exploding in tragicomical absurdities not seen since the glory days of Soviet economic planning.
- Even worse, every time some new subtype of bird was introduced, you would have to write yet another (essentially identical except for the parameter type) version of this same method, always repeating yourself *ad nauseam*. We therefore say only “¡No pasarán!” to any such approach.

- Of course, you could completely toss out all **type safety** and write the method `pluck` in a Pythonic manner to accept any object whatsoever as argument, since you intend to only ever call it with birds. But **intentions are not magic**; what would happen if somebody gave this method, say, a `String` or a `BankAccount` object anyway, and the method then tried to call the method `fly` for that object? The only possible outcome would be a **runtime crash**, which, as per Murphy's law, would usually occur at the worst possible moment during program deployment.
- **Compile-time type checking** exists in all such programming languages **to prevent runtime crashes** due to the program trying to do something logically impossible in the problem domain, such as asking a checking account to `fly`, or trying to `deposit` a hundred dollars inside an owl.
- The compile-time type checking of Java is strong enough to make it **impossible** to create a program that crashes at runtime because it calls some method that doesn't exist in that object. (A Java program can still easily crash for various other reasons that are algorithmically impossible to detect in general during compilation.)
- **Dynamic languages such as Python or Ruby**, seeing that they have no explicit type system at the source code level, cannot make a similar guarantee. (Some Python **static analysis** tools can use the Python **type hints** for this purpose.) This is the price we pay for the extra flexibility of **duck typing** in line with the dynamic philosophy of those languages.
- Decent use of **unit testing** alleviates this problem so that it is not really any problem in practice. As any self-respecting Pythonista will tell you, if your Python program actually crashes at runtime due to typing error, that only shows the poor quality of your unit testing, instead of revealing anything about the quality of the language.
- (Also, contrary to an even more commonly repeated misconception, **Python is compiled**. Yes, Virginia, it really is. This compilation, just like it has taken place in every serious interpreter in every “interpreted” language within your lifetime, merely happens silently as the textual source code is read in, instead of the programmer having to explicitly initiate this compilation process. This compiled bytecode is cached in `*.pyc` files for future use.)

3.3. *Inheritance between classes*

- A class can be defined to **extend** some **superclass**. This causes the subclass to **inherit** all the **fields** and **methods** of the superclass, almost (although not really; we will come back to this issue) as if they had been copy-pasted into the class body.
- Proper inheritance should always model an **is-a relationship** between the problem domain concepts that these two classes represent. Therefore it is sensible to say `class Car extends Vehicle`, but it would be utter nonsense to say `class Car extends Engine`.
- Every car **has** an engine, but **the car itself is not an engine**, since its **public interface of methods and resulting behaviour are completely different**.
- The public interface is a tool of **abstraction** that lists the **essential features** of the system that cannot be changed without changing the essence of the system, as opposed to its **accidental implementation details** that can be freely changed without affecting the outside users of that class as long as the public interface does not change.
- The **has-a relationship** is better modelled with ordinary **composition** with object fields. This can be either one-to-one (such as between a car and its engine) or one-to-many (such as between a car and its wheels or its passengers).
- Since a Java object cannot change its type after it has been constructed, **inheritance should only be used to express variation within the kind that remains constant during the lifetime of the**

object, such as the species of the animal. **Dynamic** aspects such as age, weight, and the physical location of each individual animal, are better modelled as object fields.

- Since an object can never change its type after construction, the class `Car` being a separate subtype of `Vehicle` also requires that `Car` has no **mutator** methods that could turn some car object into an `Airplane` or a `Submarine` in the spirit of 1970's James Bond movies, but every car that is built will solidly remain a car until the day that it ends up in the (heap) compactor.
- In problem domains where higher level things built up from smaller parts can be reconfigured to create new things in the style and spirit of Legos, high level classes represent these configurations as **compositions** of low level objects, and the objects constructed from those classes are supposed to exist only as long as that configuration exists. (In a sense, ordinary classes also do this, with the raw memory bytes being these Lego building blocks.)
- The individual Lego blocks are immutable, but the exact same individual pieces can be put together in many different ways to create either a `Spaceship` or a `Castle`. The individual pieces are exactly the same in both structures, but their different arrangement makes their emergent behaviours different.
- Furthermore, somebody (perhaps you!) might use these very same Lego pieces to build up some entirely new kind of useful structure that we could not even imagine during the original design and our language therefore does not have a word for.
- In the Java type system, **the subclass objects are at the same time also superclass objects**. For example, every `Hawk` object is at the same time a `Bird` object. `Hawk` and `Bird` are still separate classes in the type hierarchy, and the subclass object simultaneously has two different types, each type used as the public face of this Janus as needed by the rest of the program.
- Variables of type `Bird` can also refer to objects of type `Hawk`, or any other subtype of `Bird`. This is necessary for type-safe polymorphic methods that can take arbitrary birds as parameters that they operate on.
- The subclass can also define additional members that do not exist in the superclass. The subclass objects will then have these additional members, whereas the superclass objects will not.
- Through a reference of type `Bird`, you may only call methods that exist in the class `Bird`, even if you knew for a fact that the actual subclass object that the reference points to has that method.
- Assuming that the subclass `Owl` defines a brand new method called `hoot`, and you have declared `Bird hedwig = new Owl();` the call `hedwig.hoot();` will then not compile since the variable `hedwig` is a `Bird`, and there is no method `hoot` in the class `Bird`.
- **The compiler can never make any decisions based on any objects for the simple reason that objects don't exist at the time of compilation!**

3.4. *Overriding inherited methods*

- The subclass definition can also **override** an inherited method simply by providing a new version of that method, with the **exact same name and parameter types**.
- **Only methods can ever be overridden, fields can not be.** This is why you should never make fields `public`, but define **public accessor methods** to read and write them. This allows properties to be **virtual** so that their values are computed on the fly based on other properties.
- Since it is easy to accidentally **overload** a method instead of **overriding** it, it is a good idea to always tack the `@Override` **annotation** in front of each overridden method. Such **syntactic salt**

helps the compiler catch these easy mistakes at compile time. (Some modern languages even have `override` as a required keyword in this situation.)

- In overriding any method in a subclass, its parameter list must remain **exactly the same** as in the superclass version. Resist the common temptation to replace a parameter type by its subtype.
- The **return type can be made more specific** in the subclass version. For example, if the superclass method returns a `Bird`, the overridden method can return a `Hawk`. Since every `Hawk` is also a `Bird`, the caller of the subclass method gets what it was promised.
- The **body** of the subclass version of the inherited method does not need to have anything in common with the body of the original superclass version. (The compiler would not be able to algorithmically enforce that any semantic properties were maintained anyway.) However, the famous **Liskov Substitution Principle** will later precisely define what constitutes spiritually proper overriding.
- If a method is declared `final`, the subclasses are not allowed to override it. They have to take that method implementation exactly as it is defined in the superclass.
- If the entire class is declared `final`, no further subclasses can be extended from that class. (This helps enforce immutability in classes such as `String`.)
- When designing a class, you should document **whether the class itself uses some of its methods that subclasses can override**, that is, whether those methods are **template methods**. This prevents the subclass writers from unintentionally breaking the behaviour of superclass methods.
- For example, all other methods of `Random` are guaranteed to internally use the template method `next` to generate the **pseudorandom** bits that they need. When extending this class, it is therefore sufficient to override only the `next` method to produce these pseudorandom bits in a new way, and all the other methods will automatically follow that leader without further ado.

3.5. Polymorphism and dynamic binding

- Since class members will be preserved in inheritance, **the subclass is guaranteed to have all the public members of the superclass**. For example, if the superclass `Bird` has a method `fly`, every subclass of that class is also guaranteed to have a method `fly`. Even if this method has been **overridden** to do something different in different subclasses, the existence of that method is uniform throughout the entire class hierarchy.
- This allows methods to be **parameter polymorphic**, that is, **methods whose parameter type is some superclass**. This one and the same method can then operate on **all past, present and future subtypes** of its formal parameter type, so that we don't always have to write yet another redundant version of this same method every time some blasted new subtype is introduced.
- For example, the polymorphic method `void pluck(Bird b)`, written, compiled, tested, documented and released only once, can be given any instance of `Hawk`, `Owl`, `Albatross`, or any other bird that did not even exist at the time when the original method `pluck` was first dreamt of in our philosophy!
- Since the method `pluck` solidly depends on a high level abstraction `Bird` that will never change, polymorphic methods are **future compatible** with all subtypes that anybody will ever create in the future, despite this method `pluck` having been written, compiled, tested, documented and released only once!

- If the body of the polymorphic method `pluck` somewhere contains the call `tweety.fly()`, this call **cannot possibly be bound at compile time**, because this method call is supposed to do a different thing when called for different subtypes of birds.
- For this reason, Java method calls are **bound dynamically** at the time when the call is actually executed. The JVM follows the reference `tweety` to the actual object in the heap and looks at its internal **type field** to choose on the fly (heh) the correct version of `fly` to be executed.
- The very same method call will therefore do a different thing depending on the subtype of the object for which that method is called. Dynamic binding allows the existence of polymorphic methods, **the greatest single benefit of having classes and inheritance in the first place**.

3.6. Abstract superclasses and methods

- Sometimes the problem domain concept that the superclass attempts to capture is so **abstract** that it is flat out **impossible** to give a concrete implementation for some particular method.
- For example, we can describe how a hawk flies, and on the other hand, how very differently a hummingbird flies. However, how does a `Bird` fly? The concept of bird is simply too **abstract** for us to give an answer, even though every bird flies.
- In our toy problem domain, penguins and ostriches don't exist, let alone newly hatched chicks or plush toy birds. (Every model must be a simplification of the problem domain, otherwise you might as well reach into the problem domain and use it directly as its own model.)
- However, if the method `fly` were not defined in the superclass `Bird`, this would allow for anybody to create a new subtype of `Bird` that cannot fly. The polymorphic methods that operate on birds then could not assume that the method `fly` exists.
- In such a situation, the method `fly` must be defined as `public` in the superclass `Bird` to guarantee that all **concrete** subclasses of `Bird` will also have this method, but also then be declared **abstract** in `Bird` so that it is given **no body at all, not even the empty body**.
- If the class has one **abstract** method, the entire class itself must also be declared **abstract**. This declaration **prohibits all object creation** using the operator `new`. (If such an object were allowed to actually exist, pray tell, what would happen at runtime when somebody eventually created that object and then called its **abstract** method that doesn't have an implementation?)
- **Abstract types are used as parameter types in a polymorphic method.** This method can then be given objects created from the concrete subclasses of `Bird`.
- **All four possible combinations abstract-abstract, abstract-concrete, concrete-abstract and concrete-concrete are legal and meaningful** when it comes to the relationship between a superclass and its subclasses.

3.7. Visibility of class members

- Every object created from a subclass contains all the fields and methods defined in its superclass, even those that were defined `private` at the superclass level. However, the original access modifier determines whether these fields and methods are visible to the methods of this subclass.
- Members declared `private` can be accessed **inside the very same class only**. The subclass methods **cannot see these inherited fields and methods**, even though they do exist inside the object! (Use the heap byte inspector, if you don't believe me.)

- The only way that subclass methods get to access the `private` members of the superclass is by doing so indirectly by calling the inherited methods that can access these members. Since these inherited methods are defined in the superclass, they still get to see and access the `private` superclass members even when working at the subclass level.
- Sometimes you would like to invoke the superclass version of an overridden method because it does something useful that the subclass version could equally use. Instead of copy-pasting that code into the subclass version of the method, use the keyword `super` as the prefix of the method call to force the superclass version to be executed.
- Members declared `public` can be accessed **from anywhere** that has a reference to the object that contains them. Any inheritance relationships have no effect on this.
- Members declared `protected` can be accessed in the **same class and subclasses**, and in all the classes that are part of the same package.
- Data fields inside a class should *always* be declared `private`. If some field needs to be accessible from the subclass methods, provide a `protected` accessor method for that field, instead of declaring that field itself `protected`.
- Requiring accessor methods to access fields that represent **attributes** also gives you leeway to later make that attribute **virtual**, that is, computed dynamically in the accessor method at the time of the call. For example, the class `Length` that stores lengths as centimetres can also provide an accessor method for inches. (Or even furlongs or printer's points, should somebody turn out to need such units in the future.)

3.8. *The universal superclass Object*

- In the inheritance hierarchy of Java classes, the universal superclass `java.lang.Object` is a superclass of all other classes in that starting from any class and climbing up the inheritance tree as high up as it goes, you will **always** end up in the class `Object`.
- Whenever a class does not explicitly extend another class, it implicitly extends `Object`. There is no mechanism to define a class that has no superclasses at all.
- `Object` itself has no superclass. Even more amazingly, this universal superclass `Object` is not `abstract`, although any actual need to create objects of type `Object` is now long obsolete.
- The existence of `Object` as a type allows **references that can point to any object whatsoever**. Of course, through such reference you still get to only call methods that exist in the universal superclass `Object`, since no other methods can be guaranteed to exist at the time of the call.
- The second advantage of the existence of the universal superclass `Object` is the guaranteed existence of a bunch of handy methods such as `toString`, `hashCode`, `clone`, `equals` and `finalize` in all Java classes and therefore in all Java objects.
- Unfortunately, most of these methods, especially `toString`, have a rather useless default implementation in `Object` that is best overridden in actual subclasses.
- When overloading `toString`, it is also a good idea to provide a public accessor method for every field whose value appears in the generated string, so that the users of this class won't need to **parse** that information from that string. Having to do so would permanently marry the user code to that particular implementation of `toString`.

- The default implementation of `equals` only checks whether the objects are the same object stored in the same memory address, same as the operator `=`.
- This is the most that this method could do knowing nothing more than that both objects are some kind of objects. In general, the higher up you go in the inheritance hierarchy, the less you can assume and say about the actual objects.
- Since there are no guarantees of when the garbage collection will release some object, you should never rely on the `finalize` method to release important resources held by that object. Rather, you should release such resources explicitly the moment that you know that you will no longer be needing these resources.
- Overriding `finalize` to do something nontrivial can and will also cause a massive efficiency hit to guarantee the correctness of garbage collection. Such overriding should only be done if necessary to guarantee that important resources held by the object are certain to be released.

3.9. Overriding the `equals` method

- By overriding the `equals` method in a subclass, you can implement a more specific notion of equality than the trivial **memory address equality**, the strictest possible equality relation.
- As a rule, your `equals` method should implement a proper **equivalence relation** that is **reflexive**, **symmetric** and **transitive**. Every such relation partitions the space of possible objects of that type into disjoint **equivalence classes**.
- However, such **semantic properties** of methods cannot possibly be enforced by the compiler, since the problem of checking whether an arbitrary method has any nontrivial semantic property is known (since the 1930's by the work of **Alan Turing**) to be **algorithmically undecidable** in the general case. (The given semantic property such as termination can often be reasoned mechanistically for some trivial methods, but no algorithm can possibly do this in the general case so that it would always produce the correct answer for every method.)
- There exist infinitely many possible equivalence relations for the same type, these relations forming a mathematical **lattice**. On one end of this multidimensional diamond shape lies the **strictest possible concept of equality** where every object is equal only to itself. On the other end lies the **loosest possible concept of equality** where every object is equal to every other object. Between these two extremes there exists an infinity of other equivalence relations, from which you are expected to choose the one that is **congruent** with the computations that you intend to make.
- Overriding `equals` in some silly way instead of making it a proper equivalence relation will cause many other classes and methods (such as the methods `add` and `contains` in the Java Collection Framework classes) that depend on these properties to silently behave in an unpredictable fashion, making all bets on the program behaviour to be essentially off.
- The parameter type of the `equals` method must be `Object`, reflecting the reality that **any object can be equality-compared with any other object**. Comparing the proverbial apples and oranges is not any kind of compile time or runtime error, but a perfectly legal and meaningful operation whose answer in this situation just happens to be "no".
- (A surprisingly common misconception even among educated people is that only questions that both (a) are non-trivial and (b) whose answer is "yes" are somehow finger-quotes "legitimate" questions to begin with. This confusion might partially stem from the more fundamental difference between the ["ask" cultures](#) versus ["guess" cultures](#).)

- If you accidentally make the parameter type to be the same as the class in which you override this method, you are not overriding but overloading! The other classes will then be calling the original superclass version instead of your intended new version of the method.
- Again, use the `@Override` annotation to reveal such errors.
- Whenever you override `equals`, you should also override the method `hashCode` so that whenever `a.equals(b)`, also `a.hashCode() == b.hashCode()`. Otherwise any collections that use hash codes to organize their stored elements (most importantly, `HashSet<E>`) will also silently produce gobbledygook.
- The hash code for an object should be computed by suitably scrambling and mixing the hash codes of **precisely** those fields whose values can affect the result of `equals`. If some field that does not affect the equality comparison is used in hash code calculation, two objects that are `equals` will end up having a different hash code. This again silently corrupts the behaviour of all data structures and algorithms that make decisions based on these hash codes.
- No matter what, many unequal objects will still inevitably end up with equal hash codes. After all, there exists an infinity of potential objects, but only 2^{32} possible `int` valued hash codes to be distributed among those objects, so those objects will just have to learn to share. Any hash table data structure could therefore in theory, with truly astronomical levels of bad luck, devolve into essentially a **linked list**.
- (A far more likely scenario for this to happen is that your **adversary** gets to tactically choose the objects to have the exact same hash code. The blind justice of nature is never that capricious.)

3.10. *Dynamic type inference*

- Since specific types of equivalence based on the value of some particular attributes make sense only within the same type, overriding `equals` properly requires a mechanism to dynamically detect whether the given object is of a particular type.
- The expression `obj instanceof Bird` checks at runtime whether the object `obj` is a `Bird`, or any subtype of `Bird`. Instead of crashing, this `instanceof` operator will also correctly return `false` whenever `obj == null`, avoiding the need to always check that first.
- Alternatively, you can use **reflection** to ask an object its type with the method `getClass()`, defined `native` in `Object` and thus inherited to all classes for all objects to recognize it.
- To access the subclass members that do not exist in the superclass such as `Object`, you need to **downcast** the reference to create a new subclass reference to point to that same object.
- (Java 15 introduced new syntactic sugar for defining a new local variable with the correct subtype in the same `instanceof` expression.)
- All downcasts are **silently checked at runtime**. This makes it impossible to use downcasts to trick the type system to call a method that does not exist in the object.
- In Java, it is never necessary to use an **upcast**, since any reference of superclass type can point to any object from any subclass type. However, if only to again get to say “Hardly ever” in a British accent, the upcast mechanism still exists in the language, since it is genuinely needed to disambiguate between the calls to two overloaded methods whenever the offered argument would be compatible with both methods' parameter types. (For example, the argument value `null`.)
- (Speaking of `null`, this constant in the Java language is the only value of its own special type that is the opposite of `Object` in that this type is a **subtype** of all other types. This allows any

reference variable of any type to be assigned to have the value `null`. Since `null` is not an object, this does not lead to any paradoxes of calling methods through such a reference.)

- In practice, using `instanceof` (or the result-equivalent **reflection mechanism**) is usually **the absolutely wrong way to do things**, especially when doing so creates the infamous object oriented programming **anti-pattern** known as the **instanceof ladder**.
- Instead of manually dispatching the execution to different branches based on the inspected type of the argument object, you should always let Java's compile time analysis and the runtime dynamic binding do that work for you!
- *"Anytime you find yourself writing code of the form "if the object is of type T_1 , then do something, but if it's of type T_2 , then do something else," slap yourself."* -- Scott Meyers
- **Liskov Substitution Principle, bumper sticker equivalent corollary: If your polymorphic method cares about the exact subtypes of its argument objects, either your method or the entire design of your class hierarchy is wrong to begin with!** Go back to the drawing board, stat! Danger, danger, Will Robinson! Arooga! Arooga! Do nothing else until you find and fix the source of this most horrendous of all possible [code smells](#) in object oriented programming.
- The Liskov substitution principle does not apply to **metaprogramming**, such as writing the `equals` method where it must be possible to talk about the type system and the exact type of some object.
- (Reflection and self-reference in general tends to conceptual paradoxes such as this one. For more of this sort of stuff, read the modern classic *"Gödel, Escher, Bach: An Eternal Golden Braid"*.)

3.11. Decorator pattern

- As powerful as method overriding is to express variability within a concept (for example, the difference between how various birds `fly()`), this technique would be absurd in some situations where it would force us to constantly repeat ourselves throughout the entire inheritance hierarchy.
- **Method overriding can only be used to express behavioural variation within methods that were originally included in the superclass.** However, it is impossible for any superclass architect to foresee all the future needs and ways that others will be using and later extending this class hierarchy, and the properties of the problem domain that this class should model and represent.
- For example, suppose that in the `Animal` hierarchy we later decide to allow the possibility for an animal to be **loud** so that a loud animal speaks in all uppercase. Unfortunately, the superclass `Animal` has already been released and countless other people have extended subclasses from it. Adding any new methods to `Animal` would break many such subclasses; some noisily, some silently.
- If we already have the class `Cat` and want to extend that to define `LoudCat`, well, that is just one class. However, we would have to write essentially the same code for all other types of animals of which we also wanted to create loud versions.
- Even worse, allowing an animal to also possibly be mirrored would force us to have four separate classes for `Cat`, `LoudCat`, `MirroredCat` and `LoudMirroredCat`. With respect to the number of such `boolean` properties that could vary independently, the number of subclasses would grow exponentially.
- A **decorator** (also called a **wrapper** in some materials) is a subclass **from the same class hierarchy** whose objects never exist on their own, but are always placed "in front" of existing objects

from that same type hierarchy. This **underlying object** is given to the decorator as a constructor argument, and stored in a **private** field.

- After creation, the decorator object is the new public face that represents the underlying object, as far as the rest of the program is concerned. This is analogous to how an attorney represents a client in legal proceedings. Please do not address my client directly; all access to the underlying object ought to be done through the decorator, as if the decorator were the original object!
- The methods of the decorator tend to call the corresponding method of the underlying object, occasionally modify the arguments down the line or the result up the line.
- Since the decorator type is itself a subtype of the original hierarchy, the decorated object can itself be further decorated to create arbitrarily complex computation pipelines, and express arbitrary variation within the same kind.
- The only real downside of using decorators is the classic **self problem**: since the object itself does not know that it has been decorated, during the execution of its methods that call some outside methods that require the object to pass itself as one of these arguments (such activity would be common in **callback methods**), this object does not know to pass the head of its decorator chain as this argument, but will pass itself.
- Variation of the self problem leads to confusing bugs when the same object is used both directly and through the decorator chain from different parts of the same program.
- A class that behaves otherwise like a decorator but comes from a different class hierarchy than the underlying objects that it decorates (usually because it has a different public interface than the original class hierarchy) is called a **facade** or an **adapter**.
- For example, `Scanner` is a handy and powerful adapter that offers a unified interface to any `String`, `File` or `Reader`, allowing us to read characters from all these sources the exact same way without having to care whether the underlying source is a string, file or a character stream.

3.12. Constructors and inheritance

- **Even though constructors are methods, they are never inherited to subclasses.** Even if the superclass has a dozen overloaded constructors, its subclasses have only the constructors that are explicitly written in these classes. You cannot therefore assume that subclass objects can be created using the same parameters that the superclass objects can be created.
- As before, if no constructors whatsoever are defined in a class, the compiler will automatically synthesize a do-nothing **default constructor** that takes no parameters. But as soon as you write even one explicit constructor yourself, no other constructors get synthesized automatically.
- When any subclass object is created, **the default constructor of all of its superclasses are executed down the line** starting from `Object`.
- The constructor arguments, if given with `new`, are passed only to the appropriate constructor of the actual class itself, even if some superclass happened to have a constructor that could take those given constructor arguments.
- Nothing guarantees that the superclass and subclass constructor, despite taking parameters of the same type, would use these given arguments in a semantically compatible fashion. (As an analogy, consider a method that takes a **temperature** and another method that takes a **height**, both methods taking an `int` parameter. Passing the argument 400 to both methods will have very different consequences!)
- Intuitively, every object starts out its lifetime as a raw `Object` and “grows” into the actual subtype during these stages of construction. The default constructor at each level initializes the fields

that were defined at that level, so the constructors of subclasses only need to initialize the fields that are defined down there.

- The subclass constructors can still re-initialize inherited `public` or `protected` fields, though, should that be necessary. (Of course you will never use any `public` or `protected` fields, but instead always provide proper getters and setters for the data kept `private`, *ja?*)

3.13. *Overloading vs. overriding: technical interlude*

- **Overloading**: the same class contains several versions of some method, with the same name but different parameter types. (Access modifier and return type can vary over these methods.)
- **Overriding**: the subclass redefines a method that it inherits from a superclass, with the exact same name and parameter list. (Access cannot be more strict, return type can be more specific.)
- In a method call of the form `x.foo(y)`, **overloading is resolved at the compile time** based on the **compile-time type** of the variable `y`. (Since objects do not exist at compile time, no decisions can be made during compilation based on the type you get from inspecting some object.)
- In a method call of the form `x.foo(y)`, **overriding is resolved at run time** based on the type of the **object** that variable `x` refers to. (Since the compiler can't know the runtime type of the object `x`, it can't possibly bind that method call at compile time.)
- An important exception to the previous rule is that if the method `foo` is `private`, the call is always bound at compile time, seeing that the intended version of the method must be the version in the same class, since `private` methods can be called in the same class only!
- This must be so, since otherwise you could accidentally override some `private` method of a superclass, and the whole concept of `private` would have no meaning or purpose.
- Calls to `final` methods can also be bound at compile time, since we know that its subclasses will not provide a different version of that method. However, this is done only for efficiency, since the program would still produce the same results even without such early binding.
- In calls of the form `super.foo()`, the call is also bound at compile time. This must again be so, since otherwise this call wouldn't work as intended in situations where the method that contains this call was being executed for some object that was created from some subclass further down the inheritance line.

Lecture 4: The Java Collection Framework

4.1. Interfaces as perfectly abstract superclasses

- In the theory of object-oriented programming, **multiple inheritance** means that a class can simultaneously extend **two or more immediate superclasses**.
- Any language that allows multiple inheritance must somehow put together the can of worms that breaks loose when the **subclass inherits the same method or field from two of its immediate superclasses**. This goes double when these immediate superclasses themselves inherited this method or field from their common superclass, creating the so-called **diamond problem**.
- For this reason, Java allows only a limited form of multiple inheritance where the class extends only one actual superclass, but then also **implements** any number of **interfaces**.
- An interface is an abstract class that **has no fields, and all its methods are abstract**. These restrictions solve the previous ambiguity lurking in unrestricted multiple inheritance.
- [The distinction between classes and interfaces exists only in the Java language in a rather redundant fashion.](#) The Java type system **makes no distinction whatsoever between classes and interfaces**, but **interfaces are syntactic sugar for abstract classes and behave as such**.
- An interface serves as a practical guarantee that **the class that implements it will have all the methods defined in that interface**. This is sufficient for writing polymorphic methods whose parameter type is that interface, since this allows the compiler to verify that the method calls made to that object inside the polymorphic method are correct with respect to the type system.
- Interfaces are the Java equivalent to the Pythonic spirit of **duck typing** in which the actual type of the argument object given to a polymorphic method is irrelevant as long as that object has the methods that the polymorphic method needs it to have. No gatekeeping by enforced credentialism is allowed in Java; as long as the object has the required methods, who cares where it got those methods from?
- At runtime, the calls to the methods defined in the interface are dynamically bound to correct subclass versions, same as is done with any other method calls. As long as the object knows its own type so that the JVM can bind the method calls to the correct versions of these methods, the caller itself doesn't need to know or even care about the object subtype!
- Whenever a class **implements** an interface, its subclasses also implement that same interface without you having to explicitly say so.
- Java 8 relaxed the definition of an **interface** slightly (well, actually quite a lot) in that methods no longer need to be fully abstract, but the interface can provide a **default implementation** for any of its methods. If the class that implements an interface does not implement that method directly or inherit an implementation from its superclass, this **default** implementation is automatically used to fill in the gap.
- Default implementations also allow adding brand new methods to an existing interface without breaking all the existing subclasses that implement the old version of the interface. This was necessary in augmenting the new advanced **computational streams** into the rest of the Java 8 language and its class libraries.
- Whenever some class inherits a default implementation from both some superclass and some interface, the implementation inherited from the superclass is used. To resolve the ambiguity in inheriting a default method from two interfaces butting their heads, the wise compromise nearly worthy of King Solomon is to require that class to implement that method itself.

4.2. The Java Collection Framework

- The **Collection Framework** defined in the `java.util` package is not just an educational example of class inheritance and polymorphism, but highly useful in practical programming for representing the data that the program needs to remember.
- Collection classes represent **dynamic aggregations** for which they provide three fundamental methods `add`, `contains` and `remove`, plus a couple of other methods built on top of these.
- At the interface `Collection<E>` that defines the behaviour that all collections have in common, the collection hierarchy branches in two directions `List<E>` and `Set<E>`.
- `Set<E>` collections behave like mathematical sets in two important aspects. First, duplicate elements are not possible, since once an element has been added as a member, adding it again any number of times doesn't make it any more of a member. Second, no internal structure is revealed to the outside world, regardless of which data structure is actually used to store the elements to make the dynamic set operations more efficient.
- The most important `Set<E>` implementations are `TreeSet<E>` based on **balanced binary search trees**, guaranteeing $O(\log n)$ worst case running times for both dynamic set operations and **order-based operations**, and `HashSet<E>` based on **hash tables**, whose three dynamic set operations are $O(1)$ on average, but no order-based operations are supported. (Theoretically the hash table operations could degenerate into an $O(n)$ linear time slowdown, but the probability of this happening in practice is far too small to lose even a wink of sleep about.)
- `List<E>` allows duplicate elements. To tell which one of the identical elements is which, it internally maintains these elements in some **linear structure** that can be **indexed** with the **accessor** and **mutator** methods that are concisely named `get` and `set`.
- In practice, by far the most common implementation of `List<E>` would be `ArrayList<E>` that stores its elements inside an ordinary Java array. To speed up adding new elements to the end, the **capacity** of this array can be larger than its current **size** so that there is slack space available for fast additions of new elements.
- Adding and removing an element inside an `ArrayList<E>` is an $O(n)$ operation, since the underlying array is not allowed to have any gaps between its elements. For algorithms that need to **add** and **remove** elements at arbitrary positions of the list, especially at the beginning which is the most inefficient case for `ArrayList<E>`, consider using `LinkedList<E>` instead.
- Since no royal road exists to quickly jump to the given position of the linked list, the random access methods `get` and `set` for `LinkedList<E>` silently degenerate into very inefficient linear time operations. Just like you wouldn't force a fish to climb a tree nor a snake to entertain people as a juggler, every data structure should be used according to its nature instead of forcing it to perform operations that are unnatural to it.
- A second downside of `LinkedList<E>` is its additional memory consumption compared to `ArrayList<E>` to store the exact same elements, since it needs to store a separate **node object** for each stored element, instead of merely maintaining an array of references to those objects.
- The advantage of a `LinkedList<E>` is its flexible ability to **add** and **remove** not just individual elements but entire sublists to and from the given position in a constant time, once you have navigated yourself into that position. Whereas an `ArrayList<E>` is akin to a rigid crowbar, a `LinkedList<E>` is like a chain of rings that can be quickly opened and closed to rearrange the list structure in ways that would be very sluggish for an `ArrayList<E>`.

4.3. Operations on collections

- The interface `Collection<E>` contains methods in accessing the elements in bulk. Unlike arrays, all standard library subtypes of `Collection<E>` come with decent implementations of both `toString` and `equals` methods out of the box.
- Same as the utility class `Arrays` for primitive arrays, the utility class `Collections` provides a bunch of `static` utility methods for collections for common operations such as **sorting**, **shuffling** and finding the **minimum** and **maximum** element.
- So that the `Collection<E>` implementations work correctly when using your class as their element type `E`, it is essential that the `equals` method of your class is a mathematically proper **equivalence relation**. This **semantic** requirement cannot be enforced by the compiler, but if is violated, the behaviour of the collection for such elements becomes unpredictable.
- For the **order-based collections** such as `TreeSet<E>`, the element type `E` must implement the interface `Comparable<E>` for the method `int compareTo(E other)` used for element comparisons. The **sign** of the integer result gives the result of the order comparison.
- The magnitude of the returned integer carries no semantic meaning, such as a large result meaning that the first element is a “lot” larger than the second one. For some orderings of elements, such as the lexicographic (“dictionary”) ordering of strings, any two different strings are “infinitely far away” from each other in that an infinity of strings lie between them!
- If the element type `E` does not implement `Comparable<E>`, the set can be given a strategy object that implements `Comparator<E>` to compute these element order comparisons.
- For the same reason, elements added in collections should never silently change their internal state in any way that can affect the `equals`, `compareTo` and `hashCode` methods. Immutable classes therefore once again offer great benefits for programming over mutable types.
- The problem of **shuffling** a list into a random order has a sneaky algorithmic pitfall in that the seemingly obvious and simple solution that pretty nearly everyone comes up with immediately when they first hear about this problem does not produce every permutation with the same probability. The correct [Fisher-Yates shuffle](#) algorithm is **unbiased** so that it produces each one of the $n!$ possible permutations with the exact same probability $1 / n!$ within the limitations of the **entropy** of the underlying **pseudorandom number generator**.
- As with many other randomized algorithms, the method `shuffle` in the `Collections` utility class takes the `rng` object as a parameter, so that this method will always produce the same result when executed with the same state of pseudorandom number generation.

4.4. Iterators

- To explicitly **iterate** over the elements of some collection, you must create an **iterator** object pointing at the beginning of that collection. An iterator acts like a smart “finger” that points to some element inside the collection, and this finger can **advance** one step at the time to the **next** element. (When using the subtype `ListIterator<E>`, you can also retreat one step.)
- Iterators in general do not allow **random access** of skipping to the desired position in the collection in constant time, since not all collection subtypes can possibly support this.
- As in Python, functions can receive iterators as their argument, and these iterators hide the actual type of the underlying entity that yields the actual elements. This way, we don't need to write multiple functions to do the same thing with different types of sources of data.

- Similar to Python, the iterator method `next` returns the current element, and simultaneously makes the iterator silently step to the next element.
- Any number of iterators pointing to the elements of the same collection can peacefully coexist simultaneously. These iterators can advance independently of each other, and are not being forced to march in some kind of lockstep, as if they were doing hard time under the Auburn system.
- Iterators allow us to write nice polymorphic methods that can operate on any `Collection` or `Iterable` without having to care about the exact subtype of that collection. This collection could even be **virtual** or **infinite** so that the iterator itself dynamically produces the elements of this virtual collection one at a time, the method using this iterator being none the wiser.
- (Since Java does not have **generators** the way Python does, iterators are the closest equivalent of Java to implement **lazy sequences** that compute the `next` element on the fly when actually needed, instead of **eagerly** constructing the entire sequence right away.)
- Each iterator object knows how the data is organized in the underlying collection, and hides all these pesky details from its users. Your polymorphic methods that receive iterators as arguments should never have to care about the exact type of the underlying collection that produces the sequence of elements.
- Iterating through an implementation of `List<E>` in the Java standard library is guaranteed to yield the elements in order of their position indices. The order in which the elements of `Set<E>` come out depends on the internal implementation of that set data structure and its iterators.
- The subtype `SortedSet<E>` adds the guarantee that iteration yields the elements in sorted order, and the further down `NavigableSet<E>` refines this concept with many useful **order-based operations such as minimum, maximum, predecessor, successor, floor and ceiling**.
- Iterators have an optional method to `remove` the element that was previously returned by the method `next`. Calling this method will succeed only if the underlying collection actually supports the element removal.
- The `remove` method can be called only once following each call to the method `next`, otherwise an `IllegalStateException` is thrown to indicate that the even though the method call is technically legal and provides arguments of correct type and semantic intent, the object is not in a state where it could receive such requests.
- Even though the handy **for-each loop** acts as **syntactic sugar** for all iterators just as it works for arrays, you must perform this iteration explicitly if you want to use the `remove` method.
- **Modifying a collection in any other way while it is being iterated is not guaranteed to work**, and depending on the actual subtype of the underlying collection, can either work correctly, cause strange behaviour, or cause a `ConcurrentModificationException` being thrown.

4.5. Strategy pattern

- Inheritance and method overriding are completely static concepts in that they are set in stone at compilation. The decorator pattern is fully dynamic, as objects can be decorated in different combinations on the fly during the program execution.
- As handy and flexible the decorator pattern is, it still is not a panacea but would be just as absurd in some situations. For example, suppose that you have some collection of elements whose elements you want to sort or analyze according to various different criteria, not merely the **natural order** baked into the element type.

- Since the objects are stored privately inside the collection that is given to you for sorting, you would not be able to decorate these objects for modified comparison purposes. Instead, the designers of the `Collections.sort` method were prepared for this often occurring eventuality by allowing a `Comparator<E>` strategy object parameter that performs the individual comparisons between objects of type `E`.
- Similar `Comparator<E>` strategy objects can also be given to `TreeSet<E>` to perform the element order comparisons needed to organize the underlying binary search tree. (The whole point of a `HashSet<E>` is that no order comparisons between elements are ever made, which allows hash tables to contain **heterogeneous** objects of multiple types.)
- An object temporarily given to another object or method to “fill in a gap” that has been left open inside it is called a **strategy object**. The same object or method can then be made to act in different ways simply by giving it a different strategy object each time. This sort of thinking illustrates the important principle of [dependency injection](#).
- The component classes of the **Swing GUI library** of the next module will also heavily use various strategy objects to tinker with the details of their behaviour. For example, `Stroke`, `Color`, `LayoutManager`, `Font`, `Border`, ...
- In the decorator pattern, both the underlying object and the decorator object come from the same class hierarchy. In the strategy pattern, the strategy object itself comes from a different class hierarchy than the object whose behaviour it modifies.
- The strategy object, same way as a decorator, can easily be programmed to do some internal bookkeeping or **logging**, which allows us to “crack open the hood” to send a little spy inside the object to take a little peek at what is going on in there. For example, a comparator designed to maintain a count of how many comparisons it has performed can measure the efficiency of a sorting algorithm from the number of comparisons that it performs.

4.6. Java Collection Framework as a teachable moment

- Even though all `List<E>` implementations support the indexing methods `get` and `set`, these methods are not generally guaranteed to operate in $O(1)$ **constant time**. The particular implementations of `List<E>` that guarantee this, for example `ArrayList<E>`, can make this fact explicit by implementing the **marker interface** `RandomAccess`. This allows polymorphic methods that receive lists as arguments to change their behaviour depending on whether the list object supports random access operations.
- A marker interface has no methods, but is essentially one immutable `boolean` value embedded directly in the type system. This makes this information available both at compile time (as opposed to storing this `boolean` value in a **named constant** inside the class) and runtime using type inspection with the operator `instanceof`.
- (Marker interfaces are now semi-obsolete, as **annotations** can be used in their place for the same purpose. Marker interfaces still have the advantage of being able to serve as a parameter type of a polymorphic method, though.)
- The `RandomAccess` interface is a horrendous violation of the principle that a polymorphic method should never have to worry about the exact subtypes of its argument objects. Yet any polymorphic method that operates on `List` arguments so that its logic cannot be implemented with sequential iteration without the methods `get` and `set` has to be prepared for this possibility.

- The only way to find out whether the given iterator supports element removal is to call its method `remove` and see what happens. Designers of Java language did not deem necessary to use a marker interface analogous to `RandomAccess` to reify this important piece of knowledge.
- Even then, nothing will tell you whether `remove` works for the given iterator in guaranteed $O(1)$ constant time (as it does for the class `LinkedList<E>`) or in $O(n)$ linear time (as it does in the class `ArrayList<E>`).
- Except as a warning example of how not to use class inheritance, we shall completely ignore the branches `Queue<E>` and `Deque<E>`. Their very existence as part of this God-fearing hierarchy is an abomination to begin with. Saying that some `Queue<E>` instance is also an instance of `Collection<E>` is no less absurd than claiming that some `Car` is also an `Engine`!
- Instead of `add`, use the method `offer` to try to add an element to the queue. This at least maintains the pretence that queues are not collections, but something else entirely.
- Queues, stacks, **deques** and **priority queues** can be immensely useful in many computer science algorithms, though, so don't go away with the mistaken conception that I opposed the existence of these concepts! **These data types just simply aren't collections**, which is why in any properly designed class library for data structures, they would form a completely separate class hierarchy of their own, same as the association maps `Map<K, V>` of the next section.

4.7. Association maps

- Each collection stores individual keys, whereas a **map** (or a “dictionary” or a “hash”) stores **associations** from **keys** to **values**. The `Map<K, V>` interface therefore uses two generic type parameters; `K` for the type of its keys and `V` for the type of its values.
- Thank goodness, `Map<K, V>` is a completely separate class hierarchy from `Collection<E>`. (Somebody had more sense back then than the creators of that blasted `Queue<E>` branch.)
- Dynamic map operations `get`, `put` and `remove` are analogous to set operations `contains`, `add` and `remove`, but operate on key-value pairs instead of individual keys. (These methods are intentionally given different names than the methods in `Set<E>` so that there is no confusion whether some line of code operates on a set or a map.)
- The underlying data structure of a `Map<K, V>` is some set of `Entry<K, V>` pairs where the set operations are performed on the key ignoring the value. Finding the associated value of the given key is guaranteed to be efficient, but finding the associated key of the given key is generally a slow linear time operation.
- (Analogy: using a phone book printed on a dead tree, it is easy to use **binary search** to find the given person by name to look up their phone number. However, the **reverse lookup** of finding the person who has the given phone number could not be done better than slow linear search.)
- The method `keySet` returns a **live view** to the underlying map as a `Set<E>`. These views share the underlying data structure of the map, so that any modification done to that data structure will also immediately affect the view.
- In an association map, **each key can be associated with at most one value**. Trying to `put` another value to a given key first removes the previous association, returning that old value as the result of `put`. To simulate some kind of `Multimap<K, V>` where each key can be associated with multiple values, implement it either as a `Map<K, List<V>>` or a `Map<K, Set<V>>`).

- In Java 8, additional `default` methods `replace`, `getOrDefault` and `putIfAbsent` are not just useful, but more efficient than first checking if the map contains the given key and then searching for that same key again in the data structure to access its associated value. Done as a bunch, these methods can now internally perform only one data structure search to locate the key.
- (Side exercise: Analyze the following optimization to speed up this common case of searching the same key twice: have the `Map` implementation cache the location of the result of its most recent search, and before searching for a key, check if that key is equal to the cached result to make such consecutive searches of the same key to work in $O(1)$ time. Yay or nay?)
- When counting how many times the given elements occur in the data that you are looping through, you normally use an array of integer counters. However, if you don't know beforehand what the keys are going to be, such as in counting how many times each word occurs in some big text corpus such as *War and Peace*), you can conjure a **counter map** from the type `Map<String,Integer>` for this bookkeeping. Then use the method `getOrDefault` to find out how many times each element has so far been encountered, and `put` that value plus one back to the map.

Lecture 5: Swing

5.1. AWT and Swing

- **GUI programming** was historically the first **killer app** of **object-oriented programming** in the eighties. The object oriented and polymorphic approach of the **Smalltalk programming language** made many things trivial that used to be mighty difficult in imperative programming.
- Most of the techniques and ideas of computer science are at least a decade or two older than even most working programmers would assume them to be, and some such ideas had to wait patiently for the hardware processing power to catch up with their potential. (This observation applies to all walks of life, not just to computer science.)
- In early versions of Java, the **Abstract Window Toolkit** class library essentially just provided **hooks** to the native GUI components of the underlying window system. This made the GUI in both Java **applets** and **applications** inconsistent across different environments, and limited to only to the lowest common denominator of all GUI engines.
- To remedy this, **Swing** was built on top of AWT, reusing and extending its non-component classes (e.g. **Color**, **Font**, **Image**) and replacing its obsolete native component types (e.g. **Button**, **TextField**) with its own component types (e.g. **JButton**, **JTextField**) that are rendered by Swing on three basic AWT components for a consistent **look and feel** on all platforms.
- Instead of hooking to native GUI components, Swing basically just **draws a picture of each component**. For the user, a picture of a component is just as good as a real component, as long as that picture reacts to all actions the same as a real component would! (“If it walks like a duck...”)
- (Religious people can notice the similarity to the concept of iconography. Yes, we call those things on your screen **icons** for a reason that most people have forgotten.)
- The computer processor that moves bytes around and performs simple computations on them. Everything on the screen is just pixels, and the higher-level division of these pixels into semantic-level entities is done by the human who is looking at the monitor. The mindless machine itself, kind of a **peacock without eyes**, is not aware of any such distinctions.
- Experienced computer users think about computers with high-level semantic distinctions. For such users the drop-down menus on top of the OS X screen are essentially different from the drop-down menus rendered by the web browser inside the web page.
- In an amusingly zen-like manner, a complete beginner who cannot see the distinction between these two types of drop-down menus has a more realistic view of the computer as a whole!
- Swing is an entire **framework**, not merely a library. To paraphrase the comedy stylings of [Yakov Smirnoff](#), “*When using a library, you call the methods. When using a framework, the methods call you! What a language!*”
- Names of Swing component classes start with a capital J, such as **JTextField** and **JButton**, to make it clear at a glance whether something is a Swing or an AWT component.
- A hefty number of common GUI component classes are included in the package `javax.swing`, and you can **extend** them to further customize their behaviour. The examples in this course always extend `JPanel`, a handy little component that doesn’t really do anything on its own, but acts as a blank canvas on which we can impose our own behaviour.
- All Swing components extend the abstract superclass **JComponent**, a massive class filled with all sorts of bells and whistles. However, Swing is well designed so that everything usually works

“out of the box” without friction. Swing also serves as a practical example of a multitude of object oriented programming **design patterns**.

- The **only setting that you must provide** for every Swing component is its **preferred size**, as Swing cannot possibly divine whether you want a little pin button or a full screen splash.
- Swing components can **contain** other components inside them. Instead of painstakingly positioning these components yourself, some **LayoutManager strategy object** positions and resizes the contained components for you. You can also set this manager to be null, and **free position** every component yourself with pixel precision with methods `setSize` and `setPosition`. (Or in real life, probably use a **GUI maker tool** to handle all the grunt work.)
- Swing components are objects in the heap just the same as strings and birds. Their only special property is the ability of the Swing engine to visually render them on the screen. Only the Swing **top-level containers** `JFrame`, `JApplet` and `JDialog` (and therefore all their subclasses) have the ability to exist on the user's screen on their own. All other Swing components must reside inside some top level container to be shown to the user.

5.2. Graphics rendering

- At any time, Swing can arbitrarily decide that some component needs to be redrawn. The components themselves have no control over this, and **must at all times be ready to draw themselves** on command the way that they currently look like at that particular moment.
- In general, **Swing is always the dictator in charge and can always do whatever it wants**. This is the only way that components written by different people in different times can work smoothly together. (If two components simultaneously tried to boss Swing around, what would happen?)
- The method `paintComponent` is called by Swing, passing a `Graphics` object as an imaginary sheet of paper to draw on. After this method has returned, Swing posts that `Graphics` object on the user screen. By definition, that is **what that component looks like at that moment**.
- Since the primitive `Graphics` class essentially has the capabilities reminiscent of some 1980's eight-bit home computer Basic, these days a more powerful `Graphics2D` subclass object is guaranteed to be given to this method as that argument. The `Graphics2D` class provides [a rendering engine with the capabilities](#) that we expect from our graphics in this millennium.
- You need a **downcast** to call the new methods of `Graphics2D` that don't exist in `Graphics`. Calling the methods that are already in `Graphics` requires no downcast, since dynamic binding is still in full effect to bind any such calls to the `Graphics2D` overridden versions.
- The coordinate system of the component has its origin (0, 0) at the **top left corner**. Note that following the usual computer graphics convention where the familiar coordinate system of mathematics textbook has been **mirrored** with respect to x -axis, **the y -coordinate values increase downwards**.
- Even though integer coordinates correspond to the pixels, rendering calculations are performed using higher precision `float` resolution. This not only affects both **antialiasing** and **colour blending** when using partially translucent colours, but will also make a difference if the coordinate system is **transformed** to be something else than the original integer pixel grid.
- Performing the position calculations with floating point arithmetic rather than integers is especially important when calculating **animation paths**, even if all object point coordinates in each animation frame then get truncated to the integer coordinate resolution for rendering. The more fine-

grained the motion, the larger the difference in the end result. Forcing all intermediate values to be integers introduces **time quantization artifacts** analogous to **aliasing** effects.

- (Mathematically, time is a third dimension same as the two dimensions of space on your screen, and behaves the same in calculations.)
- **Geometric shapes** are represented as objects from various subclasses of `java.geom.Shape`. They form yet another entire class hierarchy of its own.
- Conceptually, all shape objects reside on the infinite two-dimensional plane. They don't even know anything about Swing that renders their visual representations somewhere. Swing components then offer our human eyes a fixed-sized peephole to spy at some finite part of this infinite two-dimensional plane.
- To control the drawing, `Stroke` and `Paint` strategy objects should be given to `Graphics2D`.
- To allow creation of more advanced shapes, the `Area` decorator allows shapes to be modified with arbitrary combinations of **affine transformations (translate, rotate, scale)** and **constructive geometry operations union, intersection and difference**. The mathematics of these operations takes place under the hood of this decorator, so that you don't have to worry about the **linear algebra** behind such operations.

5.3. Nested classes

- Java allows classes to be **nested** inside each other. Nesting a class inside another class gives this nested class special privileges that it would not have if these classes were written separately as **top-level classes** inside the same project.
- If the class B is nested inside class A without using the modifier `static` (therefore making the class B an **inner class**), **objects of B cannot exist on their own**, but can only exist in the **context** of some object of the **outer class A**.
- Therefore, a good way to determine whether B should be an inner class of A rather than a top level class is to simply ask "Would it make any sense for an object of B to exist on its own, or is it innately tied to the framework of class A from which it cannot logically be extracted from?"
- The practical consequence of class nesting is that an object of B can only be created inside the instance methods of A. The object `this` in the execution of the current method call will become the context object of the B objects thus created.
- In the heap memory, each inner class object of type B contains **an extra unnamed reference** to this context object. Through this hidden reference, the methods of B can access all members of A (even those that are `private`) directly as if they were its own, with special syntax needed.
- Nesting a class even further inside a method also allows the **local class** methods to access all effectively `final` local variables of the surrounding method.
- Why only effectively `final` variables? Because the object constructed from the local class can still escape from inside the method and **continue to exist after the local variables of that method have vanished from the stack**. Therefore the JVM cheats a little and stores **defensive copies of the local variables** inside that local class object. If these variables are `final`, it doesn't matter whether you use the original variables or the defensive copies.
- Making a nested class `static` (or equivalently, defining it as a local class inside some `static` method) allows objects constructed from that class to exist on their own without any context ob-

ject as the mothership. Of course, the nested class methods then cannot access the outer class instance methods, since no context instance exists.

- If the nested class `B` contains a field with the same name as another field in the outer class `A`, the special syntax of `A.this` inside the methods of `B` refers to the context object. In these methods, the reference `this` always points to the nested class object itself.
- Theoretically, the level of nesting could continue arbitrarily deep, thus neatly illustrating the **zero-one-infinity** principle usually talking about data, but existing just the same in programming language design. However, in practice this level of nesting never exceeds two.
- (Similar infinite ladders of which we only ever use the first few steps exist elsewhere, for example, the number of dimensions of multidimensional arrays. Courses of **theory of computation** will later introduce many such infinite ladders, such as the **LL(*k*) parser hierarchy**.)
- Non-static nested classes are not allowed to have `static` fields, unless they are `final` named constants initialized at compile time. See the Stack Overflow thread “[Why does Java prohibit static fields in inner classes?](#)” for more discussion.
- The compiler will turn each class into a separate bytecode file. Since the **dollar sign** character cannot appear inside a Java identifier, the filename uses dollar signs as separators for nested names, such as `Foo$Bar.class` for the nested class `Bar` defined inside outer class `Foo`.
- An inner class can also be **anonymous**, albeit with a rather strange syntax. The compiler will name these anonymous inner classes with integers, such as `Foo$1.class`, `Foo$2.class` etc.
- These days, the **lambda operator** that we will discuss in chapter 11 of these notes can be used to turn anonymous inner class definitions into one-liners by having the compiler fill in all that rigmarole boilerplate of the class and method names and signatures.

5.4. Object factories

- If all constructors of some class are declared `private`, no new objects of that class can be created anywhere from outside that class. Sometimes this is precisely the restriction that we need.
- Instead of defining a `public` constructor, such a class rather defines one or more **static factory methods** to create objects of that type. For example, see the Swing strategy class `Border` and the corresponding factory class `BorderFactory` for these methods.
- A factory method has several advantages over an ordinary constructor. For example, the factory method can quietly return an object of some subclass type, or return an arbitrarily complex **decorated** version of the object, the caller being none the wiser about any of this but happily using the object as if it were an object of the original type.
- For **immutable types**, the factory can keep track of all objects that it has already created. If an object identical to the one being requested already exists out there, the factory can simply return a reference to that existing object, instead of redundantly creating another identical object.
- For mutable types, the class can also allow for the outside world to give back objects that are no longer needed, and store these objects internally in an **object pool**, conceptually visualized as a **recycling center** of sorts. The next time the factory needs to create a new object, it can recycle an old object from the pool, and re-initialize its fields to the new values.
- This mechanism can be quite valuable in a class that implements some branching data structure such as a **binary tree** or a **hash table** that constantly has to create and release little **node** objects that make up these structures with the **payload data** of elements stored in the structure.

- Especially in applications that constantly create and release small objects, using the object pool may even prevent the application from running out of heap memory when the new little object creation and release is so frequent that the JVM garbage collector cannot keep up. Even though garbage collection algorithms have improved massively over the past decades, there is no harm in organizing your program logic to not rely on these algorithms.
- The factory itself can also be an entire class hierarchy, with the superclass being an **abstract factory**. The canonical example is some game where the game engine receives a factory that creates new enemies in the game. Different levels of the game can then be defined using different subclasses of the enemy factory.
- An abstract factory is a special case of a strategy object, to be consulted in situations where a new object is needed. The algorithm that creates these objects therefore does not have to care about the specifics of how these objects come to life and what are the specific rules underneath that control this creation.

5.5. Event listeners

- When you want to react to some event that takes place, **polling** has its time and place, but it is usually rather inefficient to keep repeatedly asking whether something has taken place.
- **Swing event listeners** are a good example of the more sublime **Observer pattern**, in which any component can register any number of **listeners** that get notified whenever an event takes place.
- Each event listener **implements the corresponding interface**, thus guaranteeing that the event listener object contains the appropriate methods to react to such events.
- The event listener objects wait silently without wasting even one processor cycle, until they are actually activated by the call of the method that corresponds to the event that took place.
- The relationship between the components and the listeners is **many-to-many**, so that the same listener can listen to multiple components, and the same component can be listened to by multiple event listeners that will all react to that event. However, it is usually the best to **decouple** the components and their listeners to be in monogamous relationships.
- The best practice to implement the event listeners is to write them as **private** nested classes inside the component class that they are hardwired to listen to. Since the listener is written to listen to one particular component, it should not be exposed as **public**. (One-liner event listeners can also be converted to **lambda expressions**.)
- For this reason, **the once common shortcut of making the Swing component itself to be an event listener is just plain wrong and false advertising**. This component cannot possibly have any meaningful ability to react to those events coming from arbitrary Swing components.
- If the event listener does something that changes the way the component is supposed to look, call that component's method `repaint` to request Swing to redraw it at the next opportunity.
- Of course, there is no law saying that the event has to originate from a Swing component, since anybody can call the event handling methods of any object. For example, the `Swing Timer` objects generate action events at every tick of their simulated metronome.

5.6. Turtle graphics (extra material)

- The concepts of a **Cartesian coordinate system and trigonometry** would be a little bit too high level math for curious children of all ages aching to learn graphics programming but getting traumatic flashbacks from high school trigonometry to implement these geometric calculations.

- **Turtle graphics** allow drawing operations to be described in a more intuitive way that even a small child could understand and play act by himself by pretending to be some sort of cartoon turtle (like that one that looks kinda like Maggie Gyllenhaal that you have to draw to get in the correspondence art school), since these operations are done with respect to **relative positioning** instead of the grid of absolute Cartesian coordinates.
- Conceptually, a **turtle** is an object that moves around on the two-dimensional plane, internally remembering its **current heading** (as an angle) and **current position** (as x- and y-coordinates). The users of the turtle usually should not need to care about this absolute position.
- In the tail of the turtle there is a **pen** that can be either **up** or **down**. When the pen is down, the turtle leaves behind a drawn trail as it moves.
- The turtle can be given two basic instructions: (1) **move** a given distance to its current heading, or (2) **turn** in place by a given angle. All higher level graphics such as drawing a square, regular polygon or a star can be done by suitably combining and repeating these two operations.
- A more advanced turtle can also remember its previous positions by **pushing** them in an internal **stack**, and be commanded to return to its previous position **popped** from the stack. This makes it possible to render branching **treelike** structures and **fractals** where multiple branches split into different directions from the same point.
- Classic turtle graphics programming in the style of the 1970's **Logo programming language** never talks about the absolute position, but after initialization, purists perform all movements using only relative position operations.
- Traditionally, turtle graphics have been used to teach **imperative programming such as conditions, loops, procedural abstraction and recursion** in a very intuitive manner where the programming errors become visible as the drawing is different from what was expected. (In retrospect, that sort of instruction would have been suitable for the course CCPS 109.)

5.7. *Turtles for the new millennium (extra material)*

- Here in CCPS 209, turtles are used to illustrate the concepts of inheritance, polymorphism, dynamic binding, decorators and other central object-oriented programming concepts. In the instructor's example class hierarchy, the interface `Turtle` is a **facade** to `Graphics2D` that **translates** turtle instructions to the previously seen rendering primitives of `Graphics2D`.
- To make it easier to extend `Turtle`, the subclass `AbstractTurtle` provides a reasonable concrete implementation for all methods except `move`. (The subclass `BasicTurtle` then extends that method to draw a solid line exactly as it is told.)
- Whenever some interface has lots of methods, it is a good idea to provide such a utility class to make life easier for those who want to create their custom implementation of that interface, and provide some kind of reasonable implementation (which doesn't need to be maximally fine-tuned or optimal) for as many methods as possible.
- For example, the utility class `AbstractCollection` for the `Collection` interface implements methods such as `addAll` or `forEach` to use `add` and `iterator`, even if those latter methods are still **abstract** so that the concrete subclasses must then implement them.
- In Java 8, you can also define **default** implementations inside that interface itself. However, this approach would not work in this example, since these method implementations also need data fields, which cannot be defined in an interface. This would work better for `MouseListener` and such, where the **default** implementation of each method just needs to exist and do nothing.

- Concrete subclasses of `Turtle` represent turtles with different abilities and behaviour, such as a `HandDrawnTurtle` that leaves behind a line rendered in simulated “hand drawn” style in spirit of the `xkcd` webcomic. All these turtles are commanded with the same methods of the superclass `Turtle`. See the `TurtleDemo` example class for a demonstration of a spiral drawn this way, the same sequence of `move` and `turn` commands passed to six different turtles.
- Again, any polymorphic method that uses a turtle to do some drawing **does not, and should not, have to care which exact subtype of `Turtle` it is given as argument**.
- In several labs of this course, you will be extending and decorating the existing turtle classes in various ways to create your own customized types of turtles.
- [Breaking down a move recursively into a series of smaller moves and turns](#) makes it easy to render various fractals such as **Koch curve** or **Hilbert curve**. Such **fractalizing** is best done with decorators so that it becomes independent of other details of rendering. See the example classes `ZagTurtle` or `AntennaTurtle` that can decorate any existing turtle.
- When you write your own **fractalizer**, make sure that the **turns in each level of recursion add up to zero**, and the **subdivided moves add up to the distance that the turtle was supposed to perform**. Neither of these rules is any sort of law of nature or man, but makes the result more aesthetic and well-behaved.
- **Lindenmayer systems** (also called **L-systems** for brevity) are [a classic algorithmic technique](#) to generate a long and convoluted series of `move` and `turn` instructions for turtles to produce elegant fractal shapes, especially those classics that simulate plants or **space-filling one-dimensional curves**. Students interested in generative systems and theory of automata can look at the example classes `LSystem` (which itself is a facade for `Turtle`, which in turn is a facade for `Graphics2D`, which in turn is a facade to... oh, who knows, something internal to the Java virtual machine!) and `LSystemDemo`.

Lecture 6: Exceptions

6.1. Pre- and postconditions

- Every method has a set of **pre- and postconditions** that describe **the intended semantics of the method by describing what it is supposed to achieve and in what circumstances**. The method **implementation** itself is encapsulated into a **black box** from the point of view of the caller who can only see the end results once the method has terminated.
- The pre- and postconditions are not expressed in the Java language, since they are **semantic properties** that reside one level of abstraction above the language itself.
- The Java language itself already describes perfectly **what will happen during execution**, whereas the pre- and postconditions describe the intention of **why those things happen**. For this reason, such conditions are expressed in natural language in the comments and in the (**JavaDoc**) documentation of the method.
- **Always comment on the why, not on the what**. If your code is so complicated that you need comments to explain it to other people (note that this crew includes all future versions of yourself), that is a surefire red flag that should make you go back to the drawing board and redesign that code without trying to be super terse and clever.
- In programming, calling something “clever” is not always a compliment. Real programming is not [Code Golf](#); you don’t get cookies for writing the source code with the fewest characters!
- **Postconditions describe what the method promises to achieve**, that is, what is guaranteed to be true when the method completes its execution.
- Every method should have **at least one unit-testable non-trivial postcondition**, since otherwise that method has no reason to exist. If you can’t write a unit test for your postcondition, how do you even know if your method achieves that postcondition?
- For a particularly notorious example of a method without any testable postconditions, check out `System.gc()`. What could possibly constitute a test and an observable failure to pass it?
- (That whole Hail Mary pass of `System.gc()` is a relic from a simpler time of uniprocessor computers and more primitive garbage collection algorithms, now pretty much [the programming analogue of those “door close” buttons in elevators](#) that are not actually connected to anything real.)
- Some methods cannot guarantee that they will fulfill their postconditions under all possible circumstances. The behaviour of the method may depend on properties of the outside world beyond the reach of the method.
- For example, a method that reads data from a file cannot possibly do anything meaningful if the file doesn’t exist; the file is not readable; if the data inside it isn’t properly formatted; or if the file contains semantically nonsensical information such as the company personnel file with a person who is billion years old.
- **Preconditions describe what the method needs from the outside world to be able to work**.
- Unlike method postconditions, the set of preconditions may well be empty. In fact, **this would be a wonderful thing**, since such a method can never fail to fulfill its purpose! (Well, unless the method implementation contains a bug, but this is why we are here in the first place.)
- **It is not the method’s fault that the outside world failed to fulfill these preconditions**. Caller’s mistakes in fulfilling their end of the deal do not constitute a bug in the method, since the method has no powers to close the proverbial barn door after the horse has already escaped.

- In principle, anything could happen if the caller fails to enforce the preconditions to hold. Typically, a method **throws an exception** to indicate that preconditions do not hold.
- The opposite principles of **design by contract** and **defensive programming** are both valid approaches in checking preconditions. The former minimizes code, the latter maximizes safety.
- Even when the superclass method is **abstract**, **it can still define pre- and postconditions** that all concrete subclass implementations of that method should adhere to.
- Because pre- and postconditions are semantic properties of the code, no compiler can algorithmically verify in the general case that the method body will make its postconditions true for all of the legal argument combinations of that method. The responsibility for this lie solely on the shoulders of the programmer, the man positioned above the machine in a symbiotic relationship.
- Be careful with the postconditions of any method, since unlike humans, the cold equations that run the computer processor executing your method logic are not affected by what you “intended” to happen. For example, the postcondition “The elements of the parameter array are in sorted order after the method has terminated” would be trivially satisfied by a method implementation that fills its parameter array with all zeros.
- Sometimes people really, really wish that the method had some additional postconditions. If the current implementation of the method coincidentally happens to satisfy those wished-for postconditions, the grim error of **relying on undocumented behaviour** will become painfully evident once the method implementation changes in the future so that the method no longer guarantees those same assumed postconditions, making all bets be off once again.
- The novel phenomenon of “[Hyrum’s Law](#)” is the observation that once an interface has enough users, they will collectively depend on every aspect of the implementation, intentionally or not. Changing the implementation will always break something for somebody.

6.2. Liskov Substitution Principle

- As a rule, **inheritance should model an is-a relationship between the two concepts** from the problem domain that those classes represent. However, since this is again a **semantic property**, no compiler can meaningfully check or enforce this requirement.
- Men process meaning, machines process bits. For our minds, the future is uncertain and the past is fixed, whereas for a machine, the future is a deterministic railroad track while the past could branch to multiple directions.
- It is always possible to write nonsense in any language, no matter how high or low level, and indeed many do. The only alternative to such anarchy would be to define such a tight syntax and type system as to warrant the term “**programming in a straitjacket**”, first historically applied to the **Pascal** programming language.
- Many classic languages such as Smalltalk or Pascal (or for that matter, these days also Java) had many good ideas that were [revolutionary for the thinking and limitations of their era](#), but all these good parts have since been appropriated into the more current languages.
- The famous [Liskov Substitution Principle](#) (from hereafter, **LSP**) formulated by Barbara Liskov helps us to determine whether some class Foo should be made a superclass of another class Bar. How come `Car extends Vehicle`, but never `Car extends Engine` ?
- LSP is not a law of either nature or man, and can be freely violated at will. However, the LSP is a [SOLID engineering principle](#) that makes your polymorphic methods future proof.
- **If you want to write class hierarchies that allow creation of meaningful polymorphic methods properly and safely, you simply must follow LSP. Period.**

- Where LSP is not followed, chaos and disharmony will eventually stick their noses inside the project. Where LSP is followed, order and harmony will organically reign.
- Just by following the LSP, you will soon magically solve problems before those problems even have a chance to emerge to bug you! The classes that you create just seem to seamlessly click together as smoothly as your childhood Lego pieces, and the general nature of these connections allow you to easily experiment with new combinations of ideas that you never would have even thought of, had you been programming in some less organized imperative language.
- Oversimplifying only a little bit, LSP can be expressed as the requirement that **whenever you override a subclass method, its postconditions cannot be made weaker** (but they can be equal or stronger), and its **preconditions can't be made stronger** (but they can be equal or weaker).
- This seems pretty simple and obvious, but same as all great ideas, it sure was not simple and obvious five minutes before somebody stated it explicitly.
- Intuitively, the subclass version of the method must work in any situation where the superclass method promises to work, and in any of those situations, achieve at least as much as the superclass method promises to achieve.
- If this principle is followed for all methods, **the subclass object can be substituted in place of a superclass object**, and the rest of the system will still work just as well as before, or even better.
- **The subclass object can be passed to any polymorphic method that expects a superclass object as its parameter**, and the polymorphic method works happily with that subclass object.
- **LSP is both necessary and sufficient for polymorphism to work! Without it, there can be no meaningful polymorphism that would gain anything real!**
- If the substitution principle can't be followed for even one method of `Foo`, this immediately means that no matter how much your intuition might insist otherwise, you must resign to the stark reality that the class `Foo` cannot be a superclass of `Bar` to begin with. You should not impose the inheritance relationship on classes just to be able to inherit some particular method.
- Even though both `Cartoonist` and `Gunman` have a method `draw`, trying to combine these two under a new common superclass `Drawable` results in this superclass having a method `draw` with no meaningful postconditions that both subclass implementations of `draw` could fulfill. Hence the method and therefore the interface `Drawable` have no reason to exist.
- [The most famous example of this phenomenon](#) is how the classes `Circle` and `Ellipse` ought to relate to each other in inheritance. Even though every circle is an ellipse, it turns out that neither class can be made to be a subclass of the other, assuming the existence of **mutator** methods that can turn a circle into a non-circle.
- This classic example shows that **the subclass relationship is a different thing from subset relationship**, even though these two concepts often coincide and are therefore occasionally mistakenly assumed to be different names for the same thing.
- Once again, **immutable objects make everything easier to reason about**. If no mutator methods exist in the class, `Circle` can be a subclass of `Ellipse` the exact same way that `Car` is a subclass of `Vehicle`... which also requires the tacit assumption that there do not exist any mutator methods to attach wings and a jet engine to some car to turn it into an `Airplane`!

6.3. The exception machinery

- The Java Virtual Machine executes bytecode instructions one at a time, having no idea of the larger reality that those bytecode instructions together create. As long as those bytecode instructions are legal, the execution continues.
- Errors exist in the level of intended **meaning** of the program, and the computer has no mind to divine this intended meaning. Therefore, it is up to the programmer to somehow recognize and indicate that things have gone wrong in the execution of that program. This is done with the pre- and postconditions of methods at the semantic level.
- What should the poor method do in the situation where it is asked to do something where it is logically impossible for that method to achieve its stated postconditions?
- Older languages would have that method return a special **error code** value (typically minus one in the spirit of classic C, assuming that an actual result of that method cannot be negative) or change the value of a global **flag** to indicate an error, from which the caller is expected to realize that the returned result is actually meaningless nonsense.
- However, from the software engineering perspective, this is a defective approach in several ways. The caller would have to remember to check this flag for errors after every call.
- When some error occurs, there is no easy way to propagate the error handling up the call chain to the level where the cause of error actually took place.
- This is especially true in library methods used in many different applications. Such a library cannot possibly know anything about its caller and what exactly there ultimately caused this error to manifest, let alone the steps needed to handle the problem.
- In Java and other modern languages, a failing method instead **throws an exception** to indicate a failure. **No result whatsoever is returned**, since **no meaningful result can logically exist**.
- An exception is **not any kind of special return value**, but **throwing an exception** is a separate mechanism from **returning a result**. Hence the intentional use of two different verbs.
- To terminate the method execution abnormally by throwing an `Exception` object, use the keyword `throw` instead of `return`. (In Python, the keyword to achieve this is called `raise`.)
- Different types of exceptions indicate problems whose underlying causes are different. This way, the caller knows how to handle these problems in different ways.
- For example, a method that reads data from a file can fail because the file does not exist, or that it exists but contains syntactic nonsense, or that it exists but contains semantic nonsense. All these situations would probably be handled in very different ways by the caller, so the exception type must be used to indicate the level of abstraction in which the error takes place.
- The method documentation should fully describe the possible exceptions that could be thrown from that method, and the preconditions whose violations that make them occur.
- Exception objects can (well, **should**) also contain a human readable **message string** that explains what went wrong, to help the programmer to further identify and analyze the problem to debug it.

6.4. Catch or release

- Whenever a method throws an exception, the caller has a choice of two possible ways to proceed; they can either **catch and handle** the exception, or **let that exception keep flying** to the previous level of the method call chain.
- Uncaught exceptions terminate the thread that threw them. JVM will then spit out the **stack trace** of that exception to the **standard error stream** of that JVM process.

- The caller could also catch only some types of exceptions and duck others to let them fly on. In general, the caller should only catch those exceptions that it can meaningfully handle at that level.
- Be especially wary of the common **exception anti-pattern** of catching all exceptions with an empty handler, this way quietly sweeping your problems under the proverbial rug. However, when the situation forces you to do this, name that exception object as **ignored**.
- To catch an exception thrown by a method, surround the block of code that contains the call in a **try** block. This **try** block can be followed by one or more **catch** blocks for different types of exceptions, their bodies being executed if an exception of the particular type gets thrown.
- With or without the **catch** blocks, the **try** block can also be followed by a **finally** block that is guaranteed to be executed before the execution leaves the preceding block, regardless of which way the execution exited the block and whether any exceptions were thrown.
- A common pattern is to acquire and use exclusive resources in the **try** block, and then release them in the **finally** block, to guarantee that that resource will be released. This pattern is common enough to allow the specific **try-with-resources** shorthand analogous to the **context managers** in Python.
- As a general rule, **you should never allow any exceptions to fly out of a finally block**. Since **a method can throw only one exception at a time** (analogous to how a method can return only one result at a time), this second exception would fly out while the original (and obviously more important) exception would vanish into the silent night.

6.5. Checked and unchecked exceptions

- Unlike most other languages, Java idiosyncratically divides its exception classes into two separate kinds of **checked** and **unchecked** exceptions.
- Checked exceptions **must either be caught**, or if the method lets them keep flying to the previous caller, **declared with the throws clause in the method signature**, so that the compiler can enforce that the caller will also catch or declare that exception.
- **Checked exceptions are used to announce that the caller did not fulfill the preconditions of the method.** Therefore the caller is held fully responsible for handling that exception, and the compiler will enforce this with the same strictness as with the rest of the type checking, such as the correctness of the method return type or the parameter types.
- (No, the annoying enforcement of checked exceptions cannot be “turned off just for this time, *maaan*” any more than the enforcement of method argument and return types. It’s an essential part of overall type safety.)
- When overriding a method inherited from the superclass, **the subclass version must not declare any new checked exceptions** that were not already explicitly declared by the superclass version. (It can declare fewer of them, though, if you are certain that some particular exception will never be thrown by this particular subclass version of that method.)
- **A method that has no preconditions can’t possibly throw any checked exceptions.** It simply has no excuse to fail, no matter what the caller and the rest of the outside world might have done before the method assumed control of execution.
- **Unchecked exceptions** of the Java library such as `ArrayIndexOutOfBoundsException` or `NullPointerException` indicate **a bug in the method itself**. These can also be caught and handled, but doing so is voluntary and not required by the compiler.
- It would make no sense to require all user code to be able to catch and handle all possible bugs that could exist in all methods that they call. Usually the caller could not do anything meaningful

anyway, not even in principle. All computation implicitly depends on the assumption that the previous computation steps have achieved their promised postconditions.

6.6. The *Throwable* class hierarchy

- The root superclass of all exception classes is the superclass `Throwable`. Only objects created from the subclasses of this class can be thrown as exceptions.
- `Throwable` branches into two subclasses, `Error` and `Exception` to distinguish between two fundamentally different phenomena that can go wrong during program execution.
- The subclasses of `Error` are used to indicate that **something has gone wrong in the JVM itself** in a way that prevents the Java bytecode from being reliably executed any further. The most common reason for this in practice is the JVM running out of stack or heap memory.
- Your program cannot do anything about JVM errors in principle, since nothing guarantees that your program statements will any longer even be executed in any sort of accurate fashion. You can try some Hail Mary pass such as an emergency save, but the odds of that are between slim and none (and as they say, Slim just left town).
- The subclasses of `Exception` are used to indicate that the method cannot proceed in a meaningful fashion. There is nothing wrong with the JVM, it is just that the bytecode that it executes was asked to do something that is logically impossible at the semantic level.
- Again, note how the entire concept of error exists at the semantic level about the intention of the program, not in the low level bytecode instructions that the JVM mechanistically executes. As far as the JVM is concerned, nothing has gone wrong in the execution of these bytecode instructions.
- The language rule to distinguish between checked and unchecked exceptions is that **all subclasses of `RuntimeException` are unchecked**, whereas all other exceptions are checked.

6.7. Assertions

- Same as in all other walks of life, when testing your code, **it is always better to detect the inevitable failure sooner than later**. There is no point in stubbornly keeping painting until it has become undeniable that you have painted yourself stuck in a corner.
- Detecting failures early makes it easier for the programmer to trace back the problem to the point that actually caused that failure, instead of having to slough back through a long diagnostic journey from the observed effect of the bug manifesting itself to its hidden underlying causes.
- Since our computers now easily perform hundreds of millions of invisible operations during the time between the bug silently occurring and the bug finally manifesting itself in the results, it sure would be nice to be somehow able to detect the bug the moment that it actually happens!
- In testing and debugging code that is under development, it might come handy to seed in **assertions** about the data that is handled by the program in its various points.
- Whenever you believe that some `condition` is logically inevitable at some point of your program, put your proverbial money where your mouth is and `assert` that condition explicitly by writing `assert condition;` to that point in code as a silent tripwire for incorrect logic.
- Whenever the execution reaches the `assert` statement, the `condition` is evaluated. If it is `true`, nothing happens. If it is `false`, an `AssertionError` is automatically thrown to crash your program right on that spot immediately and noisily.

- In one of the most perplexing blunders of Java design, **assertions are disabled by default**. To turn them on, you need to use the `javac` command line option `-ea`.
- If your original reasoning was correct, you will never hear about that assertion again. However, if your original reasoning was flawed (or perhaps it was correct, but some later changes to the code made it incorrect), you will find it out the next time the problem actually manifests itself.
- **You should only ever assert things that are 100% under control of your code** so that them being false would logically entail a genuine bug in your code. **You should never assert anything about the outside world** or the method arguments that it passes to you. Your method cannot possibly have any divine powers to dictate how the outside world must behave.
- All this is, of course, also heavily dependent on where you draw the boundary between the “system” and “the outside world”. Note also that if some subsystem of your code is intended to be taken out and reused as a library inside different systems, it should not assert anything about the original system inside which it was first developed.
- Modern IDE’s understand the assertions and can use them to offer **refactoring** suggestions. For example, the assertion that some reference cannot be `null` may allow simplification of some conditional structures. Some other assertions may help establish that some later condition is always true or false and can therefore be eliminated altogether.
- Assertions should be used only in the development stage and **turned off when delivering the release intended for the end users**. For this reason, the `condition` of an `assert` should never do any actual work for the program, such as calling some method that has useful **side effects**.
- **JUnit** defines its own slew of assertion methods that won’t crash the entire program, but these methods throw a custom exception to terminate that test and tell the JUnit’s internal bookkeeping mechanism that that particular test somehow failed.

6.8. Some exception patterns and anti-patterns

- The exception handling mechanism in the JVM is not optimized for speed. Furthermore, each exception object builds up and carries the entire stack trace with it, so throwing and handling an exception can be a relatively expensive operation.
- **Do not use the exception mechanism as an extralinguistic control structure**, sort of a “return on steroids” that can jump back over multiple levels of method calls in one conceptual step.
- For the same reason, you should generally not use the famous Pythonic programming idiom of “It is easier to ask forgiveness than permission” in Java code, but rather follow the more conservative principle of “Look before you leap”.
- Python uses exceptions such as `StopIteration` as a legitimate control structure, but this approach goes heavily against the philosophy of the Java language.
- Instead of rolling out your own custom exception types talking about your low-level concepts, try to use [the existing exception subclasses in the standard library](#) whenever possible. Standard exception types such as `IllegalStateException` and `IllegalArgumentException` carry well-understood messages about the situation that they describe.
- The number of different exceptions thrown by a method should not depend on the number of different ways that the method can fail, but **on the number of different ways that the caller can fix the problem**. If two different problems are repaired the exact same way, they should be indicated with the exact same type of exception.

- When throwing an exception, always throw the most specific exception type that is available to you. Doing otherwise might be termed “passive-aggressive programming” in which the program tells you that something is wrong, but won’t tell you what that was. (“You *should* know.”)
- As they say, “No” is a complete sentence. Answering “no” is not a failure that needs to be indicated with an exception, when “no” is the proper and correct answer to the question that was being asked by calling the method. For example, a method that is supposed to return the list of prime numbers between the given `start` and `end` values should return an empty list when that interval contains no prime numbers, instead of throwing some kind of exception.
- Also, things that must **always inevitably happen during the execution of your program** no matter what (for example, reaching the end of a text file whose contents are read into your program) by definition are not exceptional, and thus should not be handled with exceptions.
- If an exception flies through more than one level of encapsulation, it should be translated to a semantically appropriate type on each level from the second level up, since otherwise that exception will expose the lower level implementation details to the outside world and break encapsulation.

Lecture 7: I/O Streams

7.1. Streams and their decorators

- In the Java standard library, the **stream classes** `InputStream` and `OutputStream` represent some underlying mechanism to move **raw unsigned bytes** from some point A (**sender**) to some point B (**receiver**). The subclasses of these classes represent actual concrete implementations for different entities that A and B can actually be.
- At the stream level, the transmitted data is a sequence of raw unsigned bytes without any higher level semantics assumed or attached. The universal streaming mechanism transports those bytes from sender to receiver without the slightest care of whatever those bytes might “mean”.
- Regardless of the higher level semantics of some data, everything (either code or data) inside a [von Neumann architecture computer](#) is nothing but bytes that represent basically whatever you say that they represent. (This is necessary for **metaprogramming** tools such as compilers, interpreters, **transpilers** and **profilers** to be able to exist in the first place.)
- **Every stream in Java is unidirectional** from the sender to the receiver. If you want to perform **bidirectional** communication (for example, over a **TCP/IP internet socket** connection), you need to establish two separate streams to opposite directions.
- If the receiver is busy doing something else when the sender is sending bytes, these bytes won’t vanish, as the stream will internally **buffer** the data until the receiver is ready to read them.
- If the receiver tries to read data before the sender has sent it, the `read` method will **block** until either the next byte or the notification of reaching the end of the stream becomes available.
- The methods `read` (in receiver’s `InputStream`) and `write` (in sender’s `OutputStream`) operate on `int` parameters, since the primitive type `byte` in Java is a **signed byte** with the range `-128, ..., +127`, whereas the raw **unsigned byte** has the range `0, ..., +255`. A four-byte `int` can easily represent and store all these values.
- Either party can `close` the stream at any time, after which no further transmission will be possible through that stream. Once the sender closes the stream, **the receiver will always receive -1 at the remaining reads** to indicate that no more data is forthcoming.
- Note that no exception is thrown when the end is reached, since reaching the end of the transmission is not an exceptional situation.
- (“*Once you have received the message, hang up the phone.*” — Alan Watts)
- A stream built on top of some exclusive resource such as a file or an internet connection should be closed when it is no longer needed. This `close` operation should be put in the `finally` block to guarantee that it will happen no matter what happens during the allocation and use.
- These days you should use the **try-with-resources** shorthand.
- **The stream class hierarchy heavily uses decorators to add higher-level capabilities** such as **compression, encryption, data filtering and conversions, and checksum calculation** to existing streams. Each such decorator needs to be written only once, and it will still work with any past, present or future stream from that hierarchy, once again vividly illustrating the power of polymorphism and dynamic binding.
- Streams are allowed to internally buffer the data to speed up their work. The method `flush` in `OutputStream` forces the stream to send forward all bytes that it is possible for it to send.

- (Special cases such as compression streams might have to wait for more bytes to arrive for it to know the value of the next outgoing byte. Some **sponges** such as **sorting** cannot write out even the first byte until they have seen all incoming bytes.)
- Calling this method should be **necessary in two-way communication**, such as over an Internet TCP/IP connection, where you have to wait for an answer to your previous message before you can know what message to send next.
- Flushing might also be necessary during perfectly normal console output, to ensure that all outputs initiated before the crash get properly written out before the process terminates.
- Closing or flushing the stream decorator will also close or flush the underlying stream, assuming that the decorator was written properly to guarantee this.

7.2. Readers and Writers

- For the common situation of transmitting **Unicode textual data** from point A to point B, the class hierarchies of `Reader` and `Writer` work **exactly the same way** as the classes `InputStream` and `OutputStream`, the only difference being that the basic unit of data being transmitted is now one **Unicode character** instead of one unsigned byte.
- `Reader` and `Writer` are **facades** to some underlying `InputStream` and `OutputStream` that actually transmits these characters encoded as bytes. Inside a computer, **everything is made of bytes**, and whatever semantics or aggregation we humans imagine on top of those bytes exists only in our minds as we voluntarily put ourselves on a leash to use some bytes only in some particular way, as opposed to machine code where “Do what thou wilt” shall be the whole of law.
- The **facade** subclasses `InputStreamReader` and `OutputStreamWriter` can be used as explicit plugs to convert characters to bytes and vice versa, whereas commonly used classes such as `FileReader` will internally create and maintain their own private byte streams.
- One advantage of using these explicit conversion facades between characters and bytes is that you get to choose the encoding scheme. In practice, the [UTF-8](#) encoding is the most popular since it needs only one byte to represent every character in the old **ASCII** range with the code points from 0 to +127 (as a happy bonus, any old text file that contains only ASCII characters is a legal UTF-8 file as is), which comprise most of the English language text.
- Curiously, the methods `read` and `write` of `Reader` and `Writer` are exactly the same as those of `InputStream` and `OutputStream`! This is because the Unicode characters are encoded as integers, and the special value `-1` is again used to denote the end of the stream.
- These class hierarchies are still separate, since these methods have different **semantics** despite having the same formal signature. **Semantics exist above the syntax**, so two classes could even have identical public interfaces, and yet behave completely differently! For example, consider two classes of `Gunman` and `Artist`, both with the method `public void draw()`.
- Since it would be annoying to write out or read in text only one character at a time, the handy decorators `PrintWriter` and `BufferedReader` are usually placed in front of an existing writer or reader.
- `PrintWriter` offers the powerful and familiar methods `print`, `println` and the C-style `printf`, whereas `BufferedReader` offers the method `readLine()` to read the input one line at a time, the special value `null` now indicating the end of data.
- `BufferedReader` allows us to read text data in the spirit of the **Unix command line** so that each data file consists of **records**, each record stored in one line. Each record is broken into **fields**

between the chosen **separator character** that cannot be part of the content of any field. This flexible encoding allows arbitrarily long data fields, and the entire data file can easily be edited by any text editor or further transformed by the plethora of Unix command line tools.

- For more advanced **tokenizing**, **parsing** and error handling of character input emanating from many different kinds of character sources, these days we tend to use the powerful **Scanner facade** that can read not only from `Reader`, but from a `String`, directly from a file, or from an arbitrary **character sequence**.

7.3. Dependencies between classes

- The **Open-Closed Principle** states that classes should be **open for extension** (inheritance) but **closed for modification**. Once a class has been successfully compiled, **its source code should never be needed for writing other classes or subtypes of that same class**.
- We have already seen this principle in action when creating our custom subtypes of standard library classes such as `ArrayList<E>`.
- One should avoid **reuse by copy-paste**, the lowest stage of software engineering maturity. In fact, the more you copy-paste your own code within the same project, the more you should feel being in the state of sin, and need serious purification and repentance by means of **refactoring** your classes and methods to deliver you from this sinful state.
- Besides, combining similar methods and blocks of code into their most general parametric form also helps you better understand the underlying phenomenon. This allows you to improve your logic in ways that would have been impossible with the original implementation. All kinds of modifications become much easier to spot and execute whenever every truth about the program is expressed in exactly one place in the code.
- The dependency inversion principle is based on the idea that **there should never be a need to recompile any class**, except when the internal implementation of that class changes to an updated better version, or new functionality is added to the class in the next version of the entire library.
- Class **A depends on** class **B**, if editing and compiling class **B** makes it necessary to recompile class **A** even though its source code has not changed. For example, if **A** contains a field or a method whose parameter type is **B**, then **A** depends on **B**.
- If **A** is a subclass of **B**, then **A** depends on **B**, but not the other way around. In fact, a superclass or an interface should never ever depend on any of its subclasses.
- To keep all classes properly closed for modification, **each class should depend only on classes that are less likely to change**.
- Otherwise, making just one small change to a single class that everyone and their brother depend on would trigger the recompilation of all classes that depend on that class around the world!
- Counterintuitively, **the higher level abstractions are more solid and less likely to change than concrete classes**, opposite to how we normally think about concepts of solidity and concreteness.
- Therefore, as instructed by the **Dependency Inversion** principle, **all classes should depend on high level abstractions, and never depend on any concrete low level implementations**. In the physical world, we prefer to build our homes on solid rock instead of mud and quicksand.
- In general, try to **program to the interface**, that is, make the parameter types of your polymorphic method as high level and general as possible. For example, if you write a method that expects an `ArrayList<E>` as its parameter, ask yourself if this parameter *really* needs to be

specifically an `ArrayList<E>`, or whether it could be just any old `List<E>`, or even merely some kind of `Collection<E>`.

- (However, a method that could handle any `Object` in a meaningful fashion is probably somehow wrong in other ways. Outside metaprogramming, such methods should not really exist in modern Java. Methods that can accept and return arbitrary objects are usually screaming to be turned into **generic** versions of those methods.)
- Once you hardwire your class to depend on low level implementation details of some other class, you **create a dependency where there is not supposed to be a dependency**, and any future change in such decisions will force you to rewrite your class. **This is why we also use encapsulation to make all implementation details private so that other classes can't get too attached to them.** This gives us the freedom to change any of these implementation details later without having to worry about how this change will radiate to the outside world.
- Once the `public` interface of some class is used by enough people, that interface must be considered to be set in stone so that it may never again change in ways that are not backwards compatible. The `public` interface of any class that is intended to be used by others therefore must be very carefully designed, since it has to be correct the first time.
- (In designing the public interface of an abstract superclass, you should always also create at least three concrete subclasses of it to ensure that all aspects of that interface can be realized in practice. You should also write at least three polymorphic methods whose parameter type is the abstract superclass, to ensure that the interface is complete so that it is possible to do things with it without any extra information about the actual subtype.)
- Implementations of methods can always be changed and optimized later, but changing the public interface would break all subclasses that implement it.
- Since Java 8, any interface can always be expanded with new `default` methods without breaking any existing code that depends on that interface.

7.4. Reflection

- The **dynamic binding** mechanism of Java and similar languages allows you to call the method `fly` with `x.fly()`; without knowing or even caring about the exact subtype of the object `x`. However, the superclass that defines the method `fly` must be available during compilation.
- The `instanceof` operator allows you to check whether some object is an instance of some class that is known to you at the compile time. However, you cannot use this operator to check the type of an object against a class that is determined at runtime.
- Various **metaprogramming tools** such as **BlueJ** or **JUnit** require a much stronger ability to crack open and inspect any arbitrary object and determine its type and contents, even if this type is from some class hierarchy that did not even exist at the time that the tool itself was compiled.
- In the Java type system, **all classes are also objects**, instances of the **metaclass** called `Class`. (Some wag should turn this into an [Abbott and Costello](#) style “[Who's on first?](#)” rapid talking vaudeville comedy act between a Java programmer and a Pythonista.)
- To prevent an infinite chain of turtles all the way down, the class `Class` is an instance of itself, having given birth to itself at the time when the Java language itself sprung to existence.
- **An object is said to be a class if it allows new instances of itself to be created.** For example, `String` is a class, since the string object `"Hello, world!"` can be created from it. However, `"Hello, world!"` is not a class, since no object can possibly be an instance of it.

- Objects can be cloned with the `clone` method, but that creates another identical instance of the same type. Cloning requires that the class implements the `Cloneable` marker interface to indicate that an identical but separate copy of an object can meaningfully exist.
- Every time the Java Virtual Machine loads up the bytecode of some class, it automatically creates the **class object** to represent that class, filling it with the information that it knows about what that class contains.
- To dynamically acquire the class object of an arbitrary object `x`, call `x.getClass()`. The method `getClass` is defined to be `native final` in the universal superclass `Object`.
- The class object for the class `Foo` known at compile time is given by `Foo.class`.
- The class object can tell you what fields, methods, constructors and nested classes it contains, using the **wide reflection** classes `Field`, `Method`, and `Constructor`. These classes have exactly the methods that a reasonable person would expect. You can essentially crack open any object or class given to you and find out everything about everything that it contains.
- Note that `private` is only a **safety** feature in Java programming, but has no **security** powers of any kind, since it can always be bypassed with reflection.
- The field, method and constructor objects even allow fields and methods of an arbitrary object to be accessed and modified dynamically, even if those members are `private`. However, this is rather slow, and same as with other reflection features, should really only ever be used in **metaprogramming**. If you know the type of some object, just access its fields and methods directly and let the compiler enforce your type safety and optimize the memory access for you.

7.5. *Serialization of arbitrary object structures*

- **Serialization** is a powerful mechanism in Java that uses reflection **to convert arbitrary objects into a series of raw bytes** that can be transmitted along streams. These objects can be stored in files and databases for **persistence**; the object outlives the computation process that created it.
- Two facades `ObjectOutputStream` (with `writeObject`) and `ObjectInputStream` (with `readObject`) can transmit not merely arbitrary objects, but arbitrarily large **object structures** in one swoop.
- Serializing an object will automatically follow all references emanating from that object, and **recursively serialize all objects hanging from it**, creating a **deep copy** of the entire structure. An entire collection with all the objects stored inside it can be serialized with one method call!
- The receiving end will automatically construct a new object structure that is **isomorphic** to the one that was serialized, keeping track of objects that it has already received to detect aliasing and directed loops of references. For this reason, serialization does not get into an infinite loop even if the object structure contains **cycles of references**, that infamous bane of all programming with pointers and references.
- (Serializing an object that was constructed from a non-static nested class will therefore also serialize its outer class context object, and from there every object reachable from that. Beware!)
- In Java 11, **serialization was officially deprecated as a security hole**. Good riddance. It was trying to solve the problem at the completely wrong level of abstraction to begin with.
- In general, when you need to store persistent data, it should be encoded into some **language-independent standard representation** such as **XML** or **JSON**. These representations use **raw text as their universal data encoding** and can therefore be meaningfully read in all languages (unlike

Java object serialization binary format that can only work in Java), and be further processed with various tools or Python scripts.

- The only restriction for serialization is that all involved objects must implement the **marker interface** `java.io.Serializable` to allow it. If the serialization is given an object that is not an `instanceof` this marker interface, the serialization fails and throws an exception.
- Some objects are so inherently tied to the particular Java Virtual Machine and the underlying platform that saving them to files for future access makes no sense, nor would transporting them over the Internet to some completely different platform!
- To force serialization to ignore a reference that points to an object that is not `Serializable`, you can define that reference to be `transient`, yet another esoteric Java keyword. (In the receiving end, that reference field in the copy of the object will be `null`.)
- Serialization only serializes the object data, but not the executable class bytecode of the class the object was constructed from. (Sure, the ability to transmit executable code would have been awesomely cool among consenting adults, but also an even more horrendous security hole.)
- To protect against the possibility of the receiving end having a different version of the class and its bytecode, the serialization also attaches a **64-bit checksum of the class bytecode** into each object that is serialized.
- Perhaps a bit surprising, all enums are `Serializable` in Java. Of course the JVM is smart enough to enforce the uniqueness of each object from that enum within that JVM process.

Lecture 8: Generics And Other Nifty Features

8.1. Minor language improvements back in Java 5

- During the compilation of a class, the compiler sees all the classes in the current package (which is the **default package** in absence of a **package** declaration) and the package `java.lang`.
- Ordinary `import` makes either one class or all classes in some package **visible to the compiler** so that you don't need to write their fully qualified names every time you use those classes. Similarly, a `static import` makes one or all **public static** members of some class visible so that you can use them directly without having to prefix their names with the name of that class.
- Of course `sin(x)` looks less clunky than `Math.sin(x)`, and `sort(arr)` looks better (at least more Pythonic) than `Arrays.sort(arr)`. (These days, `ArrayList<E>` also comes with an internal `sort` method.)
- Method overloading can be used to create versions of methods for different types of parameters. However, **variable length arguments** (or **varargs**) allow the same method to handle any finite number of arguments, within some internal bounds of the compiler and JVM.
- A method can be defined to take an arbitrary number (that is, **zero or more**) **varargs** of a given type. This vararg parameter must always be the last parameter in the method parameter list, and is denoted syntactically with an **ellipsis** (...) after the parameter type.
- This mechanism works pretty much the same way as `*args` in Python. Since Java does not have **keyword arguments**, it has no analogue of `**kwargs` in Python.
- (As a sidebar, Python 3 now also has an actual **ellipsis operator** in the language as a built-in constant symbol. It currently does nothing in the core language, but it can be handy **syntactic sugar** for various current and future extensions such as **numpy**.)
- Inside the method body, the varargs are automatically collected into an array of the same name. This array can even be empty, as it should be when the caller has no varargs to give. (In programming, **zero is always a possibility** that our code must be prepared to handle.)
- The varargs can also be given as an array, which allows one vararg method to call another vararg method and pass its arguments to it, no matter how many were originally given.
- In the Java standard library, vararg methods are perhaps most often seen as **factory methods**. The canonical name for such methods is `of`.
- Perhaps the two most commonly used vararg methods in the present Java standard library are `printf` in `PrintWriter` and `asList` in `Arrays`.

8.2. Autoboxing

- For every primitive type, Java standard library defines a **wrapper class** to represent **a value of that type as a proper object**. For example, `Integer` for `int`, `Double` for `double`, etc.
- Wrapper classes also have **static** methods and constants for common operations with values of that particular type, such as `Character.isWhitespace` or `Double.isInfinite`.
- Especially since the Unicode standard defines tens of thousands of characters with a multitude of properties, you should always use the utility methods of `Character` to check for these properties to guarantee full portability and internationalization. There is no point trying to reinvent these complex wheels within wheels yourself.

- Before Java 5, conversion back and forth between the primitive and the wrapper types had to be explicit. From Java 5 onwards, the compiler silently inserts this conversion when needed when the compilation would otherwise fail due to a type error of using a primitive when a wrapper was expected and vice versa, so the primitive and the wrapper types can now be used **almost** interchangeably... if it weren't for a couple of nasty little pitfalls!
- **Wrapper classes are immutable**, and therefore very inefficient when you need to iterate through a large range of values one at a time. If you use a primitive type, the same small group of memory bytes are always reused to represent the next value, whereas using the wrapper class necessitates the creation of a new object each time to represent the next value.
- (Every immutable type such as `String` has this exact same problem, which is why we don't make all data immutable, unlike those bearded weirdos in the **pure functional programming** school of thought. If your complicated data structure really needs to be immutable, a **linked list with insert and remove from front** is the best structure. **Trees** are also good, assuming that you can share identical branches.)
- Equality comparison `x == y` does not perform the intended value comparison when `x` and `y` are both wrapper objects. The more fundamental language rule of Java that **the operator `==` always compares the memory addresses of the objects** takes precedence over autoboxing.
- To make this even more confusing, the class `Integer` **internally caches** the boxed versions of values in the range from `-128` to `+127`. This works out swimmingly when you, for example, create a `Map<Something, Integer>` to keep count of how many times you have seen different values of `Something`, since up to 127 possible `Integer` objects are shared through this entire map, creating a new `Integer` object only when some count reaches 128.
- (This is a good example of the **flyweight** design pattern, where identical immutable parts of a large collection of objects are shared instead of creating redundant copies. The canonical example of productive use of this flyweight pattern would be a **rich text editor** with font and other info potentially per character, but most of the time identical for all characters, unless you are writing some sort of ransom note.)
- In comparing `x == y` where only one variable is a primitive, the wrapper is sensibly unboxed.
- A primitive field `boolean x` will be automatically initialized to `false`, whereas a reference field `Boolean y` will be initialized to `null`. Unboxing `null` will fail and throw an exception.

8.3. Annotations

- By definition, a **comment cannot affect the compilation or runtime behaviour of the code**, since the **compiler always skips every comment** without looking at its contents. (More specifically, each comment is silently converted into a single whitespace character.)
- **The JavaDoc tool** is the opposite in that it skips the code and reads the specially marked comments to generate the HTML documentation pages.
- **Annotations** are kind of “smart comments” that can be attached to classes, fields and methods, and even to other annotations in the case of **meta-annotations**.
- Annotations exist during both compilation and runtime, and unlike comments, can affect both the compilation and the program runtime behaviour, the latter via the `getAnnotations` method.
- Annotations are themselves types. In addition to **the standard annotations** defined in the language, you can easily define your own annotation types with `@interface`.

- Syntactically, annotation types are distinguished from classes by prefixing their name with @, the name “at” of this symbol coincidentally being an abbreviation of “annotation type”.
- In practice, `@Override` is surely the most commonly used annotation. It can be applied only to methods, and it causes compilation to fail noisily if the annotated method does not actually override the superclass method. Such a mishap typically happens because of a typo in the method name, or when you accidentally use a more specific parameter type in the parameter list, and therefore accidentally **overload** the method instead of overriding it as intended.
- Metaprogramming tools such as **JUnit** use class and method annotations to decide what to do with those classes and methods. For example, JUnit recognizes the test methods from their `@Test` annotation. (Before annotations came in, each test method name had to begin with the word “test” for it to be recognized by JUnit reflection mechanism, which made automated testing clunky and prone to silent typos.)
- Even an ordinary program can use reflection and class loading to load up new classes on the fly and then use annotation information to determine what to do with these classes.
- To combat the famous “[billion dollar mistake of software engineering](#)” of `null` references, some tools associated with Java 8+ introduce a handy **type annotation** `@NonNull` that allows the compiler to verify that the value of a variable thus annotated cannot never be `null`. Such a variable must be initialized with something that is also `@NonNull`, and every assignment afterwards must have a right hand side that can be proven to be `@NonNull`.
- (This is similar to the distinction between the guaranteed non-null and possibly-null types `Foo` and `Foo?` in **Swift**, although more annoyingly verbose.)
- Various other annotations have been proposed [more or less seriously](#).
- Java 8 also introduced the generic type `Optional<T>` that is a wrapper that either contains an object of type `T`, or denotes the absence of such an object in a controlled fashion. A modern method whose return type is `T` should never return `null`, but return an `Optional<T>`.

8.4. Basic generics

- So far, the rule for static type checking of Java has been that **all data that the program talks about must be given an explicit type at the compile time**. The compiler enforces that data is used only in ways that are meaningful for that particular type. This prevents programs from crashing at runtime because of type violations.
- The downside of this type safety is that since we don’t want to repeat ourselves, the Collection Framework methods must operate on `Object` parameters and return values, which then necessitate result **downcasts** that by their nature break down the type safety.
- In a **generic** class, some types used in the code (for example, method return type or the type of a field) are written using **placeholder names** “to be filled in later” at a time when that class is actually used. The placeholder names are given as **type parameters** in angle brackets after the class name, conventionally using a single capital letter, such as `ArrayList<E>`, to distinguish them from the names of actual classes.
- No actual class should ever have a single-character name. You may therefore always assume a single-character name to be such a placeholder name when reading through somebody’s code.
- The choice of the letter can be used to convey semantic information about the purpose of that type parameter. For example, `E` stands for element, `K` stands for key, and `V` stands for value.

- The compiler will still enforce the discipline that inside that class, these unknown types are used in a consistent and legal manner based on what is known about them. For example, we don't need to know what actual type `E` is instantiated with to know that it surely must be legal to assign from one variable of that type into another variable of that same type.
- The user code can then apply **type arguments** to create as many different **type instantiations** of the generic type as are needed. The two different type instantiations `ArrayList<String>` and `ArrayList<Integer>` of the same generic class are two separate types that cannot be assigned to each other either way.
- In the instantiation `ArrayList<String>`, the method `get` has the return type of `String`, so **no downcast is needed** when assigning the result of this method to a `String` variable. In the instantiation `ArrayList<Integer>`, that same method will have the return type `Integer`.
- The **diamond operator** `ArrayList<>` tells the compiler to fill in the type argument from the surrounding context. This saves a lot of annoying typing especially when the type argument is itself a complicated generic type, something like `ArrayList<List<Set<String>>>` for an arraylist whose each element is a list of sets of strings.
- Of the four possible combinations generic-generic, generic-nongeneric, nongeneric-generic and nongeneric-nongeneric, **all four are good and meaningful for a superclass and its subclasses**.
- **Enums, throwables and anonymous inner classes** cannot be made generic at all, but doing so would make zero sense anyway.
- An individual **method can also be made generic**, independent of the generics of its surrounding class. When you call a generic method, **the compiler can in most situations infer its type arguments**, so you don't need to give them in the call but simply call the method the normal way.

8.5. Erasure and its consequences

- Java generics very much resemble C++ **templates**. In practice, most of the time in practice there is no real difference in their use in both languages.
- However, an important difference between Java and C++ is that in C++, the compiler produces a separately compiled version for each type instantiation of the generic type. This can result in **code bloat** in the generated binary executable. On the other hand, this allows the compiler to customize and optimize each separate type instantiation to the maximum extent.
- In Java, the **distinction between different type instantiations exists only at the compile time**. Once the compile-time type checking is complete, **all generics are erased** so that the separate instantiations actually become one and the same **raw type** produced by the compiler.
- The **raw type** is the generic type without any type arguments.
- Since the code was verified during compilation to obey the “invisible fences” of each type declaration, removing these invisible fences does not change the fact that this code will still obey those very same fences at runtime! Dynamic binding of Java method calls based on the object's runtime type makes each separate type instantiation use its correct subclass method at runtime even after all the information about the different type arguments has vanished.
- (**Seriously**, think carefully about that previous item. **Its message is very important**. For the exact same reason, the primitive types don't need to actually exist in the runtime once the compiler has produced the correct machine code instructions to operate on such variables.)
- Even if you instantiate the same generic type with a hundred different type arguments in your program, the size of the resulting bytecode file remains unaffected.

- However, because of the way that type erasure works, a generic class `Foo<T>` cannot (1) use the type parameter `T` in any `static` member in any role, or (2) use `new T()` (nothing even guarantees that `T` has a `public` no-argument constructor), or (3) use `x instanceof T`. At runtime when the type `T` no longer exists, what could these runtime operations possibly do?
- Generic classes can have `static` fields as long as their type is independent of `T`, since one copy of that `static` field will be shared by all type instantiations of the same generic class, and must therefore work with every actual type instantiation of that generic class.

8.6. More on raw types and reification

- For every generic type, the corresponding **raw type** exists in the type system both during and after compilation, since it is needed for interaction with all legacy code that predates generics. This is unfortunate but necessary, unless we could reboot the language from scratch.
- (Many things in Java would be different if the language were redesigned from scratch by dropping the backwards compatibility. This was done with the creation of **Scala** and **Kotlin** as byte-code-compatible redesigns of Java to make the language more suited for this millennium.)
- Types that exist at runtime in Java are said to be **reifiable**.
- However, **raw types should never be used in new code**, and the compiler issues the “unchecked” warning if you do that, since **raw types break all type safety**. Anything goes!
- In the type system, **raw type can be assigned back and forth between any type instantiation of that generic type**. This allows us to easily create, say, an `ArrayList<Bird>` reference pointing to an object whose type is `ArrayList<BankAccount>`.
- To avoid the impossible situation where a bank account is expected to fly or a sparrow is expected to accept a deposit of 500 dollars, **the Java compiler adds a silent downcast (and the corresponding runtime check)** to every return value assignment of a generic method. The downcasts that we had to write explicitly before generics are still automatically generated there but hidden in the bytecode, wasting a handful of processor cycles.
- Method `Collections.checkedCollection` returns a decorator that checks the type of the added object at time of addition, instead of at possibly much later time when that element is accessed with an unsuccessful downcast.
- In Java, **you cannot create a primitive array whose element type is any generic type**, even if that generic type is fully instantiated. In Java heap, each array object must know its element type, but obviously this would be impossible when the element type does not exist at runtime!
- These days we can write a generic method `void <T> foo(T[] arr)` for that same purpose.
- (Java arrays must know their element type at runtime because of the most idiotic design decision of Java language of making its arrays **covariant**. Since `String` is a subtype of `Object`, the type `String[]` is a subtype of `Object[]`. See the next section.)

8.7. Inheritance relationships over generic types

- Keeping the eye carefully on the prize, the purpose of inheritance is to allow writing polymorphic methods that can operate on all subtypes of some abstract superclass, instead of having to painstakingly write a separate version of that method for each of those subtypes. We sure don’t want to give up this powerful ability just to have generic types in return!

- It is not obvious that we would even have to give up anything. Same as how we can have a polymorphic method `void foo(Bird bird)` that can be given any `Hawk` or a `Sparrow` instance, could we not write a method `void foo(Collection<Bird> birds)` to operate on any collection of any kinds of birds, so that it could be given an `ArrayList<Albatross>` or a `HashSet<Hawk>` instance as an argument?
- The method itself compiles just fine. However, calling it with an argument of either of the previous types fails to compile with a mystic error message about incompatible types.
- **Liskov Substitution Principle** reveals the counterintuitive reason behind this paradox: even though `Hawk` is a subclass of `Bird`, and `ArrayList` is a subclass of `Collection`, their generic combination `ArrayList<Hawk>` cannot be a subclass of `Collection<Bird>`!
- You can add an instance of `Sparrow` into a collection of birds, but you can't add that instance into an arraylist of hawks. **Therefore the latter collection cannot possibly be a subclass of the former**, since LSP requires the subclass to be able to do everything that the superclass promised to be able to do in its documented pre- and postconditions!
- At least `ArrayList<Bird>` is a subclass of `Collection<Bird>`. No problem there, as long as the generic type arguments are **identical**.
- The problem is that `Collection<Bird>` is able to contain all kinds of birds all together like in some newspaper cartoon, whereas `ArrayList<Hawk>` can only perch hawks inside it. These two must therefore be separate types. Trying to pass one as the other is no less absurd than trying to pass, say, a `BankAccount` instance as a `Sparrow`.
- Because arrays predate generics in Java, Java arrays were designed to be **covariant**, meaning that `Object[]` is a supertype of `Bird[]`. This makes zero sense and really, really, really ought to not be this way, but it was made as a compromise to allow polymorphic methods that can be given any array object as argument, instead of having to write a new version of that method for every possible array element type.
- If generics had been in Java since day one, they would be so much more sensible and better. Alas, you go to war with the troops that you have, not the troops that you wish you had.

8.8. Bounded types

- To implement the polymorphic methods as nature intended, we need a type that carries the meaning of “A homogeneous collection of some particular but presently unknown kind of a bird to be decided later”, instead of the meaning “A heterogeneous collection that can contain all kinds of birds simultaneously” carried by `Collection<Bird>`.
- **Bounded types** allow us to define exactly such types, with either a **wildcard** question mark or a type parameter name bounded either above (with `extends`) or below (with `super`) to restrict which types are possible as instantiations of that type parameter.
- `Collection<? extends Bird>` is a collection of some particular subtype of `Bird` that we don't know at the time other than that it will be exactly one particular type. (Perhaps that type will be an owl, but “whoo” knows?)
- A polymorphic method whose parameter type is such a bounded type can be given an arraylist of albatrosses, a hashset of hawks, or any particular instantiation that anyone could ever think of. The problematic method call will no longer cause a compile time type error.

- The bounded type provides enough information about the unknown type so that the compiler can allow calling the methods of `Bird` to the elements of the argument collection. Dynamic binding ensures that the correct subclass version of that method is executed at time of the call.
- Unlike with a parameter of type `Collection<Bird>`, the polymorphic method cannot add an object of type `Owl` to the `Collection<? extends Bird>`, because nothing guarantees that this unknown subtype would specifically be an owl instead of, say, a sparrow or an eagle.
- To allow such addition, use the lower bound type `Collection<? super Owl>` instead.
- All bounded types behave as abstract classes in that they **can only be used as types of variables and method parameters**, but they cannot be used to create actual objects.
- However, even though you can't create an object of type `ArrayList<?>`, you **can** create an object of, say, type `ArrayList<List<?>>`. The element type of this arraylist is **fully known**: a homogeneous list of some unknown element type. (Yeah, read this item again a few times.)
- A type bound can even be recursive, such as `List<T extends Comparable<T>>` for lists whose element type is something that is order comparable with itself. This is a way to require the elements of the list to have some capability, and otherwise they can be as diverse as you want.
- Hardcore Java generics enthusiasts can try their hand in deciphering the intention behind the doubly recursive type `Enum<E extends Enum<E>>`.

	<code>ArrayList</code>	<code>ArrayList<?></code>	<code>ArrayList<Object></code>
Can create objects of using new	Y	N	Y
Supertype of <code>ArrayList<String></code>	Y	Y	N
Exists at runtime	Y	N	N
Guarantees runtime type safety	N	Y	Y

8.9. Enumerated types

- Normally, once a class has been compiled, any number of new objects can be constructed from it. However, sometimes that class corresponds to some concept in the problem domain that is inherently **closed** so that the possible objects are fully known and set in stone at compile time, and it would be a grievous semantic error to attempt the creation of new objects.
- For example, consider a class whose objects represent the seven days of the week. There should exist exactly seven objects for these seven days, and never more.
- To define a class that represents this kind of closed concept, use the keyword `enum` instead of `class` in its definition to define it as an **enumerated type**. Enumerated types are classes for which the `new` operator has been disabled so that trying to use it is always a compile time error.
- The old way before Java 5 to implement this sort of a class was to make its constructor `private` to ensure that no new objects can be created from the outside code, and then define a `public static final` named constant inside the class body for each legal value of this type to refer to the object that represents that particular value.
- An even older way in languages before the era of object oriented programming was to use `int` as this enumerated type, and then simply define that 0 stands for the first particular value, 1 stands

for the second value, and so on. However, encoding enumerations as integers provides zero type safety since any `int` variable can be assigned any integer value regardless of our noble intentions that only the first n natural numbers should be used.

- Furthermore, looking at an `int` variable we cannot tell whether it is semantically meant to be an actual integer or to encode an enumerated type. (For example, addition makes sense for integers of the first kind, but makes no sense for integers for the second kind.) Having multiple enumerated types in the same program also allows arbitrary “cross pollination” between these types to give birth to mysterious bugs that the compiler cannot prevent.
- Many methods in the Java standard library still are relics of this kind of thinking; their **option constants** are given as integer vales declared as named constants inside that class. Had these enumerated types been part of the Java language since day one, these option constants would have surely been enumerated types instead.
- In the beginning of the `enum` body, you must first list the names of the objects to be automatically created by the JVM when the class bytecode is loaded in memory. This is actually just syntactic sugar for those `public static final` named constants that the compiler turns these freely floating names into when it is compiling an `enum`. Since these objects are really **named constants**, they are conventionally named in **all capitals** the same as other named constants.
- The enumerated type is still a real class and can define arbitrary `static` and instance fields and methods, even constructors. The constructor arguments are given in parentheses after the name of each individual value.
- Inside every enumerated type, the compiler will also automatically generate a useful bunch of additional instance and class methods such as `ordinal()` and `values()` that allow these enumeration objects to be converted to and from integers for the purposes of old-timey operations, and iterated one at a time with a for-each loop.
- Enumerated types tend to be immutable, but this is not a law of either nature or man. Even if the set of objects is fixed, those objects might still change their state during execution.
- As the example of using an `enum` to represent the planets of the solar system illustrates, be careful in making assumptions about some problem domain concept being inherently set in stone. Even definitions that all of humanity agrees on today might still change in the future, subject to the [Lindy effect](#).

Lecture 9: Concurrency Fundamentals

9.1. Concurrent execution of threads in JVM

- It is useful to distinguish between the two terms **concurrency** and **parallelism**; the former meaning that **multiple logical computations are being executed partially interleaved** (as opposed to **sequentially** waiting for one computation to end before the next one is started), and the latter meaning that these multiple computations are **simultaneous in the physical sense**.
- All parallelism is concurrent, but not necessarily the other way around. Concurrency is a meaningful concept even inside a classic single-processor machine sealed from the outside world.
- The **Java Virtual Machine** is one of the **processes** running inside your computer. The execution of JVM simulates an imaginary computer architecture that executes **bytecode instructions**, this imaginary computer safely residing in a **virtual sandbox** inside your real computer.
- Of course, this “real” computer could be just another virtual machine, with no theoretical upper limit how many levels of virtual machines could be piled up on top of the actual physical computer that is run by the physical universe... which may or may not itself be a computational simulation running inside some higher level universe!
- As Descartes famously realized, you can be certain of one thing that you exist, but can never truly know which level of The Matrix you currently exist in. Is reality a simulation?
- The data memory inside the JVM is divided into the shared **bytecode** and **object heap** areas shared by all threads, and **local variable stack areas** separately for each thread, jealously guarded against improper use by the JVM.
- The JVM loads up the class bytecode files in the bytecode area as needed. There are no **introspective** bytecode instructions that could access or modify the contents of the bytecode itself. It is important for both safety and security that a Java program cannot modify itself during its execution, except by expanding itself by loading more classes in a controlled fashion through the JVM services.
- The execution of the bytecode program is divided into separate **threads**. Each thread has a **program counter** that points to the position in the bytecode that the execution of that thread is currently at.
- Each thread must have its own **stack of local variables** because different threads are usually executing different parts of the same bytecode program.
- **All threads share one and the same object heap** so that the communication between Java threads takes place by **reading and writing shared objects**.
- Concurrency using **shared memory** under the same physical hardware makes concurrency simpler than general **distributed computation** where the concurrent entities are executed in separate computers connected in a **network** that introduces an unpredictable delay to **message passing** between the concurrent entities.
- Even something as elementary as knowing the exact number of participants might not be available to any particular entity. (Does the computer that you currently use to read this document need to be aware of the total number of computers connected to the Internet?)

9.2. Java execution model

- The JVM starts up by creating the **main thread** that starts its execution from the main method of the program, and the **garbage collector daemon thread** that looks for unreachable objects in the heap and releases the memory used to store them.
- **When only daemon threads are running, the JVM process terminates.** Unless you have created additional threads yourself, this rule effectively means that the JVM terminates once the program's `main` method returns, which is exactly what you are accustomed to so far.
- New threads can be dynamically launched during the execution of the program. This should not be done willy-nilly in real programs, since launching a new thread is a relatively expensive operation, seeing that it requires setup work such as allocating the **thread local stack** in the memory available to the JVM process.
- In the classic single-processor system, the JVM keeps one of its threads **active** while the rest are paused and standing still. The JVM executes the active thread by repeatedly reading in the next bytecode instruction and simulating its effect on the stack and the object heap. The execution then automatically continues from the next instruction, unless the current instruction was some kind of (either conditional or unconditional) jump.
- At any time, the JVM is free to perform a **context switch** to make any other thread active. **The threads and the Java program code that they execute have no control whatsoever over this.** The context switches can happen at any time, as far as the threads are concerned.
- Threads are not guaranteed to get **execution time slices** in any kind of **round robin** fashion.
- For our slow human eyes and brains, it seems like all these threads are truly executed in parallel. Analogous to watching a movie at 25 60 still images per second, our slow human visual system cannot tell apart the individual execution time slices.
- This sort of **preemptive multitasking** works much better than the more primitive **co-operative multitasking** where each thread always has to explicitly give up its turn.
- Even in a correct program, it would be difficult to decide where to place the instructions to give up your turn. When bugs are present, some malfunctioning thread stuck in an infinite loop will also freeze all other innocent threads.
- Some programming languages allow threads to be broken into finer co-operative **coroutines** or **fibers** that operate on a co-operative multitasking discipline, so that each coroutine has to explicitly give up its turn. Good use for this technique are various **producer-consumer systems**.
- (In modern Python, coroutines are best handled with **generators** that **yield** elements of some virtual data stream one element at a time, passing the control back and forth between the generator object and the object that consumes these elements.)
- In the computer operating system, different **processes** operate the same way, just one level of abstraction higher in a preemptive multitasking model. (Hopefully in this day and age nobody uses an OS that uses co-operative multitasking for its processes.)
- The semantic difference between “threads” and “processes” is that **threads share memory**, whereas each process runs in its own separate memory protected from other processes. No process can read or write the memory of other processes either due to error or malice, but all communication and control takes place through the operating system services.

9.3. *Threads in Java*

- It is extremely easy to launch a new thread in Java: simply create a new object from `Thread` or any custom subclass of `Thread`. This class is special in that the actual threads inside the JVM process always correspond one-to-one to the instances of `Thread` in the JVM object heap.
- Concurrent programming would become really annoying if every new thread always started executing from the `main` method. Fortunately, each thread is given a `Runnable` object as a constructor parameter whose `run` method it will then execute.
- To actually make the thread to start executing the `run` method, you need to call its method `start`. Once started, the **daemon status** and **priority** of that thread are set in stone.
- From the `run` method, the execution can go to whatever other parts of the program that it can reach through the method calls. Once the execution returns from the top level call of `run`, the thread terminates, and cannot be restarted.
- Launching and managing threads yourself is a bit too low level thinking for high level object oriented programming. Instead, you rather create one instance of `ExecutorService`, a handy utility that juggles threads behind the scenes so that you don't have to.
- You can submit any number of `Runnable` tasks to the same `ExecutorService` instance. They get scheduled and executed in some **asynchronous** order by the threads that are privately held by the service.
- At the submission of each task, you immediately get back a `Future<T>` object. This object acts as a **receipt** that you can later use to inquire about the status of your task, cancel the task, etc. Unlike an ordinary return value of type `T` that contains the result of a completed computation, the `Future<T>` object represents the result of a **future computation that is yet to complete**.
- The interface `Runnable` and its method `run` are deficient for modern coding, since you can't even return the result from a `void` method, nor can you override this method to throw any checked exceptions, since the superclass version didn't declare any.
- A modern interface `Callable<T>` with better method `T call()` throws `Exception` can be used in place of `Runnable` for tasks given to an `ExecutorService`.
- After you submit a `Callable<T>` task, the `Future<T>` receipt can also be used to query the exception thrown by the failed task, which is not possible with `Runnable`.

9.4. *Critical sections and race conditions*

- A computational operation is **atomic** if the outside world cannot observe any intermediate stage of data inside an atomic operation. As far as the rest of the world knows, an atomic operation either looks like it hasn't even started, or that it has completely finished, with nothing existing in between those two points in time.
- Very few operations in Java are truly atomic. Assignment between `int` values is one of them.
- Even an assignment from an eight-byte `long` to another `long` might happen only halfway through, with a context switch taking place in the middle so that the `long` variable becomes a chimera that contains four bytes from the source and four bytes of its original value!
- When two threads access and modify the same object with **non-atomic operations** that could be interrupted in the middle by an unlucky context switch, the resulting execution interleaving can cause the object and the surrounding computation to be left in an inconsistent state.

- (The related term **failure atomicity** means that if an exception is thrown inside a non-atomic operation, the object returns to the state it was before that operation began. A failure atomic operation therefore happens either completely or not at all. Failure atomicity is rather important in **distributed databases** that contain, for example, banking information.)
- **Immutable objects** are always **thread safe** since they cannot be modified, and thus need no synchronization regardless of how many threads access them concurrently. Immutable data is therefore well suited for concurrent programming.
- Redundant identical copies of immutable objects might occasionally get created, but this affects only the memory efficiency, not the correctness (assuming that your `equals` method does not use `==`). These days we will cheerfully trade a bit of extra memory for guaranteed good time and program correctness.
- The combined effect of two unfortunately interleaved threads **typically manifests as lost or missing data or values**. Such **concurrency bugs** resulting from **race conditions** can be extremely difficult to find and debug.
- Java Collections and **especially I/O streams** are **not thread-safe** (except for some particular implementations that explicitly guarantee this by their design), so that they have no built-in protection against race conditions and can thus behave strangely in concurrency.
- It is not possible to make some operation atomic simply by telling JVM not to make context switches during that operation, since the threads have no control over context switches. (Even if such a global ban on context switching were somehow possible, it would just eat the profits of parallel execution.)
- Since the JVM executes bytecode instructions instead of Java statements that are broken down into a series of bytecode instructions that can then be arbitrarily combined and rearranged by various compiler optimizations, a context switch can very well happen “inside” a statement! **You therefore cannot magically make some complex operation atomic by writing it as some clever one-liner!**
- To prevent race conditions, you need to identify all **critical sections**; blocks of code that should be executed by at most one thread at the same time so that while one thread is executing the critical section, no other thread may enter it.
- (Technically, we could allow two threads to enter the same critical section provided that they are operating on separate objects. However, the tiny efficiency gains available from such much more fine-grained locking are not usually worth the effort of complicating your code.)
- **Only you can identify what constitutes a critical section**, since they depend on your intended semantics of the program. No compiler can identify critical sections for you, even in principle.
- Even worse, the one and the same critical section is not necessarily one continuous code block inside the program, but can exist in multiple pieces in separate methods, even in separate classes!

9.5. *Mutual exclusion locks*

- Having identified a critical section, you should declare a **lock object** of some suitable subtype of the interface `java.util.concurrent.Lock` (usually `ReentrantLock`). This lock must be visible to all threads that will access the critical section.
- This lock object would be best named `mutex` or similar to instantly document its purpose.
- A lock can be **held by at most one thread at any time**. A thread can **capture the lock** by calling its method `lock`, and release it by calling `unlock`. If another thread (or “threat”, from an unin-

tentionally amazing misspelling in an old final exam) tries to capture a lock that is held by another thread, the method `lock` will **block** until the time that the lock becomes available.

- The JVM will simply context switch to execute some other thread, since it would be pointless to execute the thread that is waiting for a lock owned by someone else.
- In the beginning of each critical section, call `mutex.lock()`; and at the end of that critical section, call `mutex.unlock()`;
- If the execution of this critical section might crash and throw an exception, make sure that you release the lock in the `finally` block, to ensure that the lock does not remain forever in the locked state so that nobody can ever again get in.
- The mutex lock effectively acts as an invisible force field (or to use a more mundane analogy, an ordinary bathroom privacy lock) that prevents other threads from entering the critical section while the execution of some other thread is inside it. Once the lock gets released on the way out, the force field vanishes and the next thread can get in, that new thread again locking the critical section for the time being.
- If the critical section is split into multiple disjoint pieces in the source code, **each piece must use the same mutex lock**. Two separate critical sections that are independent of each other can use separate locks to avoid needlessly blocking each other.
- The Collection Framework classes are typically not synchronized against concurrent access. The static method `Collections.synchronizedCollection` (similarly for `List` and `Set`) decorates any underlying collection with a decorator object whose three dynamic set methods `add`, `contains` and `remove` are internally synchronized within this method.
- (Iteration cannot, and should not, be synchronized this way. Why this is so, is left as an exercise for the reader.)
- The package `java.util.concurrent` offers list and set implementations specifically designed to tolerate and even thrive under concurrency.
- A **reentrant lock** can be **held by the same thread multiple times**, which makes a difference if the critical section is **recursive**. The lock maintains an internal counter of how many times its current owner has captured it, so that its holder must then release it that many times for that lock to become available for other threads.
- (In computer programming, a system is said to be **reentrant** if multiple activations of that system can exist simultaneously. This forbids storing data used by that system in global or static variables, of which **there can always be exactly one**. As a rule, you should never use global variables anyway, except as `static final` named constants.)
- When designing synchronization, it is better to lock a little bit too much (so that possibly once in a blue moon some thread is forced to wait a while when it really wouldn't need to) than to not synchronize quite enough and by doing so break the correctness of the entire program!
- However, forcing too much synchronization on a program may cause it to **deadlock**.
- Nobody ever said that concurrent programming was going to be easy. **It is a whole another dimension** that extends orthogonally from the flat sequential programming. Fortunately, the high level tools from the next chapter will make it somewhat easier.

Lecture 10: Concurrency Controls

10.1. Condition variables

- Mutual exclusion is but one of the two types of synchronization that holds the reins of concurrent programs. The ability to pause a thread until some condition has been satisfied is just as important, since it is necessary for methods that **block** until they get the green light to proceed.
- For example, a method that needs some exclusive resources such as files has to wait until that resource becomes available, if the resource is currently being held by another thread or process. This wait must continue for however long it takes for the resource to become available.
- Any time you call a blocking method, you assume the risk of your thread getting stuck in carbonite for an unpredictably long time. The method documentation must make this clear.
- Some libraries can offer alternative **timeout** or **fail-fast versions** of those methods that throw an exception instead of blocking, so that the caller can decide how to proceed.
- Blocking methods, as a rule, throw `InterruptedException` in their signature (see below).
- **Blocking is infectious** in the sense that if your method calls any blocking method, then your method is also blocking, from the point of view of its callers.
- A **busy wait** (also called a **spin wait**) that uses an empty loop `while(!CONDITION) {}` is a **horrendously inefficient** and bone-headed way to implement blocking on a condition.
- **Condition variables** (in older literature, also called **monitors** when paired with a mutual exclusion lock, as the Java `Condition` objects always are) are a simpler and much better way to implement waiting so that a thread that is waiting for some condition does not waste any processor cycles in running around in circles like a dog chasing its own tail.
- Whenever you identify a condition from your program that some thread is waiting for, create an instance of `Condition` to represent that condition. This instance must be visible to both the waiting thread and to whoever else will make that condition to become true.
- Modify the busy wait loop to be `while(!CONDITION) { condition.await(); }`, so that the JVM will not waste any time executing any thread that is in a known waiting state, but can context switch away from it immediately.
- It is your responsibility to identify every single part of your program that could theoretically make that condition true for some awaiting thread. (Same as with recognizing critical sections, the compiler cannot help you in this, even in principle.) In all those places, you should call the method `signal` of that condition variable to wake up one waiting thread.

10.2. Crowding inside a condition variable

- **False alarms** that `signal` some waiting thread to wake up cause no problem in principle, since the condition of the while-loop that surrounds the `await` will merely cause that thread to go back to the waiting state again. No harm, no foul, *no problemo*, as my fellow kids like to say to sound hip and relevant to these times.
- (That reference is now over three decades old and must therefore be either hip or quaint, depending on the ebb and flow of the culture, perhaps even switching its existential state between semesters. Hasta la vista, baby!)

- A thread having to give up its turn at the head of the queue and then going back to wait at the end of the queue after a false alarm can have **fairness** implications; see the next section.
- If multiple threads are waiting inside the same condition variable, you should make sure that whichever of them gets woken up will also be able to actually get out to productively use its new-found freedom. This immediately entails a separate condition variable for each separate condition that your threads could be waiting for.
- If no thread is waiting for some condition inside that condition variable, **the unused signal does not get stored for the future awaits**, but simply vanishes as if it had never even occurred. This can cause surprising starvation and deadlock situations to emerge.
- Each condition variable is always built on top of an existing mutex `Lock`. The thread must own that lock before it is allowed to call `await`. Therefore, **a thread can only await for a condition inside the critical section that uses that particular mutex lock**.
- When a thread enters the waiting state, **it gives up that mutex lock** (but not any others that it might currently be holding), so that other threads can enter that same critical section, which may be necessary to make that `condition` true. There is no risk of corrupted data due to race conditions, since the waiting thread is not doing anything, and thus is not reading and writing any data concurrently with any other threads.
- To actually wake up from the `await`, the thread must first recapture that mutex lock before it is allowed to proceed. If some other thread currently holds that lock, the thread woken up has to wait in line with other threads for that lock to become available to it. This guarantees that at most one thread at a time can be active inside that critical section.

10.3. Fairness in concurrency

- As with many other terms in computer science and engineering, in the theory of concurrency the term “fair” has a rather different technical meaning than in our everyday language.
- In real life, we would not call it “fair” when a restaurant waiting line or supermarket checkout queue could at whim call some snazzy customers to the front of the line past all the boring normos. And yet in computing, such a queue could be perfectly “fair”.
- In concurrency, a system is **fair** if no thread can ever **starve** by having to wait forever and never receive the requested service. The actual waiting time is irrelevant in this binary theoretical sense: some thread having to wait a billion years (still $O(1)$) but no more still makes that system “fair”.
- **First-in-first-out (FIFO)** service by itself is obviously fair in this sense, assuming that the entity performing the service does not **deadlock**.
- In concurrency, a **deadlock** is created by a **cycle of waits** where each party is stuck waiting for the next one up the cycle to do something, typically resulting in too much mutex locking. Concurrent systems should be designed so that no such cycles can emerge, but this (like many other issues of concurrency) is an extremely nontrivial problem in general!
- Unlike the primitive object locks in the Java language that are not guaranteed to be fair, the objects of `Lock` (and other high level concurrency utilities in `java.util.concurrent`) can be made fair (in fact FIFO, since the Java concurrency API terminology is technically incorrect here) with a `boolean` constructor parameter.
- It is important to understand that **fairness does not propagate upwards**: building a concurrent composite from individually fair parts does not guarantee the fairness of the entire composite system. (This may have larger societal implications for other departments to discuss.)

- Classic exercises in concurrency such as the **dining philosophers** or the **producers and consumers** illustrate the difficulties in designing concurrent systems that make deadlock and starvation impossible by design, and still be efficient and live.

10.4. Asking another thread to terminate

- Every thread will terminate by itself when it returns from its **top level call**, the **main** method for the main thread and the **run** method for others. But what if some thread ought to terminate earlier than that?
- For example, consider the graphics thread of some real time action game that runs in a loop that draws the game screen 25 60 times a second. When the player quits the game, this thread should terminate along with all other threads of that program, so that the JVM process can terminate.
- Since calling the method **start** will start a thread once it has been created, it is probably not that difficult to guess which method will stop the thread. However, this method **stop** in **Thread** has long been officially **deprecated** and should never be used, since relying on this method is highly dangerous and error prone in any nontrivial program.
- Any resources held by the stopped thread are not released, its **finally** blocks are not executed, and anyway the thread is terminated in the middle of whatever operation it was currently at, possibly leaving the data in an inconsistent state.
- Instead of shooting the thread with the elephant gun of **stop**, we ask it nicely to stop at some convenient moment in the near future. The code executed by that thread is written to occasionally check the value of some **boolean** field (in all my code examples, this field is named **running**) shared by both the thread and the thread that might be asking it to stop.
- To stop a thread designed this way, just assign the value **running** to be **false** and leave the thread alone so that it will eventually clean up after itself as it terminates in proper order.
- Just one problem remains with this simple scheme: what if that thread happens to be stuck inside some condition variable, waiting for some condition that will never become true, since another thread that would have eventually made that condition true has already terminated?
- Whenever you set **running** to **false**, you should also call the special method **interrupt** to all threads that care about **running**. This method gives the thread a friendly “nudge” that is enough to **get it out from any wait** that the thread might currently be in. If the interrupted thread is not currently in any waiting state, the interrupt does nothing to the execution, except that it sets an internal **boolean** flag that will prevent it from entering a wait in the future.
- A good analogy to **interrupt** would be the way an incoming casino dealer gives the previous dealer a light shoulder tap to let her know that the shift change is due and she should finish up the current deal, but not start another deal after that one.
- Instead of returning normally from that wait, the method that was interrupted awake throws an **InterruptedException**. This particular exception subtype happens to be checked, which is why all blocking methods in Java will declare this exception in their **throws** clause.
- The checked nature of this exception subtype is also the reason why the **run** methods must catch and handle it, since the **run** method overridden in the subclass of **Runnable** cannot let out any checked exceptions.
- A thread can also check (using the method **isInterrupted**) if it has been interrupted, and could theoretically be programmed to react to an interrupt whichever way you want. (After all, all

syntactic things in programming really mean only whatever we say that they mean.) However, the most common and standard use for `interrupt` is to request termination.

10.5. Semaphores

- **Semaphores**, the great granddaddy of all concurrency control techniques, were historically the first high level concurrency control mechanism, invented by the late great **Edsger W. Dijkstra**.
- Semaphores are **universal** in the sense that all concurrency control could theoretically be done using nothing but semaphores. Standard exercise in concurrency textbooks is to implement other high level concurrency controls (e.g. `Condition`, `CountDownLatch`, `CyclicBarrier` or `BlockingQueue` in the Java terminology) using only semaphores.
- **A semaphore is an integer counter that can be initialized with any initial value**, positive or negative or zero. A semaphore is said to have a number of **permits** inside it. There does not exist any kind of `Permit` class, we merely conceptually think of this integer as “permits”.
- Two basic atomic operations are called **acquire** and **release**, so that **acquire** will block until one permit becomes available (that is, the counter becomes positive) and then decreases the counter, whereas **release** unconditionally increases the counter without blocking.
- A semaphore with initial value of one can work as a mutex lock, with the method **acquire** corresponding to **lock**, and the method **release** corresponding to **unlock**.
- A semaphore used as a mutex with initial value higher than one will permit that many threads to enter an **expensive section** simultaneously. This is a useful technique for **load control**.
- An important difference between semaphores and other controls such as locks and condition variables is that a **semaphore has no concept of owner**. At any time, **any thread can try to acquire or release a permit to the semaphore**. This flexibility is the source of the power of semaphores, but imposes some extra responsibility on the programmer.
- Surprisingly complex concurrency controls can be implemented using one or more shared semaphores with cleverly chosen initial values and logic to acquire and release them. (For those of you interested in this kind of stuff going potentially way beyond this course, check out the excellent [“The Little Book of Semaphores”](#) freely available online.)

10.6. Additional concurrency tidbits

- Before Java 5, locking of critical sections was done using the keyword **synchronized** that wrapped a block of code or an entire method to be synchronized with the particular object.
- Yes, Virginia, **every object whatsoever can act as a synchronization lock** in the Java virtual machine. This causes every object to use more heap memory than it really ought to.
- Even more strangely, every object in the JVM heap is not only a synchronization lock, but an entire thread queue that can act as a condition variable! The methods **wait** and **notify** are inherited to every Java object from the universal superclass `Object` for this purpose.
- The class `Thread` also contains some more or less useful **static** methods for thread control. For example, to make some thread wait for some precise number of milliseconds, have it call the method `Thread.sleep(ms)` with the number of milliseconds to sleep given as a parameter.
- Threads can be given a **priority**, a suggestion that the JVM should give the high-priority threads more execution time than it will give for the low-priority threads. However, these **priorities can**

only be used to optimize the speed of a concurrent system that is already correct for every possible execution interleaving, and they cannot make an incorrect program to be correct.

- Changing the thread priorities of a faulty concurrent system cannot help, since the JVM is not guaranteed to obey these priorities anyway. Even if it did, there is no guarantee that at some critical juncture with two waiting threads, the JVM would always first wake up the one with a higher priority, since that just might be the time for the lower-priority thread to get its turn to go first.
- Even the lowest priority threads must *sometimes* get their turn to execute. Since the JVM has no knowledge of the high level structure of the original program but mindlessly executes instructions one at a time, it has no way to recognize an important juncture.
- A single-threaded program is correct if it produces the correct answers for all possible inputs. **A concurrent program has to do the same for every possible interleaving of its execution threads.** This makes concurrent programs and their behaviour vastly harder to reason about and debug than ordinary single-threaded programs.
- Concurrency bugs can be very annoying in that **they are usually not reproducible**, since they tend to manifest themselves only for a small portion of all possible interleaved executions, which are generally not reproducible even within the exact same hardware platform.
- **Modifying a concurrent system to detect its bugs can change its timing and other behaviour** so that the bugs mysteriously vanish in that particular testing environment, and reappear when the detection is removed. This phenomenon, humorously analogous to the **Heisenberg uncertainty principle** in physics, is thus called a **Heisenbug**.
- Declaring a field to be `volatile` forbids the compiler from assuming that the field does not change unless it is explicitly assigned to. Such fields must always be read and written directly from memory at every use, and their values cannot be **cached** in the processor registers. Declaring a field to be `volatile` prevents various compiler optimizations such as **hoisting invariants** outside a loop, and whose correctness tacitly depends on the assumption that the fields cannot change unless explicitly assigned to.
- The package `java.util.concurrent.atomic` contains **atomic versions of boxed primitives** so that important non-atomic operations such as “increment and test” exist as internally atomic methods.

10.7. Concurrency and Swing

- Swing **runs in a separate event dispatch thread** that is not a daemon, and thus continues to exist even after the main thread and all other threads have terminated. Typically in a Swing program, the main thread merely runs the `main` method to set up the initial components before terminating, after which this Swing thread takes over.
- The event dispatch thread works as an infinite loop that processes the internal operations of AWT one at a time.
- **The event dispatch thread comes into existence when the first top level container becomes visible**, and terminates when no top level containers are visible. This makes the `setVisible` method call of `JFrame` way more important than it might initially seem, almost like an ignition key to the roaring Swing engine itself.
- Note that the Swing event listener methods get executed in this AWT event dispatch thread. If your event listener does something that takes a very long time, this will freeze the entire Swing until that method terminates.

- If your event listener really needs to do something that takes a long time, it should create a new background task for that purpose, and then return immediately to allow Swing to continue its operation smoothly.
- If you create additional threads that concurrently modify Swing components, these modifications could interfere with the ordinary Swing operations, since Swing components are not guaranteed to be thread safe.
- To give Swing a `Runnable` task to execute in its event dispatch thread, use the utility class `SwingUtilities` and its methods `invokeAndWait` and `invokeLater`. Since Swing is guaranteed not to be doing anything else while it is executing tasks submitted that way, such tasks can safely modify all Swing components to their heart's content.

Lecture 11: Computation Streams and Lambdas

11.1. The increasingly important laws of Moore and Amdahl

- The famous **Moore's law** states that computing speed (or more accurately in the physical sense, **frequency**) doubles every 18 months.
- From its inception in 1965, this prediction has held up astonishingly well. Outside the field of astronomy and especially in any fields of study concerning humans and their messy interactions, normally one would not observe numerical predictions that have unfolded in reality for over five decades as neatly as Moore's law, and never see such accurate predictions of how the great unknown future will unfold.
- However, this exponential growth cannot possibly continue forever, since there are both physical (speed of light, electronic noise, quantum uncertainty) and economic hard limits that keep pushing back our attempts to speed up execution.
- Doubling the processor speed does not necessarily cut the running time in half, unless the reading and writing bytes from and to memory, often the bottleneck of machine code execution in actual computers as opposed to our Platonic virtual machines where each operation takes the same abstract $O(1)$ time, is similarly accelerated.
- (Plus, once the computer has to interact with slow and smelly humans, machines will spend most of their time tapping virtual fingers waiting for us to finally make a decision.)
- Algorithms that access RAM bytes in the same memory page to minimize swapping and perform several calculations on the bytes they read from RAM into the processor registers before having to write the results out are said to have good [data locality](#). For this reason, looping through a large 2D array using nested loops is faster with rows in the outer loop and columns in the inner, compared to having columns in the outer loop and rows in the inner.
- A less painful way to speed up many computations is to **parallelize them over multiple processors and processor cores**. Unfortunately, computational problems vary greatly in how far they can theoretically be parallelized.
- [Amdahl's law](#) from the year 1969, not quite as famous and eerily prescient as Moore's law but still just as important even though people don't seem to realize it, points out that every computation consists of an **inherently sequential part** and **parallelizable part**. Using P idealized processors, the total running time will be $T = T_S + T_P / P$.
- At one end of this continuum lie the **inherently sequential problems** for which $T = T_S$. These problems consist of a strict series of **stages** so that the stage $n + 1$ cannot begin until the previous stage n has been completed.
- Analogy: building a skyscraper, with the pouring of concrete to build each floor being a separate stage. (Also the old chestnut of "Nine women can't team up to make a baby in one month.")
- The execution of the given arbitrary sequential program cannot in general be parallelized by algorithmic means, except in rare special cases. The lower level language some program is written in, the more difficult its parallelization becomes while guaranteeing the preservation of its higher level semantics.
- At the other end of this continuum reside the **embarrassingly parallel problems** for which we have $T = T_P / P$, at least in the idealized case. Such problems consist of a multitude of small tasks that are all independent of each other, allowing them all to be potentially evaluated in parallel. Typical examples would be all kinds of combinatorial **needle-in-haystack** searches and the **graphics rendering pipeline**.

- Graphics hardware speedup has obeyed its own accelerated version of Moore's law where the doubling happened every **six** months, which creates an eightfold increase over 18 months. Many other problems (such as **bitcoin mining** and **machine learning** with **neural networks**) have been found not to require the universal computation framework provided by the computer CPU, but can be pipelined and parallelized for **GPU hardware** that can then be piled in racks and warehouses to build a "poor man's supercomputer".
- Most actual problems and the algorithms to solve them fall somewhere between these two extremes in that they can be parallelized to some extent. An important part of the **science of algorithms** is developing new ways to parallelize important computations.
- Parallelization works out ideally when the parallel processor cores do not need to communicate with each other, synchronize the access to any shared data, or wait for some previous tasks to finish and give them the results that they need to proceed. Synchronization of shared memory, especially if done in a crude "stop the whole world while I do my little thing" fashion, can slow down parallel computations significantly and even eliminate the time savings theoretically achieved by parallelization.
- But so it is in many other walks of life. For example, in software engineering, the famous **Brooks' Law**: *"Adding more manpower to a late project only makes it later."*
- Another dimension to parallelize computations that consist of a large number of tasks, each task a strictly sequential series of stages (e.g. microprocessor machine code instructions, or computer graphics rendering pipeline) is to **pipeline** them so that each stage feeds results to the next one, and can start its next task immediately after passing its results down the line.
- This will not speed up the processing of one sequential task, but will greatly speed up the processing of, say, a million tasks of that type.
- Analogy: a factory assembly line, or a cafeteria line with separate stations for various foods, drinks and payment. Or the way Subway and similar fast food joints pipeline and parallelize their assembly of sandwiches with such small parallelization optimizations.

11.2. Fork-Join framework

- Java 5 improved the state of concurrent programming in Java by bringing in all sorts of high level concurrency controls such as **Lock**, **Condition** and **ExecutorService**. In this course, we will not even cover the language level concurrency techniques which are idiosyncratic to Java anyway, all concurrency controls should be as high level as the rest of the program.
- Ideally, the compiler and the execution environment should do the parallelization for us. Unfortunately, this requires language to be sufficiently high level to allow the necessary inferences about its semantics, otherwise parallelization is not guaranteed to maintain these semantics.
- Java 6 or 7 were not very big updates, but at least Java 7 introduced a higher level concurrency framework of a **Fork-Join framework** that hides away all the details of individual threads.
- A **ForkJoinPool** is a generalization of **ExecutorService** in which you can submit computations. These tasks have to be subtypes of **ForkJoinTask<T>**, or more usually its utility subclasses such as **RecursiveTask** where some methods are given a handy and meaningful default implementation.
- Just like with an **ExecutorService**, you only need one instance of a **ForkJoinPool** in your entire program.
- The **ForkJoinPool** keeps track of all thread maintenance and **load balancing of submitted tasks** to different threads, using as much computer resources to perform these tasks as the under-

lying machine provides. A Java program written today, if by some miracle executed by somebody in the far future, will fully use the power of future computers that will hopefully have hundreds or even thousands of parallel cores.

- (To further mangle a famous but apocryphal quote, “Six hundred and forty processor cores should be enough for anybody!”)
- Any `ForkJoinTask` can **fork** new tasks to existence to the pool, or **join** another task waiting for it to terminate. In the theory of concurrency, every time you see the word “join”, just mentally substitute “wait for termination”.
- It is your job to write your `ForkJoinTask` methods so that you try to happily **fork** as much as you possibly can, since the more you **fork**, the more parallelism you will get. In fact, every time you write a **fork**, you should mentally go “Yay! More parallelism for me!”
- Sometimes some particular task cannot proceed until some previous task has finished, in which case you have to write a **join**. Every time you **join**, you decrease parallelism, and should perhaps mentally go “Boo! Better check again whether that join is really necessary!”
- When solving an inherently sequential problem, you would never get to write a **fork**, so using a `ForkJoinPool` would be pointless. With an embarrassingly parallel problem, you never have to write a **join**, except possibly in the very end just before you return to say that everything is complete.
- A **semaphore** initialized to value $1-n$ is an excellent technique to wait for n separate tasks to finish. The **acquire** operation for waiting will block until precisely n calls to **release** have taken place, regardless of the order in which these tasks complete their work.
- For problems that lie between these two endpoints of the continuum, you will write some forks followed by some joins. For example, the **parallel merge sort** can fork the subtasks of sorting the two halves of the array, but merging these sorted halves cannot begin until both halves have been sorted. (**Parallel quicksort** can similarly fork the subtasks of sorting the two partitions, with only one wait at the top level for every recursive part to finish.)
- Since the pool needs to use both time and memory for task scheduling and maintenance, there necessarily exists some cutoff point so that **when a task is small enough, breaking it down to smaller tasks to be executed in parallel would cost more than the savings of parallelization**.
- Even when using a fork-join framework, with more complicated tasks sometimes you still need more fine-grained synchronization using locks, conditions and semaphores to protect the critical sections of your code, especially if your forked tasks access and modify shared objects.
- The Fork-Join framework also has no magical superpowers to divine the critical sections of code and data from your task submissions. Critical sections still depend on your intended higher level semantics, and are not inherent to the code itself.

11.3. Lambdas as syntactic sugar for anonymous classes

- **Statically compiled languages** tend to make a hard distinction between code and data, so that code is set in stone at compilation, whereas data can be created and modified dynamically. On the other hand, **dynamic and functional languages** use **homoiconic** representations for both so that **code is data and data is also code**, depending on which way you currently look at it.
- In the theory of programming languages, a **first class object** is (roughly) something that can be assigned to a variable or passed to and from a function or method. For example, strings, integers and birds are first class objects in Java, but **methods and blocks of code are not**.

- (The term “object” is used here in an older sense that predates the term “object-oriented programming”. To dispel this confusion, some material talks about “**first-class citizens of the language**” instead.)
- Strategy objects are often defined on the fly using an **anonymous class** that extends some **functional interface** that has only one method, the **functor object** thus wrapping that method inside a first class object that can be given to another method.
- **To pass a method to another method, pass a strategy object to the method.** This is ultimately just as good as passing the method as a real object, as if we were working in some higher level functional language such as Python 3.
- Defined in the package `java.util.function`, there are a whole bunch of such functional interfaces such as `Function<T,R>`, `Predicate<T>`, `BinaryOperator<T>`, ...
- When you write the inner class that contains the method you want to pass as argument, you need to write a whole bunch of boilerplate to keep the compiler type checking happy. However, while tiring your fingers with such tedious typing chores more suited for our mechanized servants, you are not really telling the compiler anything that it doesn’t already know!
- Using the **lambda** operator `->` introduced in Java 8, you can essentially ask the compiler to fill in all that tedious boilerplate for you, so that you only need to type in the actual method body.
- For example, to create a strategy that extends `Predicate<String>` and checks whether its parameter string is at least five characters long, write `(x->x.length(>4))`. **The compiler will fill in everything else**, even the `return` keyword for such a one-liner!
- If the body has multiple statements, the curly braces are necessary. If the method takes more than one parameter, their names must be listed in parentheses, e.g. `(x,y)->x+y` for the anonymous function that returns the sum of the values of its two parameters.
- What is the particular name of the method in `Predicate` or `BinaryOperator` that we were supposed to override in those examples? Answer: Who cares! The compiler knows all that stuff anyway, so let that guy look it up for us so that we can allocate our precious and limited brain cycles on more pressing matters.
- The term “lambda” refers to the historical syntax of defining an **anonymous function** that is itself data, with that greek letter used to denote the parameter variable. This dates all the way back to the thirties with **lambda calculus** of Alonzo Church.

11.4. Functional programming

- Ever since the early days of programming, programming languages have traditionally split into two major camps. One started from the machine and aimed high for mathematics, the other started from mathematics and aimed downwards for the machine.
- Most sources tend to be dualistic and consider these **imperative** and **functional** programming paradigms a black and white binary, even though their reality is actually a continuum that transitions smoothly from one extreme to the other.
- The further towards the functional style you ascend in this continuum, the more you describe the problem to be solved, rather than describing the computational steps to solve that problem.
- In the imperative end of this continuum, each operation describes exactly what the physical computer actually does. The purest imperative programming is done in **assembly language** whose operations correspond directly to the operations of the processor. Most **assemblers** allow macros, eventually rising to the level of the **C programming language**, the lowest level high-level language still in common use.

- As languages ascend to talk about higher level abstractions of computation, they acquire more functional programming features. Different languages can therefore be roughly placed on this continuum based on how “high-level” they are.
- **We should not confuse “high-level” with high quality:** there are good and bad imperative languages, just as there are good and bad functional languages. Every tool has its purpose, and **no tool can possibly be perfect for every purpose.**
- (Even for working “close to iron”, [there are now Go and Rust](#) to replace the old workhorse of C.)
- For any language and other programming tool, the proof of the pudding always ultimately lies in **whether real programmers can productively use it to solve their actual problems.**
- In fact, today we can get the best of both worlds by mixing languages of different levels to do what each one does best. For example, many games have their computation intensive and low level engine parts written in low level languages close to the computer hardware, whereas the higher level logic is written with Python or some similar high level language. This separation of concerns makes the game itself easier to script, extend and modify.
- Especially dynamic languages that can treat their own code as data and vice versa, instead of the code being set in stone at compilation, make for easy scripting that would allow even entirely new sets of rules to be added to the game dynamically! This is the approach and philosophy behind **data-driven programming**.
- Purely functional programming languages have been created as academic research projects, but in their strictness of immutable data and other functional programming aspects, they are too unwieldy and [just too damn weird](#) for real programmers to solve real problems.
- As another pithy expression put it, nobody has yet written an operating system kernel using a functional programming language.
- (Even writing silly arcade game of **Pac-Man** or **Space Invaders** that is supposed to run in real time, something that was achieved by a bunch of stoned out hippies back in the 1970's by soldering connections directly onto the motherboard, becomes [non-trivial when using a pure functional programming language](#) even with the best of our current tools. Something has therefore clearly gone very wrong with this entire picture.)
- We should also never forget that **no programming language can ever escape the physical reality of the underlying computer and its actual mindless low level operations**, no matter how high up in outer space its concepts arise. In fact, losing sight of this underlying reality means losing sight of the true cost of your computations, akin to the [Poor Shlemiel tale](#) for writing nested loops when one would have sufficed, or in the field of software engineering, becoming an [architecture astronaut](#) whose designs essentially become highly convoluted ways of stating that “There exists some stuff, and then something happens.”
- Old time functional programming languages were also unpopular due to their strange syntax (or in the case of Lisp, [having essentially no syntax](#)), but modern languages enable the functional programming ideas in a syntax that is more familiar to the imperative programming masses.

11.5. Meeting in the middle of the two paradigms

- It remains to be determined where exactly the “sweet spot” between purely imperative and purely functional programming lies. However, as the modern languages such as Python, Ruby, Rust, Kotlin etc. from this century demonstrate, this spot is clearly higher up the continuum than the spot occupied by the major nineties languages such as C++ or Java.
- (Back in the day, even those two languages were considered “high level”, except by a bunch of [bearded Lisp and Smalltalk](#) weirdos who just [smirked and kept doing their own thing](#). See the

section “The Blub Paradox” in the linked article from 2003 for a better description of the above notes, or the famous [Greenspun's Tenth Rule](#) of how programmers whose mindset is stuck to low level languages inadvertently end up crudely simulating higher level languages as soon as their project size grows past a certain threshold. See also “[Frequently Rediscovered Technologies](#)”.)

- To cross the line from a static to a **dynamic** language, the language needs to have a built-in **eval**, the ability to **evaluate any dynamically generated string as an expression of that language**. Having the eval furthermore allows the language to have a **REPL, read-eval-print-loop** for interactive use of that language. (Java 9 took the huge step of introducing both of these.)
- Functional programming **makes no distinction between code and data**, but either one can always be treated as the other. New code can be composed dynamically by combining existing functions in various ways. (Java 8 lambdas took a big step in this direction, even though the byte-code is still set in stone at compile time.)
- Furthermore, same as in mathematics, **functional programming treats all variables as final and all data as immutable**. This has multiple upsides, as we have seen many times already, but also the huge downside already encountered in the discussion of the autoboxing mechanism.
- Combined with pure functions with no **side effects** (other than possibly input and output to the outside world), immutability guarantees **referential transparency** in that the same expression will always evaluate to the same result, and any subexpression can always be replaced with a result-equivalent subexpression without affecting the value of the entire expression.
- Reading through an imperative program you can reason which computational steps can be taken, and in which order these steps happen. Functional programming languages are more free to reorder the evaluation, and more importantly for modern processors, **parallelize** it using the knowledge of the high-level structure to guarantee the correctness of this parallelization.
- Functional programming allows **partial evaluation** of some function to create another, more specialized function, which would be impossible in imperative programming languages with their strictly **eager evaluation**.

11.6. Computation streams

- Java 8 took a huge step towards the functional programming paradigm by introducing **computation streams** (not to be confused with I/O streams that we saw earlier, although conversions are trivial) using the **map-reduce framework**.
- Computations expressed in pure functional fashion with map-reduce streams **parallelize automatically** over multiple processors. Most famously, Google’s search algorithms are internally written using this kind of framework, so that each individual query that would take well over a minute to execute sequentially in a single machine is broken down and parallelized to execute in a split second over thousands of servers in the massive server farm.
- A computation stream begins with some **supplier** of elements, and continues with various operations (most importantly **map** and **filter**) that perform computations on those elements. These operations can now be concisely written by using lambdas to define their strategy objects.
- A stream can be **sequential** or **parallel**, the latter being more efficient when the order of elements is irrelevant and the individual operations performed on them are **stateless**. A sequential stream guarantees to maintain the order of the elements in the stream and can therefore use **stateful** operations that depend on this order.
- An operation is **stateless** if its **result depends only on the current element coming from the stream**, but not the past history of previously processed elements. A stateless operation does not

need to remember the past in its data fields, and can therefore be defined with a lambda, and computed in parallel for each element separately.

- For example, primality testing of integers is stateless, whereas maximum, average, sum etc. have to remember something of the past history to calculate the answers.
- A stream can be infinite, but is then cut down to a finite stream using a `limit` stage.
- In the other end of the stream, there is some kind of **consumer** that collects the elements and generates some kind of end result out of them, often using `reduce` to turn a stream of elements into a single element as the final answer.
- Many classes in the Java standard library (the entire Collection Framework, `File`, `Random`, ...) have been retrofitted to be able to stream their contents on demand. Iterators can also be converted back and forth to streams, to combine the power of both approaches.

11.7. Lazy evaluation

- **Lazy evaluation** is an important concept in **functional programming languages** where it can reside more comfortably and can even be baked into the very language semantics.
- Imperative programming languages are by their nature **eager**. For example, all arguments of a method call are fully evaluated before the method call itself is executed. However, some lazy evaluation ideas can also be applied when programming in the classic imperative style.
- Lazy evaluation simply means that **no computation is performed until it is certain that its value will definitely affect the end result**. This can potentially speed up computations by eliminating redundancy.
- As a funny (but most likely apocryphal) tale of the definite pinnacle of lazy evaluation, the author recalls once hearing about a professor who, instead of grading the final exams, simply failed everybody without even looking at the exams, and then actually graded the exam only for the small handful of students who came to his office to complain about their grades.
- I am not a lawyer nor do I play one on TV, but the difference between the civil law and common law legal systems feels analogous to the difference between lazy and eager evaluation.
- (In real world systems, different incentives and costs for different participants can distort the system in similar ways to produce arbitrarily perverse outcomes, as is illustrated by another, hopefully also apocryphal, anecdote of the commander of an aircraft carrier group who was trying to sleep in his cabin but was disturbed by the sunlight coming in through the porthole, so instead of just getting up to close the curtain, stayed in his bunk and phoned in the order for the entire carrier group to change its bearing.)
- Even in imperative languages, the **short circuiting boolean operators** and **if-else** are an example of lazy evaluation.
- In computer science, the word “lazy” does not have the same negative connotations that it has in everyday language. In fact, an algorithm that is “lazy” in our everyday sense, that is, **does precisely that what is actually needed to produce the correct answer and not even one thing more**, would rather be called **elegant** or **optimal**! Somebody who could act like that in real life would surely be revered as a great zen master.
- As a real world analogy of lazy algorithmic thinking, imagine a raffle where the participants have to send in a postcard that must also contain the correct answer to the skill testing question. After the deadline, the winning card will be randomly picked from inside a spinning barrel.
- The eager evaluation version would check each postcard the moment that it arrives, and put the card in the barrel only if it had the correct answer. This guarantees that whichever card gets ran-

domly picked will be legal, but also wastes time looking at all the cards that did not win and whose contents therefore cannot affect the end result.

- The lazy evaluation version will dump all postcards straight in the barrel without even looking at them. When the random winner is picked, only that card is checked for having the correct answer. Keep picking another random winner until the correct answer is found.
- Java **computation streams guarantee lazy evaluation** in that even though a stream is defined forward from supplier to consumer, its **evaluation is executed backwards so that each stage requests the previous state for an element only when it really needs an element**.
- The computation stream does nothing until some consumer is attached to its end to start requesting elements. Before that, the stream is like lightning caught in a bottle, an object that can be assigned to variables, passed as an argument to a method, and returned as a result.
- Even if the producer or any intermediate stage produces a logically infinite stream, only a finite number of elements of that stream will ever actually get generated, precisely as many as the following stages actually need, no more and no less.
- If a field in an object is expensive to initialize but is needed only in rare situations, it can be **lazily initialized** to `null` (or some other convenient placeholder value, typically 0 or -1) and initialized the first time when its value is actually needed. For example, the expensive `hashCode` method in `String` is initialized lazily this way, and then **cached** for future use, knowing that **the hash code of an immutable object cannot possibly change**.
- In general, lazy evaluation can also make the computation have unpredictable pauses and use more memory in situations where the accumulated “debt” of delayed operations comes down for us to pay all at once. This makes inherently lazy languages unsuited for **real-time applications** whose response time must be predictable within hard physical limits.
- Lazy initialization can be vulnerable to **race conditions** in situations in which two threads try to use the same lazily initialized field simultaneously.

Bonus Lecture 12: Integers and Floats in Java

This entire topic is no longer part of this course, but there is no harm in including these notes here for any student who has come this far and still wishes to keep going. Properly dealing with fixed-size integers and floating point numbers while understanding their limitations is an important and essential skill to avoid their silent pitfalls in all programming languages, not only in Java. Especially all the Pythonistas, spoiled with having the luxury of unbounded integers in the core language and the `Fraction` type for exact representations of rational numbers in the standard library, should take heed that not every language is a similar utopia where numbers actually work the way they are supposed to work. Besides, many an unwary Pythonista has still fallen by the wayside with unintentional floating point calculations...

12.1. Positional number systems

- Two fundamental principles: (1) An object is not the same thing as its name or representation, and (2) There is nothing magical about the number 10, except the universal agreement once upon a time made by the humanity of using base ten to represent integers, basically because we usually have ten **digits** in our hands.
- We represent integers in a **base 10 positional number system** so that each digit gives the magnitude of the power of 10 corresponding to that position. It is not enough just to see the individual digit, you must also see its position to know how much it is.
- Contrast this to **Roman numerals**, where X is always ten and M is always one thousand no matter where it is positioned. (A slight exception to this rule is that if a smaller unit comes before a larger one, its value is subtracted instead of added. For example, XI is 11, but IX is 9. However, canonically a smaller symbol can only be subtracted from the next larger one, so that 99 cannot be written out concisely as IC, but must be given in the for XCIX.)
- Every positive integer has exactly one unique representation in the positional number system, assuming that the leading zeros are ignored. (Proof by induction. No, newbie, this claim is not “trivially true”. Hand-waved proofs such as “Every number has a unique representation in binary” are just begging the question.)
- By convention, the digits are written out left to right starting from **the most significant digit**. Many algorithms that operate on numbers would be easier to write as one-pass loops if the least significant digit came first, but convincing the entire humanity of the advantages of the reversed approach seems like a futile effort.
- Since there is nothing magical about number 10, we could use any other integer as base, and pretty much **everything would work exactly the same way**, including all numerical algorithms on sequences of digits that you already know (addition, multiplication, long division...)
- When using base n , the possible digits are always $0, \dots, n - 1$.
- **Computers use internally binary base 2** for easier representation in electronics hardware.
- Binary numbers and the familiar base 10 integers are still the same integers as mathematical objects, just represented in a different way. The integer objects exist in the timeless platonic space of mathematical truths, and we humans use different ways to talk about these integers, the same way as different natural languages use different words to talk about the same things.
- **Mixed radix systems**, such as our familiar 24/60/60 representation for time, or the **imperial system of measures** for weights and lengths, use a different base for different positions. As long as the arithmetic is done properly with respect to the base of the current position, generalizing to a mixed radix system doesn’t really change anything either.

- Bases 12 and 60 are actually more flexible than the base 10, since they can be easily divided into more useful parts than 10. Dividing things by either two or three is no problem in base 12.
- In principle, even negative numbers, fractional numbers, complex numbers or [Fibonacci numbers](#) could be used as pretty weird bases, although mostly only for **recreational mathematics** and similar amusement purposes for those of you with such inclinations. (See the works of the late great [Martin Gardner](#).)

12.2. Bits and bytes

- The computer memory consists of a very large number of bytes, each byte comprising 8 bits. Everything that exists inside a **von Neumann architecture**, both code and data, must be somehow stored as bytes, since bytes are the only thing that physically exist in such an architecture, as far as we can reach down from our programming languages.
- An individual byte can store **an unsigned integer** in the small range from 0 to +255.
- **Bytes don't know what they represent**, or whether they have been grouped to represent larger structures. All semantics for the bytes are always imposed from the outside by whoever is currently using them. The exact same byte and its value in the computer memory can mean different things depending on whether that type is used as a machine code instruction, a part of a four-byte `int`, a part of an eight-byte `double`, or some larger data structure.
- We saw this same principle earlier with byte streams that assume or require nothing about the semantics of the raw bytes that they transport.
- To represent integers too big to fit into a single byte, we group together 2, 4 or 8 bytes to gain a wider range.
- The smaller the base, the longer the representation of a particular number. To make binary numbers easier to read for humans, **hexadecimal base 16** is used to pack a **nybble** (a half-byte of four bits) into one **hexadecimal digit**. Letters A to F are conventionally used for the additional digits 10 to 15 for which do not have a numerical symbol for.
- Since the bases 2 and 10 do not have the same prime factors, changing one digit of the number inside one representation can change multiple digits (in the worst case, even all of them) in the representation in the other base. However, since $16 = 2^4$ exactly, each hexadecimal digit corresponds precisely to one nybble, and everything comes together perfectly.
- In Java and almost all other programming languages, hexadecimal integer literals are conventionally denoted with the prefix `0x`. Java 7 introduced the option of writing integers as **binary literals** with the prefix `0b`.

12.3. Signed integers and integer arithmetic

- To allow the representation of both positive and negative numbers, **signed integers** use the highest order bit to store the sign of that number, 1 being negative and 0 being nonnegative.
- Again, you cannot simply look at a byte or a group of bytes and determine whether it is signed or unsigned. Every byte means whichever way you decide that it means.
- To avoid having both positive zero and negative zero, which would make no sense for integers, the **two's complement signed representation** encodes $-n$ by **negating the bits of the unsigned representation of n , and adding one**.
- Since zero is the odd man out in the middle, there is one more value to the negative direction than to the positive direction. For example, signed `byte` has the range -128, ..., +127.

- **Java language standard requires** that `byte`, `short`, `int` and `long` absolutely positively must be stored as **two's complement signed integers** using 1, 2, 4 and 8 bytes, respectively. This guarantees that **integer arithmetic in Java will always produce exactly the same results in every platform that runs a standard compliant JVM installation**.
- Java does not have unsigned integer primitives in the core language. You know, “for simplicity”, just like the language still does not allow **operator overloading**. (Because it is much “simpler” for everyone to read and write `a.add(b.mul(c))` than it is to read and write `a + b * c` like we would write that everywhere else.)
- **Integer arithmetic is always exact**, so that `2 + 2` is exactly 4, not one iota more or less. However, **integer arithmetic can overflow silently**, with no exception thrown and no way to detect later that an overflow took place.
- The class `Math` offers **checked integer arithmetic methods** that throw an exception at overflow. (As an aside, if you had to write these methods yourself, how would you solve the chicken and egg problem of detecting whether adding two `int` values will overflow?)
- The utility class [`java.math.BigInteger`](#) represents signed integers whose size is limited only by the available heap memory, so such a number could have millions of digits. Arithmetic operations are performed with specialized algorithms (such as [`Karatsuba algorithm`](#) for integer multiplication) that are blazingly fast even for humongous numbers.

12.4. Bitwise arithmetic

- Bitwise arithmetic operators are analogous to the logical operators for truth values, but performed separately in parallel for each bit of the two operands. Their importance is that they allow us to read and write individual bits of a number without affecting the rest of the bits in that number.
- Important use in **bit vectors** and related data structures that pack 32 bits into a single `int` value, using every part of every byte to store useful information. In Java, `EnumSet` and `boolean[]` are internally implemented as bit vectors.
- **Bitwise or**, denoted by `x | y`, produces a number where the *i*:th bit is on if it was on in **at least one** of the numbers `x` and `y`. (Just like for the logical or, this is an **inclusive or**.)
- **Bitwise and**, denoted by `x & y`, produces a number where the *i*:th bit is on if it was on in **both** of the numbers `x` and `y`.
- **Bitwise negation**, denoted by `~x`, produces a number where each bit is the **opposite from** what it was originally.
- **Bit shift operators** `<<` and `>>` move the bits left and right the given number of steps. For the right shift, `>>` maintains the highest order **sign bit** whereas `>>>` always brings in a zero. The choice between these two operators depends on whether the bit patterns that you apply these operators on are semantically intended to represent signed integers.
- To turn on the *k*:th bit of `x`, use `x = x | (1 << k);`
- To turn off the *k*:th bit of `x`, use `x = x & ~(1 << k);`
- To read the value of the *k*:th bit of `x`, use the expression `x & (1 << k) != 0`
- The previous operations generalize from one bit position to any group of bit positions using a **mask** that is combined from the individual masks.
- **Bitwise xor**, denoted by `x ^ y`, produces a number where the *i*:th bit is on if it was on in **exactly one** of the numbers `x` and `y`. Note that the caret does not denote exponentiation in either Java

or Python languages. Java does not have exponentiation operator in the language at all, whereas Python exponentiation operator is the double asterisk `**`.

- Since the **bitwise xor does not erase information and is its own inverse**, it has nice applications such as **one-time pad cryptography**, quite beautiful in their elegant simplicity.
- Whenever you have a for-loop counting from 0 up to $2^k - 1$ for some positive integer k so that the variable `i` contains the current number, inside the body of the loop you can visit the value of $i \oplus (i \gg 1)$ instead of `i`. This conversion will still visit each number from 0 to $2^k - 1$ exactly once, but will do so in a more exciting [Gray code](#) order where exactly one bit changes its value between two consecutive numbers and the other bits stay the same.
- **Munching squares** is a classic display hack that can be implemented various ways, but the general idea is that the colour of the pixel (x, y) is some function of $x \oplus y$. the simplest of which is determining a threshold value t , perhaps as an animation over time, and colouring the pixel (x, y) in that animation frame if and only if $x \oplus y \geq t$. Variations of this technique can be combined and transformed to produce outcomes steeped in **glitch art aesthetic**.

12.5. Floating point

- In everyday life, we represent decimal numbers (in the sense of having a fractional part) using **fixed point decimal**, where the decimal point denotes the start of the negative powers of the base.
- For example, the two numbers 1.2345 and 1234.5 have the exact same **mantissa** of digits, but the latter is a thousand times larger than the former, as we can immediately see from the position of the fixed decimal point.
- Fixed point notation has two important weaknesses that make it unsuitable for representing decimal numbers in computations. First, it gives us **uniform precision** for the entire representable range, which is kinda silly, since the further away we move from zero, the less precision we need. Second, fixed point notation allows for a **relatively small maximum** value for the fixed length representation.
- In **scientific notation**, each number is rather expressed as the product of sign, mantissa and the power of the base. To make this representation unique for every representable value, the sign must be either +1 or -1, the mantissa has to be at least one but strictly less than the base 10, and the exponent must be an integer.
- The exponent tells us the **order of magnitude of that number**. The mantissa then pinpoints that number linearly within that order of magnitude, as if these representable numbers were lines of a uniform ruler.
- The previous example numbers 1.2345 and 1234.5 both have the same mantissa 1.2345, but the first number has the exponent of 0, whereas the second number has the exponent of 3.
- In Java and almost all other programming languages, scientific notation can be used to write constant literals in the form `1.234e-7`, the letter `e` standing for “times to the power of ten”.
- Again, the number 10 has no magical properties that make anything work. **Floating point is simply scientific notation using base 2 instead of base 10. Nothing else changes.**

12.6. IEEE 754 floating point numbers

- **IEEE 754 standard** defines the **single precision** encoding where each floating point number is stored 32 bits by using 1 bit for the sign, 8 bits for exponent **biased** by +127 to guarantee that it is nonnegative, and 23 bits for the bits of the fractional part of the mantissa over the one.
- A handy trick to perform this conversion by hand for number x is to break it down into a sum of powers of two. Divide this sum by its highest term to produce the mantissa, and then multiply this highest term back to serve as the exponent. This way you can read the bits of the mantissa directly from the powers of two that appear in your breakdown of the mantissa.
- Instead of doing this by hand like some peasants, use sites such as [IEEE 754 Floating Point Converter](#) to play around with different numbers and their floating point representations.
- Similar **double precision** (hence, the type `double`) encoding for 64 bits uses 1 bit for sign, 12 bits for exponent and 51 bits for the mantissa. (There exists even 128-bit **quadruple precision** for special applications.)
- The standard also defines special bit patterns for +0, -0, +Inf, -Inf and NaN.
- When the exponent has its smallest possible value with the number being as close to zero as we can get, **subnormal** numbers have the mantissa values placed **logarithmically** along the real line interval, instead of the marks of a uniform ruler as is done for higher exponents. This makes arithmetic of extremely small positive numbers more useful and accurate.
- The `double` type can represent all positive and negative integers up to 2^{53} , and furthermore, all integers that are a product of some such number and some power of two up to 10^{23} .
- The Java language standard requires that a `float` has to be stored **in at least 32 bits**, and a `double` has to be stored **in at least 64 bits**. However, this representation **is not required to be the IEEE 754 standard**, but other encodings can be used in special platforms. Floating point arithmetic in Java is therefore not guaranteed to produce identical results on different platforms, unlike integer arithmetic.
- Definitely the least used keyword in the Java language, `strictfp` forces the JVM to internally use the IEEE 754 standard for handling `float` and `double` values for a field, method or class. If this standard does not happen to be supported in some particular piece of iron, the JVM has to simulate the results on software, being as slow as if this were the 1980's again.
- Terminology sidebar: **simulation** is done in software, whereas **emulation** is done on hardware. (The obvious philosophical question for anybody who has ever gotten high and rewatched *The Matrix*, the movie now older than many of the students taking this course.)
- `StrictMath` is otherwise the same as `Math`, but its methods are `strictfp`.

12.7. Imprecision of floating point

- There exists an infinity of integers to both positive and negative directions, but we use only a small part of this infinity centered around the zero. Since the integers in that small part are “all there” so that nobody will ever discover a new integer between 7 and 8, binary representation allows us to represent all of them with arithmetic that is exact, unless there is an overflow.
- (As a lighthearted aside, can you explain why you are so sure that nobody will ever discover a new integer between 7 and 8 without resorting to any kind of circular reasoning, begging the question, or proof by intimidation by energetic hand waving?)
- Decimal numbers are more difficult since there exists an **infinite continuum of numbers** between any two decimal numbers. So even if we restrict our attention to some small part of the

infinite real line, there still exists the same infinite continuum of different numbers in there. The real number line is **holographic** in that the small part is the same as the entire thing, if you want to get all fancy about it.

- Fortunately, the floating point encoding is perfectly symmetric around the zero so that the number x is representable in floating point if and only if its negation $-x$ is representable. (This is not true for two's complement signed integers for which the negative side has one more value than the positive side, due to the zero being the odd man out in the middle.)
- No matter what the encoding, at most 2^{32} different numbers can be represented in 32 bits. There is no going past this fundamental limitation imposed on us by basic combinatorics. So, which real numbers *should* we represent in our encoding to make the system maximally useful?
- A rather serious problem with floating point is that even many seemingly very simple numbers don't have an exact representation in the base two floating point. (Even zero, the most important number of them all, had to be represented in a special way.)
- In base b , the fraction $1/n$ has a terminating representation if and only if all the prime factors of n are also prime factors of b . For example, when b is the familiar base 10, we can represent $1/10$ exactly as 0.1, but we cannot exactly represent $1/3$, which would be 0.333...
- In base two, we can't exactly represent **any** of the seemingly simple numbers 0.1, 0.2 or 0.3. All these would become an infinite repeating series of negative powers of two.
- **Every program whose logic depends on equality or order comparisons of computed floating point values is broken by design**, and cannot possibly be fixed without a complete redesign of its entire logic.
- The only acceptable ordering comparison is comparison to zero. Any other comparison of floating point appearing in your code means you should use either `BigDecimal` or some kind of fraction type to perform your computations.
- **Do not use floating point to represent actual currency** or other decimal values that have to be exact in our everyday base ten. Always use `BigDecimal` to represent such values.
- A `BigDecimal` object consists of a `BigInteger` mantissa and an `int` exponent given in base 10, and all arithmetic is performed with this representation. This base can represent familiar numbers such as 0.1 exactly so that we can add ten cents to some dollar amount.
- Unlike in the primitive floating point numbers, this mantissa is not necessarily normalized, to account for differences in **precision** so that 0.3 and 0.30000 are different numbers that produce equal but different precision results in arithmetic operations.

12.8. Floating point arithmetic

- Even when x and y are exactly representable in floating point, the results of simple arithmetic operations such as $x + y$, $x - y$, or $x * y$ might not be exactly representable.
- Whenever floating point arithmetic would produce a result that cannot be exactly represented, the operations produce the representable number closest to the true result. If two results are equally close this way, the choice between these two depends on the current **rounding mode**.
- If some long iterated computation is **chaotic**, even a small difference in some intermediate temporary result can blow up to become a huge difference in the end result.
- The dark art of organizing your computational steps to minimize the total effect of these rounding errors is known as [numerical analysis](#).
- Contrary to our intuition, floating point multiplication is easier and more accurate than floating point addition, especially when its operands are of vastly different orders of magnitude!

- Floating point hardware can internally use more bits of precision to guarantee the accuracy of the last bit of the mantissa of the multiplication result, even if the resulting mantissa is then truncated to 23 or 51 bits.
- Especially if x is several orders of magnitude larger than y , computing $x + y$ becomes imprecise because for addition, y has to be **scaled up** to have the same exponent as x , which will necessarily cut off that many lower order bits from its mantissa.
- It can even be that y is too small to make any difference at all, so that $x + y == x$. In this case, adding y to x repeatedly even a trillion times inside some loop changes nothing!
- Floating point addition is **commutative**, so that $x + y$ always equals $y + x$.
- Floating point addition is **not associative**. The result of adding three numbers x , y and z in the order $(x + y) + z$ might not produce the same end result as the order $x + (y + z)$.
- When adding lots of numbers separate numbers of different signs and orders of magnitude, you should try to arrange the order of additions to keep the operands of each individual addition within the same general magnitude. (Alternatively, you can keep track of the current cumulative error with [Kahan's algorithm](#).)
- Equality comparisons between two floating point values x and y cannot be reliably performed by testing whether `Math.abs(x-y)` is less than some strategically chosen epsilon value. First of all, a relation defined this way is not an **equivalence** but a **tolerance** relation, a totally different thing to begin with. Since tolerance relations are not transitive (see the classic **Sorites paradox**), using this relation as an equivalence silently breaks the logic of your program.
- Second, the choice of appropriate epsilon depends of the scale of x and y .